



# Chapter 7: Entity-Relationship Model

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use

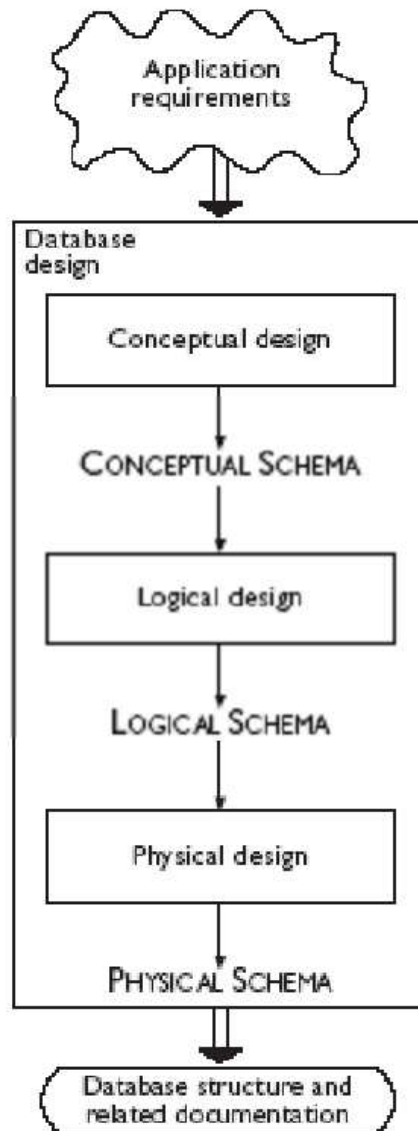


# Overview of the Design Process

- Creating a database application
  - **Design of the database schema**
  - Design of the programs that access and update the data
  - Design of a security scheme to control access to data
- Two major pitfalls to avoid in designing a database schema
  - Redundancy
    - ▶ repeating information → data inconsistency
  - Incompleteness
    - ▶ difficult or impossible to model certain aspects of the enterprise



# Main Phases of Database Design



- Requirements collection and analysis
  - Understanding the needs of users and enterprises
- Conceptual design
  - Choosing an abstract model like E-R Model
  - **Conceptual schema**: descriptions of the data requirements, entities, relationships, and constraints
- Logical design
  - Converting the abstract model to implementation model
    - ▶ E-R model to relational model
- Physical design
  - Specifying physical features of the database
    - ▶ File organization, index structures (Ch. 10 & 11)



# E-R Model

- Proposed by P. Chen in 1976
- Simple and powerful tool for the database design
  - Many database design tools draw on concepts from the E-R model
- A *database* can be modeled as:
  - a collection of **entities**
  - **relationships** among entities



# Entity Sets

- **Entity** – an object that exists and is distinguishable from other objects
  - Example: specific person, company, event, plant
  
- **Entity set** – a set of entities of the same type that share the same properties
  - Example: set of all persons, companies, trees, holidays
  
- **Attribute** – descriptive properties possessed by all members of an entity set
  - Example: people have *names* and *addresses*



# Entity Sets – *instructor* and *student*

instructor\_ID instructor\_name

76766	Crick
45565	Katz
10101	Srinivasan
98345	Kim
76543	Singh
22222	Einstein

*instructor*

student-ID student\_name

98988	Tanaka
12345	Shankar
00128	Zhang
76543	Brown
76653	Aoi
23121	Chavez
44553	Peltier

*student*



# Relationship Sets

- **Relationship** – an association among several entities

Example:

44553 (Peltier)	<u>advisor</u>	22222 (Einstein)
<i>student</i> entity	relationship set	<i>instructor</i> entity

- **Relationship set** – a mathematical relation among  $n \geq 2$  entities, each taken from entity sets

$$\{(e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n\}$$

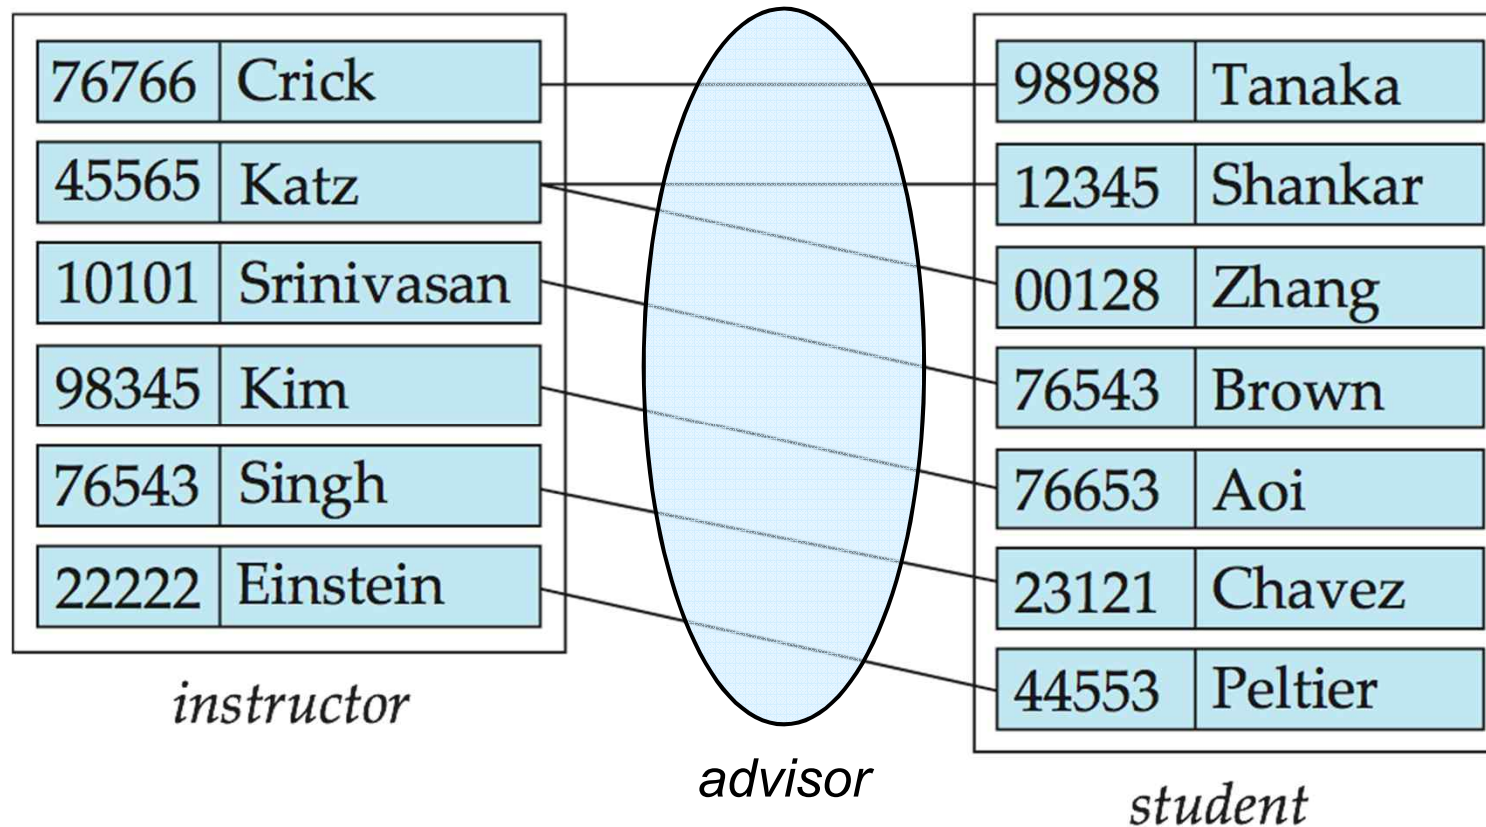
where  $(e_1, e_2, \dots, e_n)$  is a relationship

- Example:

$$(44553, 22222) \in \text{advisor}$$



## Relationship Set – *advisor*

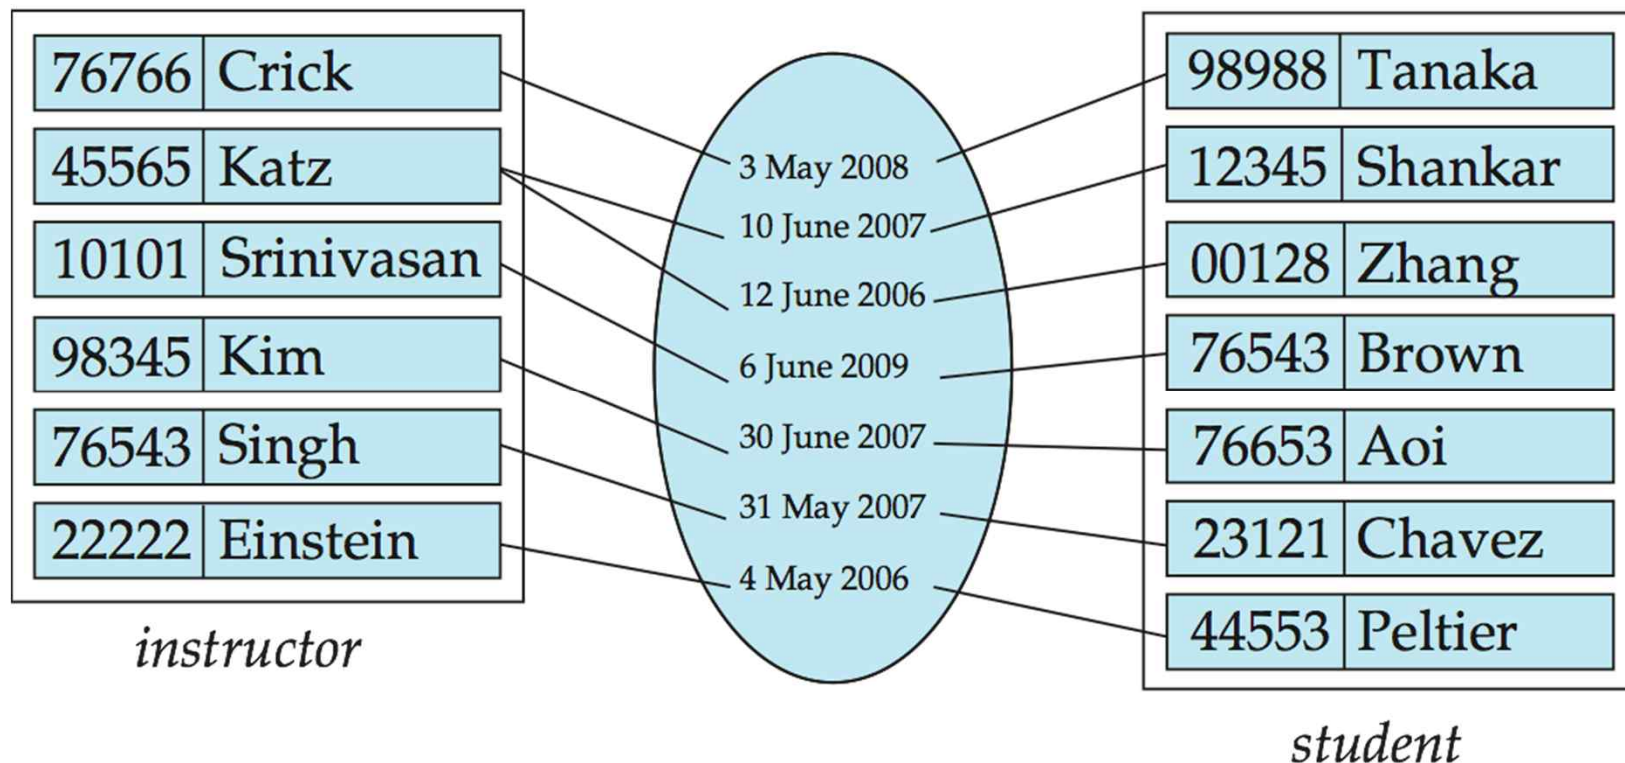






# Attribute of Relationship Set

- An **attribute** can also be property of a relationship set






# Natural Language Sentences to E-R Model

- Rules of thumb for mapping natural language descriptions into E-R model:
- Noun
  - Common noun → entity set
  - Proper noun → entity
- Verb
  - Transitive verb → relationship set
  - Intransitive verb → attribute for entity
- Adjective → attribute for entity
- Adverb → attribute for relationship




# E-R Diagrams



- Rectangles represent entity sets
- Diamonds represent relationship sets
- Lines link entity sets to relationship sets
- Attributes are listed inside entity rectangles
- Underline indicates primary key attributes



# Relationship Sets with Attributes

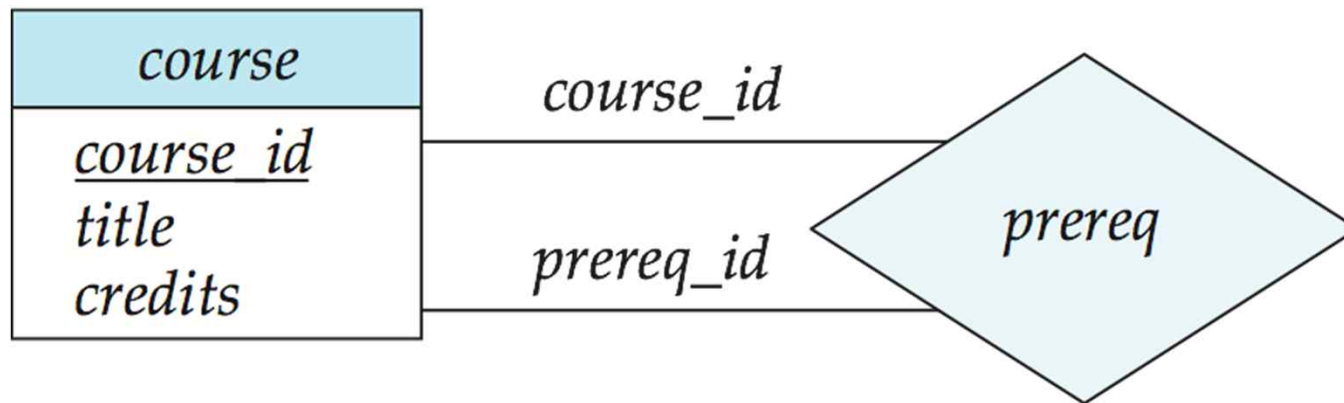


- **Dashed lines** link attributes to the relationship sets



# Roles


- Entity sets of a relationship need not be distinct
  - Each occurrence of an entity set plays a “role” in the relationship
- The labels “*course\_id*” and “*prereq\_id*” are called **roles**.






# Degree of a Relationship Set

- **Degree** of a relationship set
  - The number of entity sets that participate in the relationship
- Most relationship sets in a database system are binary



- You can define non-binary relationships





# Attribute Types

- **Simple** and **composite** attributes
  - Simple attribute: can not be divided into subparts
  - Composite attribute: composed of multiple subparts
    - ▶ Example: name = (first\_name, middle\_initial, last\_name)  
address = (street, city, state, zip\_code)
  
- **Domain** – the set of permitted values for each attribute
  - Null value: a special value meaning “missing” or “unknown”
    - ▶ Some attributes are not allowed to have null values



# Attribute Types

## ■ Single-valued and multivalued attributes

- Single-valued attribute
  - ▶ Each attribute has a single value for an entity
  - ▶ Example: *ID*, *name*, *address*
- Multivalued attribute
  - ▶ An attribute may have more than one value for an instance
  - ▶ Example: *phone\_number* = {7287, 7288}

## ■ Derived attributes

- Can be computed from other attributes
- Example: *age*, given *date\_of\_birth*

<i>instructor</i>
<u><i>ID</i></u>
<i>name</i>
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
<i>address</i>
<i>street</i>
<i>street_number</i>
<i>street_name</i>
<i>apt_number</i>
<i>city</i>
<i>state</i>
<i>zip</i>
{ <i>phone_number</i> }
<i>date_of_birth</i>
<i>age</i> ( )






# Mapping Cardinality Constraints


- Express **the number of entities** to which another entity can be associated via a relationship set
  
- For a binary relationship set, the mapping cardinality must be one of the following types:
  - One-to-one
  - One-to-many
  - Many-to-one
  - Many-to-many



# Mapping Cardinalities



One-to-one




One-to-many


Note: Some elements in  $A$  and  $B$  may not be mapped to any elements in the other set



# Mapping Cardinalities



Many-to-one




Many-to-many

Note: Some elements in  $A$  and  $B$  may not be mapped to any elements in the other set




# Mapping Cardinality Constraints in E-R diagram

- Line types between the relationship set and the entity set
  - Directed line ( $\rightarrow$ ): at most “one” (including 0)
  - Undirected line ( $\text{—}$ ): “many” (including 0)




One-to-one



One-to-many



Many-to-one




Many-to-many



# Participation Constraints


- **Total participation** (indicated by double line): every entity in the entity set participates in at least one relationship in the relationship set
  - Example: participation of *section* in *sec\_course* is total
    - ▶ every *section* must have an associated course
- **Partial participation**: some entities may not participate in any relationship in the relationship set
  - Example: participation of *instructor* in *advisor* is partial





# Cardinality Limits on Relationship Sets


- Cardinality limits can indicate more complex participation constraints





# Cardinality Constraints on $n$ -ary Relationship

- We allow **at most one arrow** out of a ternary (or greater degree) relationship
  - E.g., an arrow from *proj\_guide* to *instructor* indicates each student has at most one guide for a project



- If there is more than one arrow, there are two ways of defining the meaning.
  - E.g., a ternary relationship *R* between *A*, *B* and *C* with arrows to *B* and *C* could mean
    1. each *A* entity is associated with a unique entity from *B* and *C* or
    2. each pair of entities from (*A*, *B*) is associated with a unique *C* entity, and each pair (*A*, *C*) is associated with a unique *B*
  - To avoid confusion **we outlaw more than one arrow**




# Keys

## ■ Keys for entity sets

- A **super key** of an entity set is a set of one or more attributes whose values uniquely determine each entity.
- A **candidate key** of an entity set is a **minimal super key**
  - ▶ *ID* is candidate key of *instructor*
- Although several candidate keys may exist, one of the candidate keys is selected to be the **primary key**.

## ■ Keys for relationship sets

- The combination of primary keys of the participating entity sets forms a super key of a relationship set.
  - ▶ (*s\_id*, *i\_id*) is the super key of *advisor*







# Keys for Relationship Sets

- Must consider the mapping cardinality of the relationship set when deciding what are the candidate keys (primary keys)


Let  $R$  be a relationship set involving entity sets  $E_1, E_2, \dots, E_n$ .

- Primary keys for binary relationship set ( $n = 2$ )
  - Many-to-many:  $PK(R) = PK(E_1) \cup PK(E_2)$
  - Many-to-one/one-to-many:  $PK(R) = PK(\text{"many"-side entity})$
  - One-to-one:  $PK(R) = PK(E_1)$  or  $PK(E_2)$
- Primary keys for  $n$ -ary relationship set
  - No arrow edges:  $PK(R) = PK(E_1) \cup PK(E_2) \cup \dots \cup PK(E_n)$
  - With an arrow edge:  $PK(R) = PKs \text{ of the entity sets not on the "arrow"-side}$
- If the relationship set  $R'$  is the relationship  $R$  with attributes  $\{a_1, \dots, a_m\}$ 
  - $PK(R') = PK(R) \cup \{a_1, \dots, a_m\}$



# Removing Redundant Attributes

- Suppose we have entity sets




- Attribute *dept\_name* in entity *instructor* is redundant
  - The attribute replicates information present in the relationship *inst\_dept*, and should be removed from *instructor*
  - BUT: when converting back to tables, in some cases the attribute gets reintroduced, as we will see.



# Weak Entity Sets

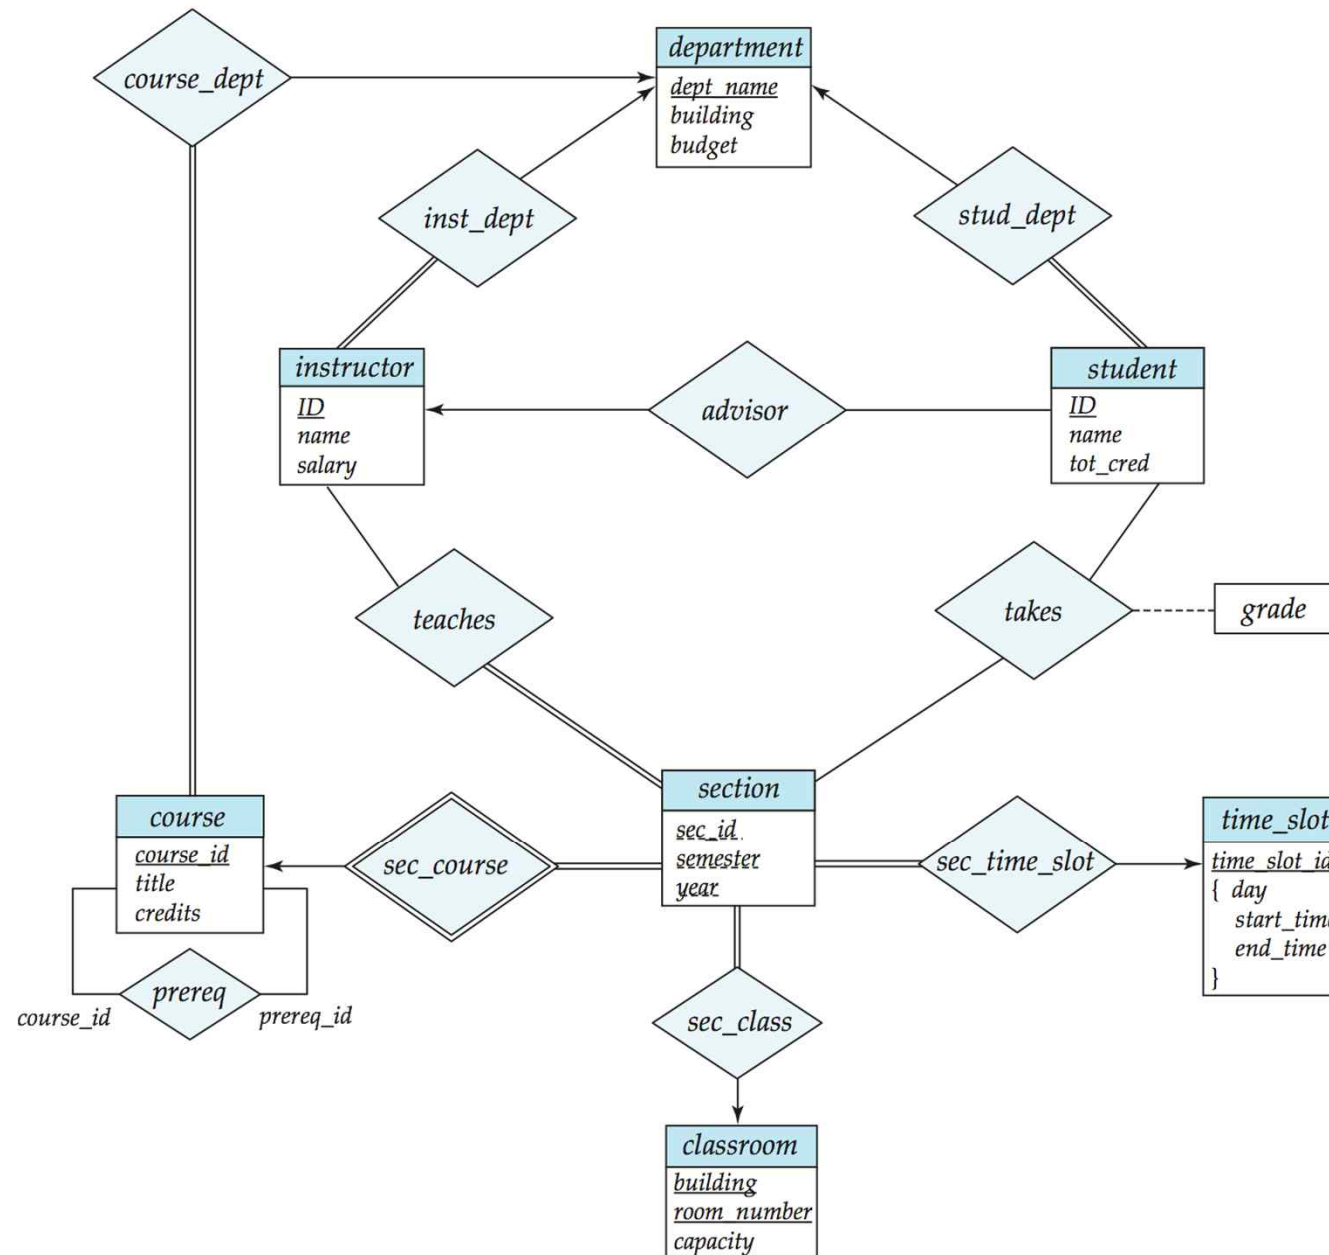
- **Weak entity set**: an entity set that does not have a primary key



- The existence of a weak entity set depends on the existence of an **identifying entity set**
  - It must relate to the identifying entity set via a total, one-to-many relationship set from the identifying to the weak entity set
  - **Identifying relationship** depicted using a **double diamond**
- The **discriminator** (or *partial key*) of a weak entity set is the set of attributes that distinguishes among all the entities of a weak entity set.
  - We **underline** the discriminator of a weak entity set with a **dashed line**.
- The **primary key** of a weak entity set = (the primary key of identifying strong entity set)+ (the weak entity set's discriminator)



# E-R Diagram for a University Enterprise





# Exercise

- Construct an E-R diagram for the following company enterprise.
  - A company has many employees. We store each employee's name, SSN, address, salary, gender, and birth date.
  - The company is organized into departments. Each department has a unique name, a unique number, and a particular employee who manages the department. The department may have several locations.
  - An employee works for one department, and we keep track of the employee assignments. We also keep track of the direct supervisor of each employee.
  - Also, we want to keep track of the dependents of each employee for insurance purposes. We keep each dependent's name, gender, birth date, and relationship to the employee.






# Reduction to Relation Schemas

- Entity sets and relationship sets can be expressed uniformly as *relation schemas* that represent the contents of the database.
- A database which conforms to an E-R diagram can be represented by a collection of schemas.
- For each entity set and relationship set there is a unique schema that is assigned the name of the corresponding entity set or relationship set.
- Each schema has a number of columns (generally corresponding to attributes), which have unique names.



# Representing Entity Sets With Simple Attributes

- A **strong entity set** reduces to a schema with the same attributes  
*student*(ID, name, tot\_cred)
- A **weak entity set** becomes a table that includes a column for the primary key of the identifying strong entity set  
*section* ( course\_id, sec\_id, sem, year )








# Representing Relationship Sets

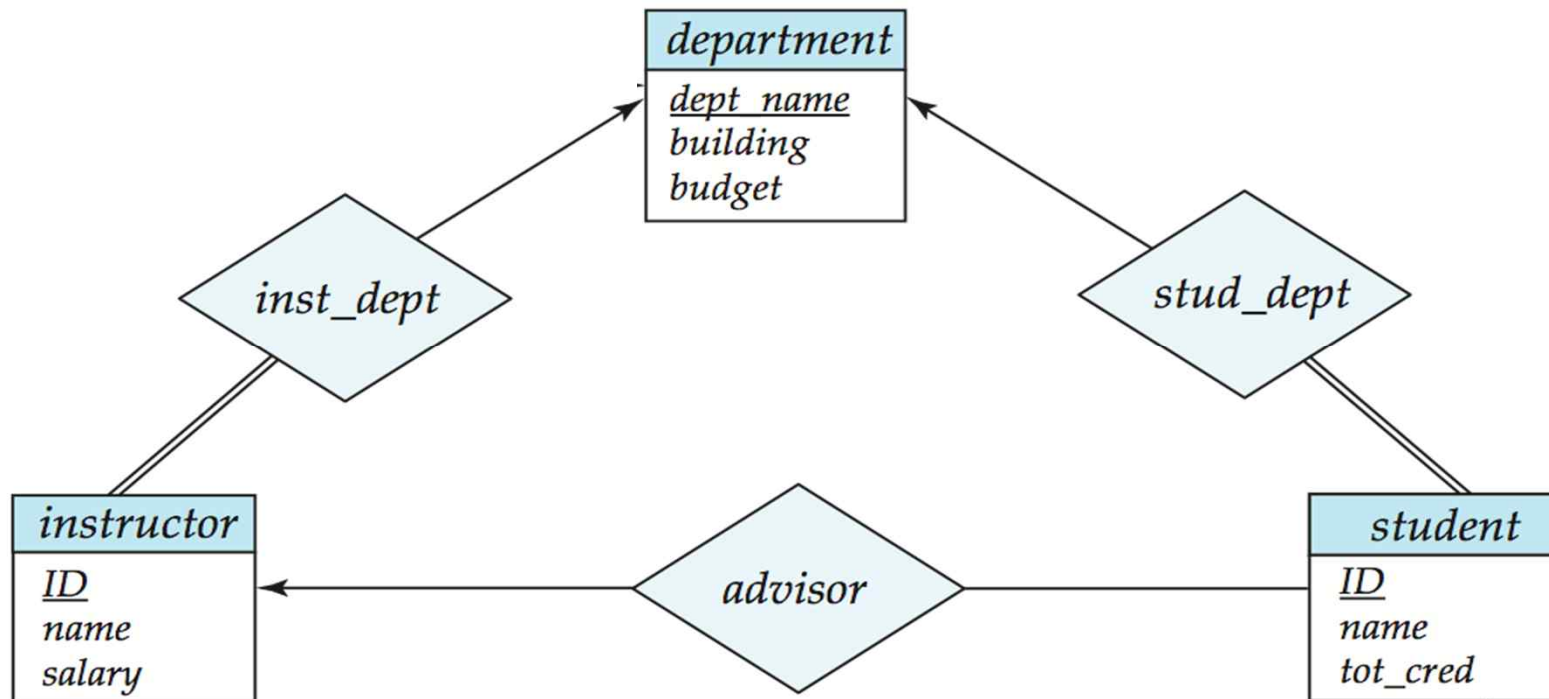
- A relationship set is represented as a schema with attributes for the primary keys of the participating entity sets, and any descriptive attributes of the relationship set.
- Example: schema for relationship set *advisor*  
 $advisor = (\underline{s\_id}, \underline{i\_id})$





# Redundancy of Schemas

- Many-to-one and one-to-many relationship sets that are total on the many-side can be represented by **adding an extra attribute to the “many” side**, containing the primary key of the “one” side
- Example: Instead of creating a schema for relationship set *inst\_dept*, add an attribute *dept\_name* to the schema arising from entity set *instructor*





## Redundancy of Schemas (Cont.)

- For one-to-one relationship sets, either side can be chosen to act as the “many” side
  - That is, extra attribute can be added to either of the tables corresponding to the two entity sets
- If participation is *partial* on the “many” side, replacing a schema by an extra attribute in the schema corresponding to the “many” side could result in null values
- The schema corresponding to a relationship set linking a weak entity set to its identifying strong entity set is redundant.
  - Example: The *section* schema already contains the attributes that would appear in the *sec\_course* schema



# Composite and Multivalued Attributes

<i>instructor</i>
<u>ID</u>
<i>name</i>
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
<i>address</i>
<i>street</i>
<i>street_number</i>
<i>street_name</i>
<i>apt_number</i>
<i>city</i>
<i>state</i>
<i>zip</i>
{ <i>phone_number</i> }
<i>date_of_birth</i>
<i>age</i> ( )

- **Composite attributes** are flattened out by creating a separate attribute for each component attribute
- Ignoring multivalued attributes, extended instructor schema is
  - *instructor*(ID, *first\_name*, *middle\_initial*, *last\_name*, *street\_number*, *street\_name*, *apt\_number*, *city*, *state*, *zip\_code*, *date\_of\_birth*)




# Composite and Multivalued Attributes

- A multivalued attribute  $M$  of an entity  $E$  is represented by a separate schema  $EM$ 
  - Schema  $EM$  has attributes corresponding to the primary key of  $E$  and an attribute corresponding to multivalued attribute  $M$
  - Example: Multivalued attribute *phone\_number* of *instructor* is represented by a schema:  
 $inst\_phone = ( \underline{ID}, \underline{phone\_number} )$
  - Each value of the multivalued attribute maps to a separate tuple of the relation on schema  $EM$ 
    - ▶ For example, an *instructor* entity with primary key 22222 and phone numbers 456-7890 and 123-4567 maps to two tuples:  
(22222, 456-7890) and (22222, 123-4567)



## Multivalued Attributes (Cont.)

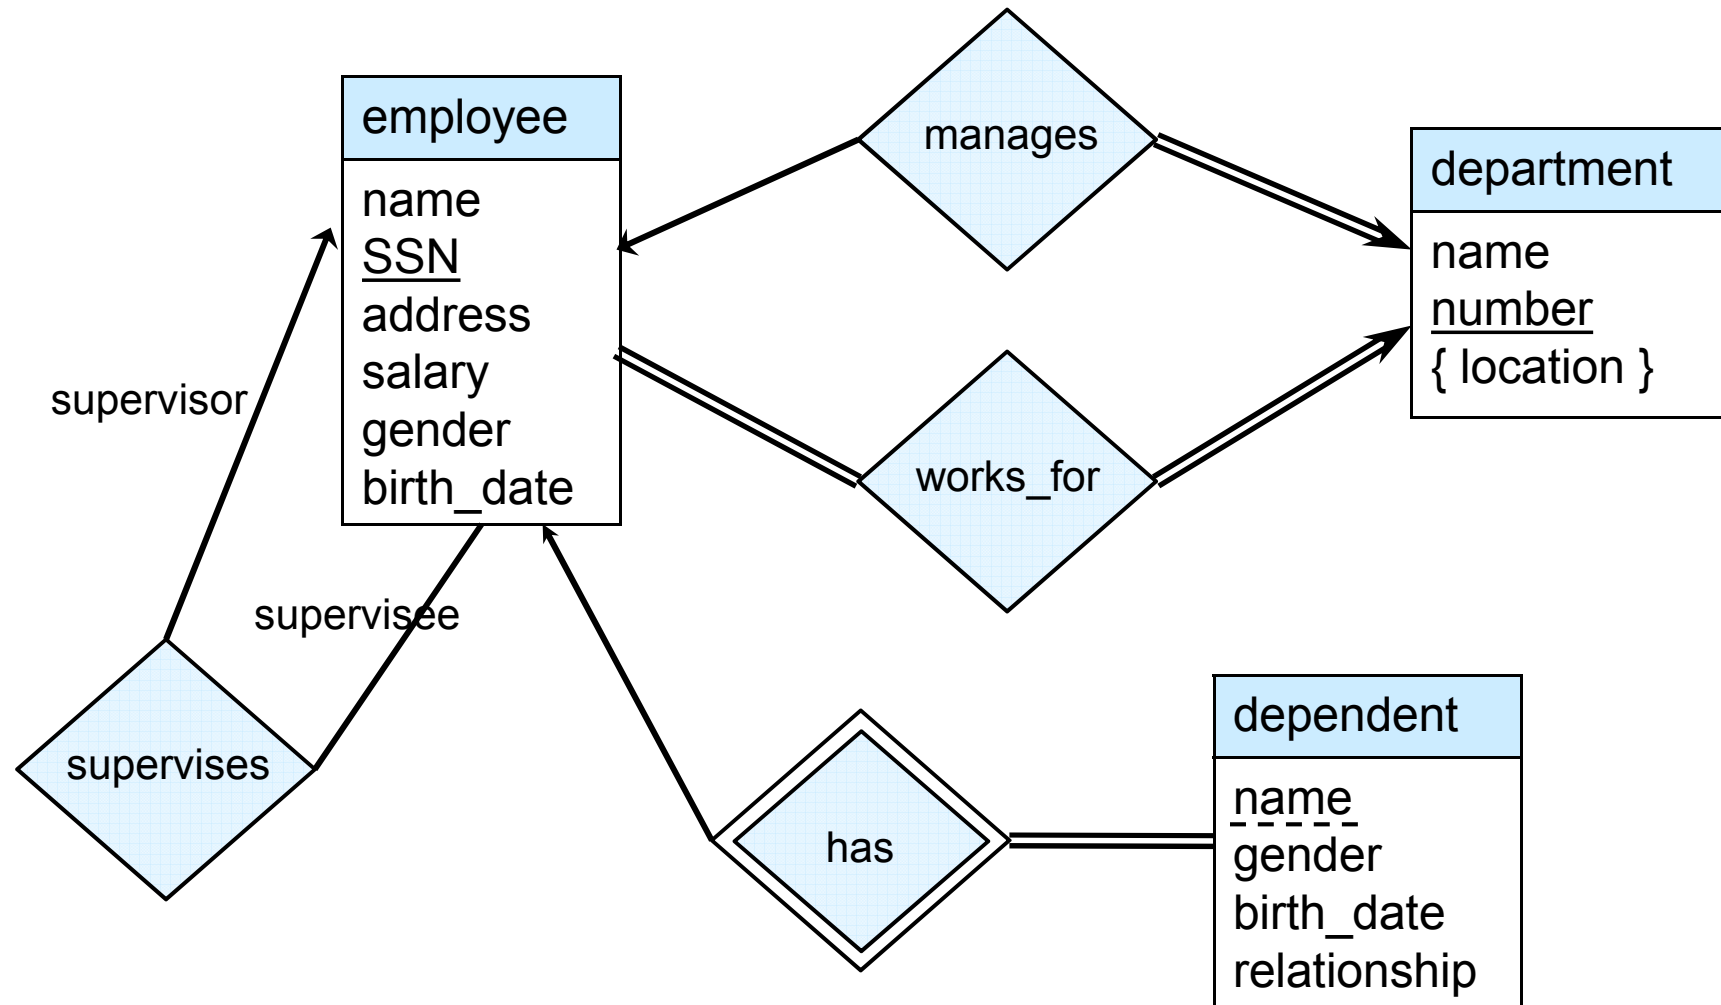
- Special case: entity *time\_slot* has only one attribute other than the primary-key attribute, and that attribute is multivalued
  - Optimization: Don't create the relation corresponding to the entity, just create the one corresponding to the multivalued attribute
  - *time\_slot*(*time\_slot\_id*, *day*, *start\_time*, *end\_time*)
  - Caveat: *time\_slot* attribute of *section* (from *sec\_time\_slot*) cannot be a foreign key due to this optimization





# Exercise

- Convert the following E-R diagram into a set of relations.









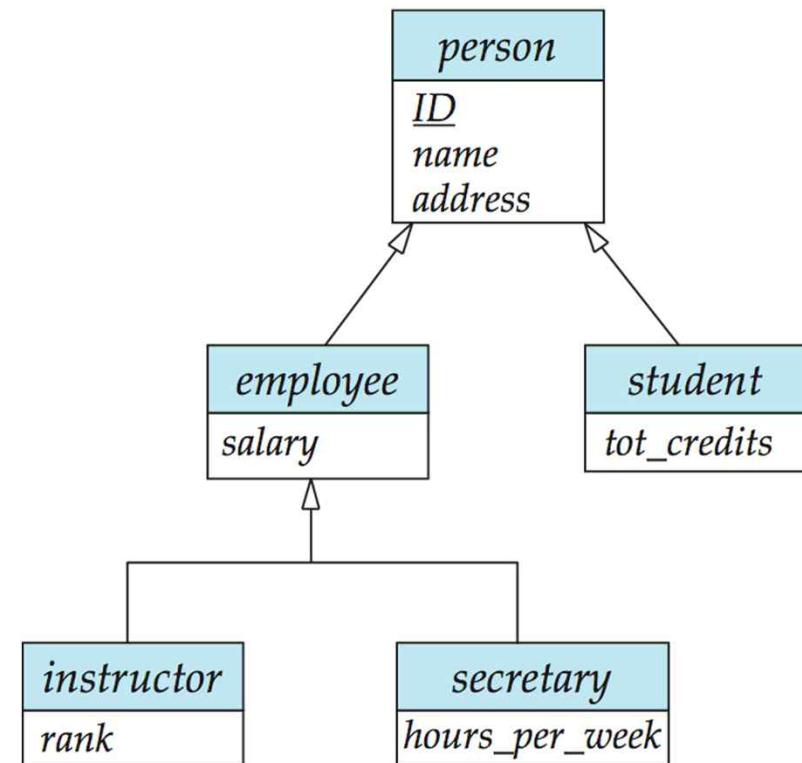
# Extended E-R Features: Specialization/Generalization

## ■ Specialization

- Designating subgroupings within an entity set
- **Top-down** design process

## ■ Generalization

- Combining a number of entity sets that share the same features into a higher-level entity set
- **Bottom-up** design process



- **Attribute inheritance** – a lower-level entity set inherits all the attributes and relationship participation of the higher-level entity set to which it is linked.
- Lower-level entity sets may have their own specific attributes or participate in relationships that do not apply to the higher-level entity set.



# Constraints on Specialization/Generalization

- **Disjoint constraint** – specifies whether or not entities may belong to more than one lower-level entity set within a single generalization/specialization
  - **Disjoint** (denoted by a single arrow)
    - ▶ An entity can belong to only one lower-level entity set
  - **Overlapping** (denoted by separate arrows)
    - ▶ An entity can belong to more than one lower-level entity set
  
- **Completeness constraint** -- specifies whether or not an entity in the higher-level entity set must belong to at least one of the lower-level entity sets within a generalization/specialization
  - **Total** (denoted by a dashed line and the keyword “total”)
    - ▶ an entity must belong to one of the lower-level entity sets
  - **Partial**(default)
    - ▶ an entity need not belong to one of the lower-level entity sets



# Reduction to Relation Schemas

## ■ Method 1:

- Form a schema for the higher-level entity
- Form a schema for each lower-level entity set, include primary key of higher-level entity set and local attributes

schema	attributes
<i>person</i>	<i>ID, name, street, city</i>
<i>student</i>	<i>ID, tot_cred</i>
<i>employee</i>	<i>ID, salary</i>

- Drawback: getting information about, an *employee* requires accessing two relations, the one corresponding to the low-level schema and the one corresponding to the high-level schema



# Reduction to Relation Schemas (Cont.)

## ■ Method 2:

- Form a schema for each entity set with all local and inherited attributes

schema	attributes
<i>person</i>	<i>ID, name, street, city</i>
<i>student</i>	<i>ID, name, street, city, tot_cred</i>
<i>employee</i>	<i>ID, name, street, city, salary</i>

- If specialization is total, the schema for the generalized entity set (*person*) not required to store information
  - ▶ Can be defined as a “view” relation containing union of specialization relations
  - ▶ But explicit schema may still be needed for foreign key constraints
- Drawback: *name, street* and *city* may be stored redundantly for people who are both students and employees



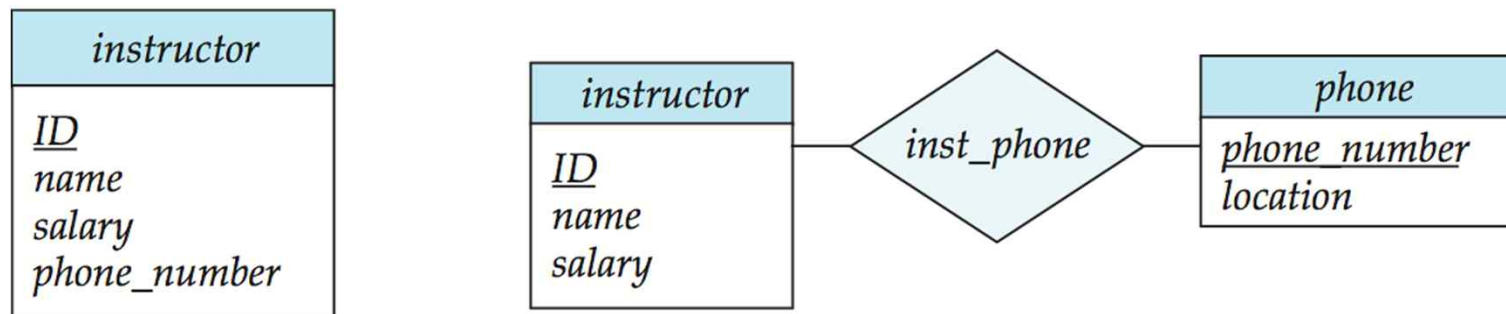
# Entity-Relationship Design Issues

- The use of an attribute or an entity set to represent an object
- The use of an entity sets or an relationship sets to represent an object
- The use of a ternary relationship versus a pair of binary relationships



# Use of Attributes vs. Entity Sets

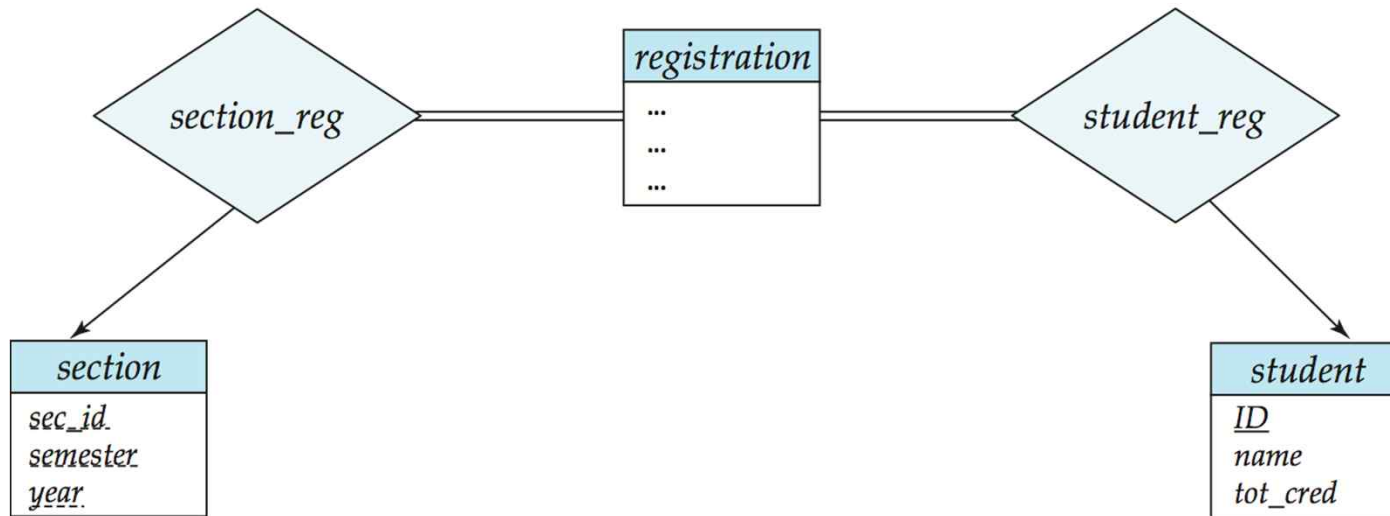
- Whether the entity must be treated as an independent entity
  - Whether to have multiple entities
  - Whether to keep extra information about the entity
- Example: use of phone as an entity allows extra information about phone numbers (plus multiple phone numbers)





# Use of Entity Sets vs. Relationship Sets

- Use of entity sets: keeping other information about the entity



- Use of relationship sets: more compact



- It is not always clear – possible guideline is to designate a relationship set to describe an action that occurs between entities



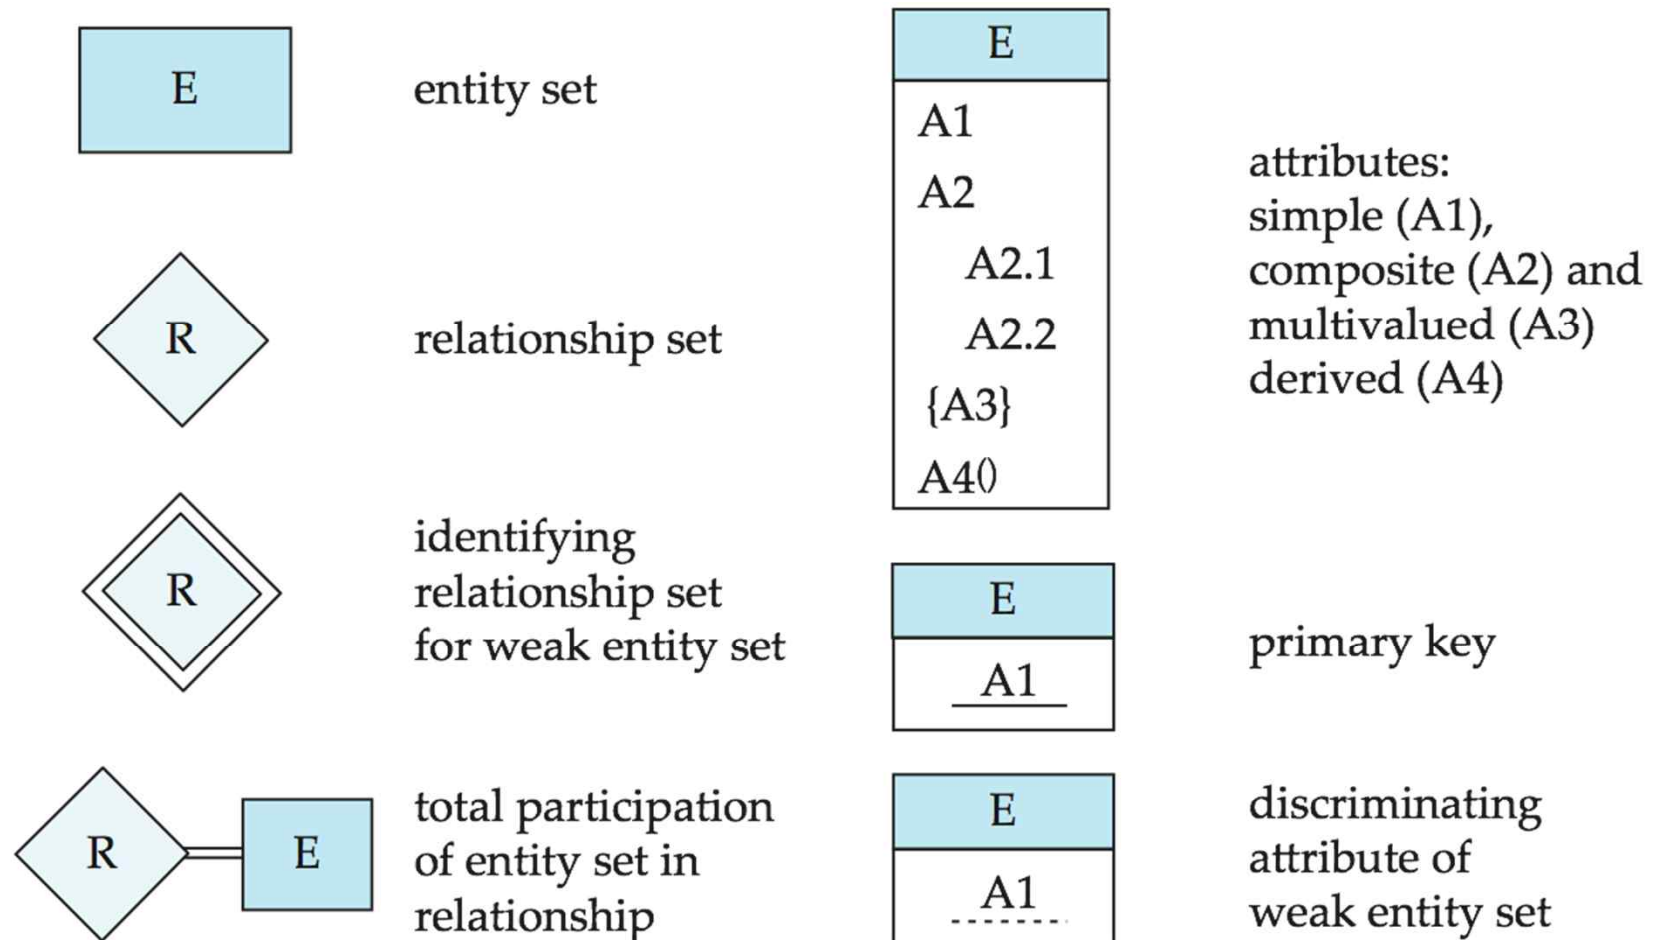
# Binary vs. Non-Binary Relationships

- In general, any non-binary relationship can be represented using binary relationships by **creating an artificial entity set**.
- Some relationships that appear to be non-binary may be better represented using binary relationships
  - E.g., A ternary relationship *parents*, relating a child to his/her father and mother, is best replaced by two binary relationships, *father* and *mother*
    - ▶ Using two binary relationships allows partial information (e.g., only mother being known)
- But there are some relationships that are naturally non-binary
  - Example: *proj\_guide*



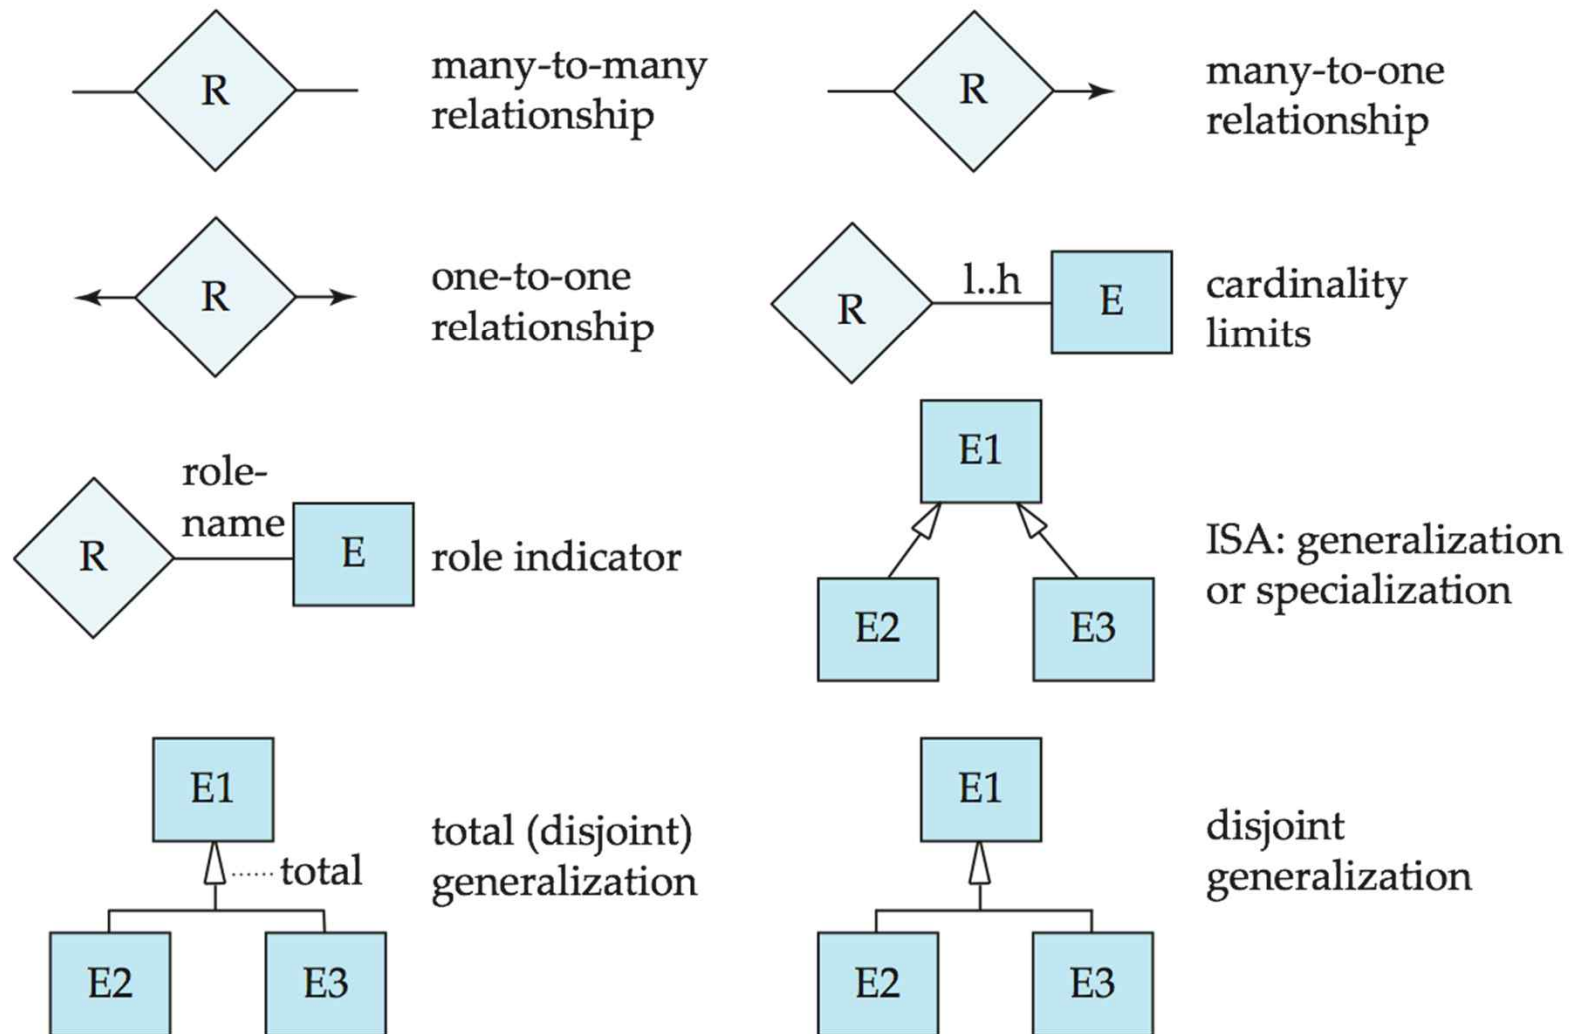


# Summary of Symbols Used in E-R Notation





# Symbols Used in E-R Notation (Cont.)





# End of Chapter 7

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use