



# Chapter 14: Transactions

## Database System Concepts, 6<sup>th</sup> Ed.

- Chapter 1: Introduction
- **Part 1: Relational databases**
  - Chapter 2: Introduction to the Relational Model
  - Chapter 3: Introduction to SQL
  - Chapter 4: Intermediate SQL
  - Chapter 5: Advanced SQL
  - Chapter 6: Formal Relational Query Languages
- **Part 2: Database Design**
  - Chapter 7: Database Design: The E-R Approach
  - Chapter 8: Relational Database Design
  - Chapter 9: Application Design
- **Part 3: Data storage and querying**
  - Chapter 10: Storage and File Structure
  - Chapter 11: Indexing and Hashing
  - Chapter 12: Query Processing
  - Chapter 13: Query Optimization
- **Part 4: Transaction management**
  - Chapter 14: Transactions
  - Chapter 15: Concurrency control
  - Chapter 16: Recovery System
- **Part 5: System Architecture**
  - Chapter 17: Database System Architectures
  - Chapter 18: Parallel Databases
  - Chapter 19: Distributed Databases
- **Part 6: Data Warehousing, Mining, and IR**
  - Chapter 20: Data Mining
  - Chapter 21: Information Retrieval
- **Part 7: Specialty Databases**
  - Chapter 22: Object-Based Databases
  - Chapter 23: XML
- **Part 8: Advanced Topics**
  - Chapter 24: Advanced Application Development
  - Chapter 25: Advanced Data Types
  - Chapter 26: Advanced Transaction Processing
- **Part 9: Case studies**
  - Chapter 27: PostgreSQL
  - Chapter 28: Oracle
  - Chapter 29: IBM DB2 Universal Database
  - Chapter 30: Microsoft SQL Server
- **Online Appendices**
  - Appendix A: Detailed University Schema
  - Appendix B: Advanced Relational Database Model
  - Appendix C: Other Relational Query Languages
  - Appendix D: Network Model
  - Appendix E: Hierarchical Model



# Chapter 14: Transactions

- Transaction Concept
- A Simple Transaction Model
- Storage Structure
- Transaction Atomicity and Durability
- Transaction Isolation
- Serializability
- Transaction Isolation and Atomicity
- Transaction Isolation Level
- Implementation of Isolation Levels
- Transaction Definition in SQL



# Transaction Concept

- A **transaction** is a *unit* of program execution that accesses and possibly updates various data items
- E.g. transaction to transfer \$50 from account A to account B:
  1. **read(A)**
  2.  $A := A - 50$
  3. **write(A)**
  4. **read(B)**
  5.  $B := B + 50$
  6. **write(B)**
- Two main issues to deal with:
  - **Failures of various kinds**, such as hardware failures and system crashes
  - Concurrent execution of multiple transactions



# ACID Properties

- A **transaction** is a unit of program execution that accesses and possibly updates various data items
- To preserve the integrity of data the database system must ensure
  - **Atomicity:** Either **all** operations of the transaction are properly reflected in the database or **none** are
  - **Consistency:** Execution of a transaction in isolation preserves the consistency of the database
  - **Isolation:** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions
    - ▶ Intermediate transaction results must be **hidden** from other concurrently executed transactions
    - ▶ That is, for every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  that **either**  $T_j$  finished execution before  $T_i$  started, **or**  $T_j$  started execution after  $T_i$  finished
  - **Durability:** After a transaction completes successfully, the changes it has made to the database **persist**, even if there are system failures



# Chapter 14: Transactions

- Transaction Concept
- A Simple Transaction Model
- Storage Structure
- Transaction Atomicity and Durability
- Transaction Isolation
- Serializability
- Transaction Isolation and Atomicity
- Transaction Isolation Level
- Implementation of Isolation Levels
- Transaction Definition in SQL



# A Simple Transaction Model [1/3]

- We ignore operations other than **read** and **write** instructions
  - We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes
- Transaction to transfer \$50 from account A to account B:

1. **read(A)**
2.  $A := A - 50$
3. **write(A)**
4. **read(B)**
5.  $B := B + 50$
6. **write(B)**

- **ACID Properties**
  - **Atomicity requirement**
  - **Consistency requirement**
  - **Isolation requirement**
  - **Durability requirement**



# A Simple Transaction Model [2/3]

- Transaction to transfer \$50 from account A to account B:

1. **read(A)**
2.  $A := A - 50$
3. **write(A)**
4. **read(B)**
5.  $B := B + 50$
6. **write(B)**

- **Atomicity requirement**

- If the transaction fails after step 3 and before step 6, money will be “lost” leading to an inconsistent database state (Failure could be due to SW or HW)
- The system should ensure that updates of a partially executed transaction are not reflected in the database

- **Consistency requirement** in above example:

- A transaction must see a consistent database
  - ▶ Explicitly specified integrity constraints such as primary keys and foreign keys
  - ▶ Implicit integrity constraints
    - The sum of A and B is unchanged by the execution of the transactions
- During transaction execution the database may be temporarily inconsistent
- When the transaction completes successfully the database must be consistent
  - ▶ Erroneous transaction logic can lead to inconsistency



# A Simple Transaction Model [3/3]

- **Isolation requirement** — If between steps 3 and 6, another transaction T2 is allowed to access [the partially updated database](#), it will see an inconsistent database (the sum  $A + B$  will be less than it should be)

T1	T2
1. <b>read(A)</b>	
2. $A := A - 50$	
3. <b>write(A)</b>	read(A), read(B), print(A+B)
4. <b>read(B)</b>	
5. $B := B + 50$	
6. <b>write(B)</b>	

- Isolation can be ensured trivially by running transactions **serially** (that is, one after the other)
  - [However, executing multiple transactions concurrently has significant benefits](#), as we will see later
- 
- **Durability requirement** — once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), the updates to the database by the transaction must persist even if there are software or hardware failures



# Chapter 14: Transactions

- Transaction Concept
- A Simple Transaction Model
- Storage Structure
- Transaction Atomicity and Durability
- Transaction Isolation
- Serializability
- Transaction Isolation and Atomicity
- Transaction Isolation Level
- Implementation of Isolation Levels
- Transaction Definition in SQL



# Storage Structure

## ■ Volatile storage:

- does not survive system crashes
- examples: main memory, cache memory

## ■ Nonvolatile storage:

- survives system **crashes**
- examples: disk, tape, flash memory, non-volatile (battery backed up) RAM
- but may still fail, losing data

## ■ Stable storage:

- a **mythical form of storage** that survives **all failures**
  - approximated by maintaining **multiple copies on distinct nonvolatile media**
  - See book for more details on how to implement stable storage
- 
- For a transaction to be **durable**, data changes need to be written to stable storage
  - For a transaction to be **atomic**, log records need to be written to stable storage before any changes are made to the database on disk



# Chapter 14: Transactions

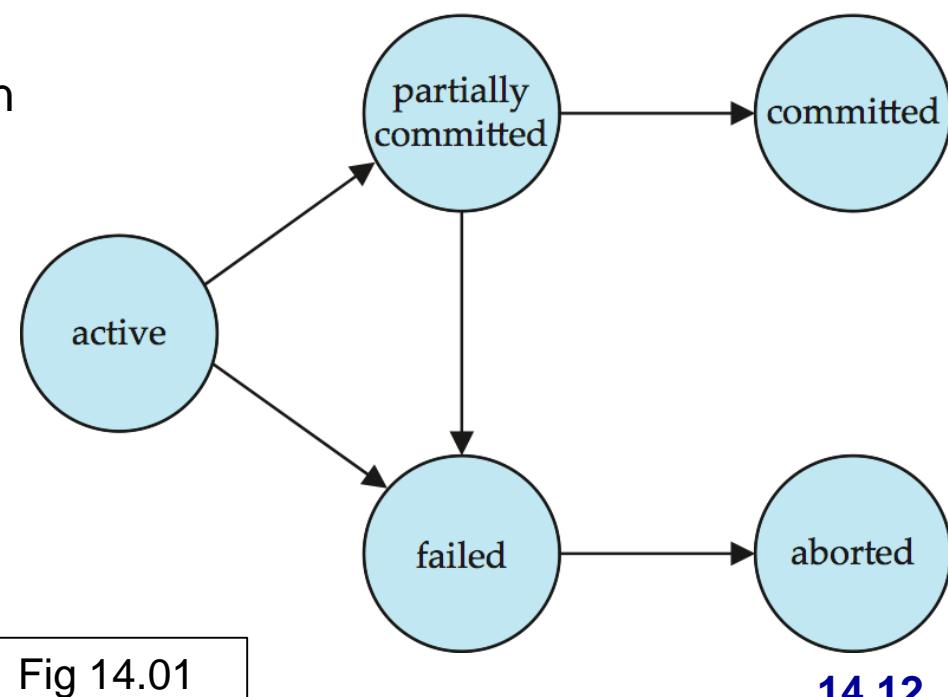
- Transaction Concept
- A Simple Transaction Model
- Storage Structure
- Transaction Atomicity and Durability
- Transaction Isolation
- Serializability
- Transaction Isolation and Atomicity
- Transaction Isolation Level
- Implementation of Isolation Levels
- Transaction Definition in SQL



# Transaction State

- **Active:** the initial state; the transaction stays in this state while it is executing
- **Partially committed:** after the final statement has been executed
- **Failed:** after the discovery that normal execution can no longer proceed
- **Aborted:** after the transaction has been rolled back and the database restored to its state prior to the start of the transaction
  - Two options after it has been aborted:
    - ▶ restart the transaction (can be done only if no internal logical error)
    - ▶ kill the transaction
- **Committed:** after successful completion

Maintaining a log is critical!





# Chapter 14: Transactions

- Transaction Concept
- A Simple Transaction Model
- Storage Structure
- Transaction Atomicity and Durability
- **Transaction Isolation**
- Serializability
- Transaction Isolation and Atomicity
- Transaction Isolation Level
- Implementation of Isolation Levels
- Transaction Definition in SQL



# Concurrent Executions

- Multiple transactions are allowed to **run concurrently** in the system
- Advantages are:
  - **increased processor and disk utilization**, leading to better transaction *throughput*
    - ▶ E.g. one transaction can be using the CPU while another is reading from or writing to the disk
  - **reduced average response time** for transactions: short transactions need not wait behind long ones
- **Concurrency control schemes** – mechanisms to achieve **isolation**
  - that is, to control the interaction among the concurrent transactions in order to prevent them from destroying **the consistency of the database**
    - ▶ Will study various CC schemes in Chapter 16, after studying notion of correctness of concurrent executions



# Schedules

- **Schedule:** a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
  - A schedule for a set of transactions must consist of all instructions of those transactions
  - A schedule must preserve the order in which the instructions appear in each individual transaction
- Definition: Serial schedule is instruction sequences from **one by one** transactions
- A transaction that successfully completes its execution will have a **commit instruction** as the last statement
  - by default transaction assumed to execute commit instruction as its last step
- A transaction that fails to successfully complete its execution will have **an abort instruction** as the last statement



# Schedule 1 and Schedule 2

- Let  $T_1$  transfer \$50 from  $A$  to  $B$ , and  $T_2$  transfer 10% of the balance from  $A$  to  $B$ .
- Serial schedules:  $\{ T_1 ; T_2 \}$        $\{ T_2 ; T_1 \}$

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

Fig 14.02

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

Fig 14.03



# Schedule 3

- Let  $T_1$  and  $T_2$  be the transactions defined previously
- The following schedule is not a serial schedule, but it is *equivalent* to Schedule 1
- We call it a *serializable schedule*

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ )	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ )
read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

Fig 13.04

In Schedules 1, 2 and 3, the sum  $A + B$  is preserved.



# Schedule 4

- The following concurrent schedule does not preserve the value of  $(A + B)$
- The following schedule is not serializable

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ )
write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	$B := B + temp$ write ( $B$ ) commit

Fig 14.05



# Chapter 14: Transactions

- Transaction Concept
- A Simple Transaction Model
- Storage Structure
- Transaction Atomicity and Durability
- Transaction Isolation
- **Serializability**
- Transaction Isolation and Atomicity
- Transaction Isolation Level
- Implementation of Isolation Levels
- Transaction Definition in SQL



# Serializability

- **Basic Assumption** – Each transaction preserves database consistency
- Thus serial execution of a set of transactions preserves database consistency
- **Definition:** A (possibly concurrent) schedule is **serializable** if it is equivalent to a serial schedule
- Different forms of schedule equivalence give rise to the notions of:
  1. **conflict serializability**
  2. **view serializability**



# Conflicting Instructions

- Instructions  $I_i$  and  $I_j$  of transactions  $T_i$  and  $T_j$  respectively, **conflict** if and only if there exists some item  $Q$  accessed by both  $I_i$  and  $I_j$ , and **at least** one of these instructions wrote  $Q$ 
  1.  $I_i = \text{read}(Q)$ ,  $I_j = \text{read}(Q)$  :  $I_i$  and  $I_j$  don't conflict
  2.  $I_i = \text{read}(Q)$ ,  $I_j = \text{write}(Q)$  : They conflict
  3.  $I_i = \text{write}(Q)$ ,  $I_j = \text{read}(Q)$  : They conflict
  4.  $I_i = \text{write}(Q)$ ,  $I_j = \text{write}(Q)$  : They conflict
- Intuitively, a conflict between  $I_i$  and  $I_j$  forces a (logical) temporal order between them
  - If  $I_i$  and  $I_j$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been **interchanged** in the schedule
  - Intuition
    - ▶ 같은 data item을 access하는 instruction 2개는 read-read 경우만 순서를 바꿀수 된다
    - ▶ 서로 다른 data item을 access하는 instruction 2개는 순서를 언제든 바꿀수 된다



# Conflict Serializability [1/2]

- If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are **conflict equivalent**
- Definition: A schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule
- Example of a schedule that is **not** conflict serializable:

$T_3$	$T_4$
read(Q)	
write(Q)	write(Q)

We are **unable to swap instructions** in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$

$T_3$	$T_4$
read(Q)	
write(Q)	write(Q)

$T_3$	$T_4$
	write(Q)
read(Q)	



# Conflict Serializability

[2/2]

- Schedule 3 can be transformed into Schedule 6, a serial schedule where  $T_2$  follows  $T_1$ , by a series of swaps of non-conflicting instructions
  - Therefore Schedule 3 is conflict serializable

**Schedule 3**

$T_1$	$T_2$
read( $A$ )	
write( $A$ )	
	read( $A$ )
	write( $A$ )
	read( $B$ )
	write( $B$ )
	read( $B$ )
	write( $B$ )

**Schedule 5 –**  
After Swapping a Pair of non  
conflicting Instructions in  
schedule 3

$T_1$	$T_2$
read( $A$ )	
write( $A$ )	
	<u>read(<math>B</math>)</u>
	<u>write(<math>A</math>)</u>
	write( $B$ )
	read( $B$ )
	write( $B$ )

**Schedule 6 –**  
A Serial Schedule That is  
Equivalent to Schedule 3

$T_1$	$T_2$
read( $A$ )	
write( $A$ )	
read( $B$ )	
write( $B$ )	
	read( $A$ )
	write( $A$ )
	read( $B$ )
	write( $B$ )



# View Serializability

[1/4]

- Let  $S$  and  $S'$  be two schedules with the same set of transactions
- $S$  and  $S'$  are **view equivalent** if the following 3 conditions are met, for each data item  $Q$ ,
  1. If in schedule  $S$ , transaction  $T_i$  reads the initial value of  $Q$ , then in schedule  $S'$  also transaction  $T_i$  must read the initial value of  $Q$
  2. If in schedule  $S$  transaction  $T_i$  executes **read( $Q$ )**, and that value was produced by transaction  $T_j$  (if any), then in schedule  $S'$  also transaction  $T_i$  must read the value of  $Q$  that was produced by the same **write( $Q$ )** operation of transaction  $T_j$
  3. The transaction (if any) that performs **the final write( $Q$ ) operation** in schedule  $S$  must also perform **the final write( $Q$ ) operation** in schedule  $S'$

As can be seen, view equivalence is also based purely on **reads** and **writes** alone



# View Serializability [2/4]

Schedule A

$T_3$	$T_4$
read ( $Q$ )	
write ( $Q$ )	write ( $Q$ )

Schedule B

$T_3$	$T_4$
read( $Q$ )	
write( $Q$ )	
	write( $Q$ )

Schedule C

$T_3$	$T_4$
	write( $Q$ )
read( $Q$ )	write( $Q$ )

- 스케줄 A는 Conflict serializable schedule이 아님, View serializable schedule도 아님



# View Serializability [3/4]

- A schedule S is **view serializable** if it is **view equivalent** to a serial schedule
- Every conflict serializable schedule is also view serializable
- Below is a schedule which is **view-serializable** but **not** conflict serializable

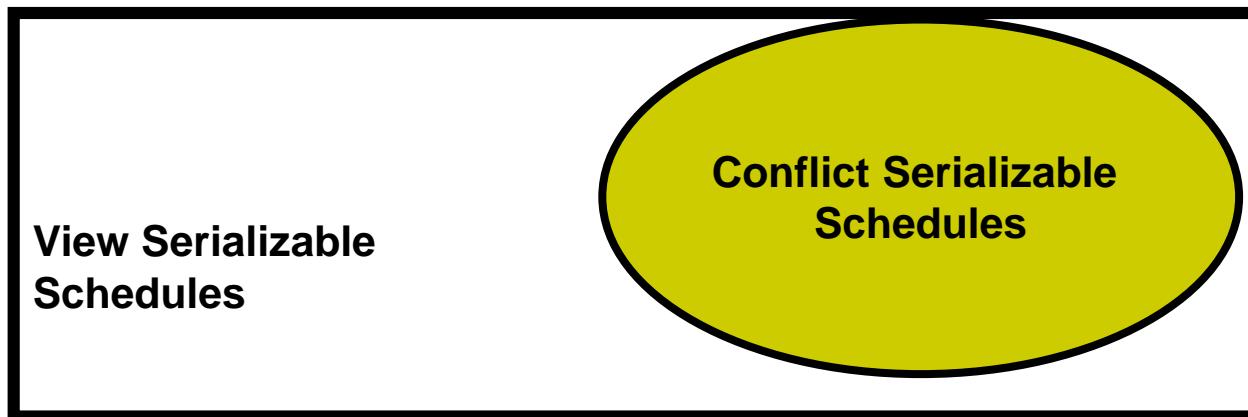
$T_{27}$	$T_{28}$	$T_{29}$
read ( $Q$ )	write ( $Q$ )	
write ( $Q$ )		write ( $Q$ )

- Non-conflicting operations를 swap할 여지가 없다
- $T_{27}$   $T_{28}$   $T_{29}$  의 어떤 combination도 만들수가 없다
- Conflict serializable schedule이 될수 없다
- 그러나 왼쪽의 schedule은  $\langle T_{27} \ T_{28} \ T_{29} \rangle$  와 View-serializability의 3가지 조건을 만족하므로 View serializable schedule이다



# View Serializability [4/4]

- Every view serializable schedule that is not conflict serializable has **blind writes**
- Intuition: 두개 schedule에서 read들이 같은값을 읽었고, final write가 같다면 Database에 반영되는 결과가 serial execution중의 한가지와 같게되는 view serializable schedule이다 → we can have more potential concurrency



이 예제에서 view equivalent가 result equivalent로 보일수 있다.  
많은 경우에 같기도 하지만 꼭 같은것은 아니다  
(다음 slide)



# Other Notions of Serializability [1/2]

- The schedule below produces same outcome as the serial schedule  $< T_1, T_5 >$ , yet is **not** conflict equivalent or **view** equivalent to it

$T_1$	$T_5$
read ( $A$ )	
$A := A - 50$	
write ( $A$ )	
	read ( $B$ )
	$B := B - 10$
	write ( $B$ )
read ( $B$ )	
$B := B + 50$	
write ( $B$ )	
	read ( $A$ )
	$A := A + 10$
	write ( $A$ )

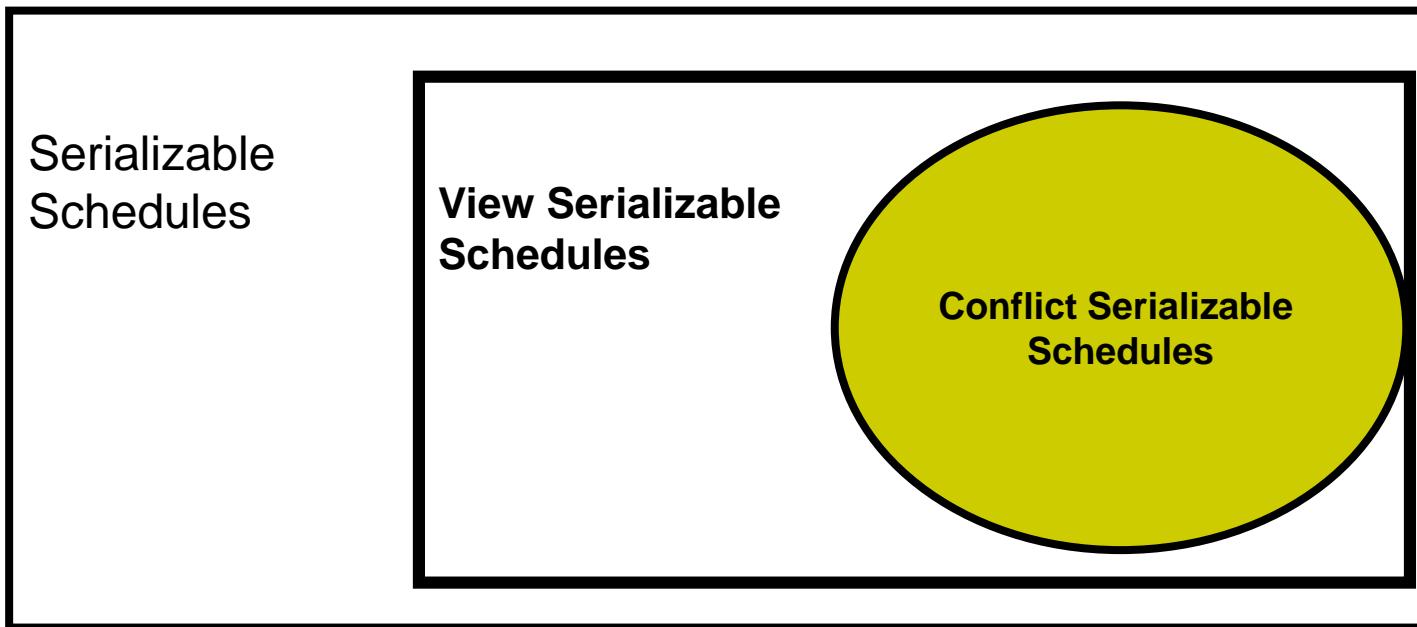
- Result Equivalent Schedule

- 우연에 의한 경우
- 실제 사용은 부적합



# Other Notions of Serializability [2/2]

- Determining such equivalence requires analysis of operations other than read and write



- Still we can pursue potential concurrency!



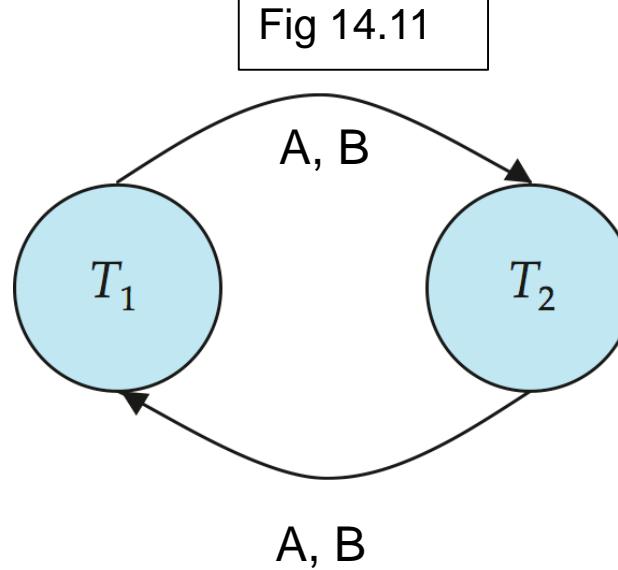
# Testing for Serializability

- Consider some schedule of a set of transactions  $T_1, T_2, \dots, T_n$
- Precedence graph** — a direct graph where the vertices are the transactions (names)
  - We draw **an arc from  $T_i$  to  $T_j$**  if the two transaction conflict, and  $T_i$  accessed the data item on which the conflict arose earlier
  - We may label **the arc by the item that was accessed**

## Example 1

$T_1$	$T_2$
read( $A$ ) $A := A - 50$	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ ) read( $B$ )
write( $A$ ) read( $B$ ) $B := B + 50$ write( $B$ )	$B := B + temp$ write( $B$ )

Fig 14.11



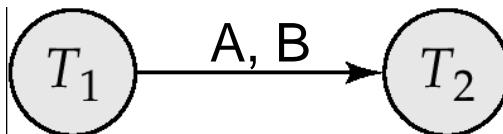
- A schedule is **conflict serializable** if and only if **its precedence graph is acyclic**



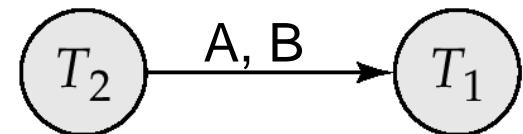
# Precedence Graph

for (a) Serial Schedule 1 and (b) Serial Schedule 2

Fig 14.10



(a)



(b)

$T_1$	$T_2$
<code>read(A) <math>A := A - 50</math> write (A) read(B) <math>B := B + 50</math> write(B)</code>	<code>read(A) <math>temp := A * 0.1</math> <math>A := A - temp</math> write(A) read(B) <math>B := B + temp</math> write(B)</code>

$T_1$	$T_2$
<code>read(A) <math>A := A - 50</math> write(A) read(B) <math>B := B + 50</math> write(B)</code>	<code>read(A) <math>temp := A * 0.1</math> <math>A := A - temp</math> write(A) read(B) <math>B := B + temp</math> write(B)</code>



# Test for Conflict Serializability [1/2]

- A schedule is **conflict serializable** if and only if its precedence graph is acyclic
- Cycle-detection algorithms exist which take order  $n^2$  time, where  $n$  is the number of vertices in the graph.
  - (Better algorithms take order  $n + e$  where  $e$  is the number of edges.)

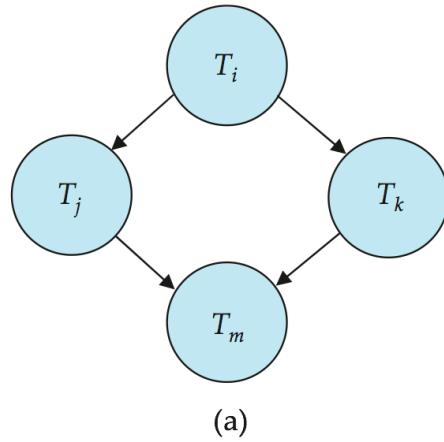


Fig 14.12

Serializability를 transaction들이 들어오는 run time에 check하고 concurrent execution을 허락하는 방법은 현실적이지 못하다!

Serializability를 보장하는 concurrency control scheme들을 이용하는것이 정답!

(참고) Cycle detection in DAG

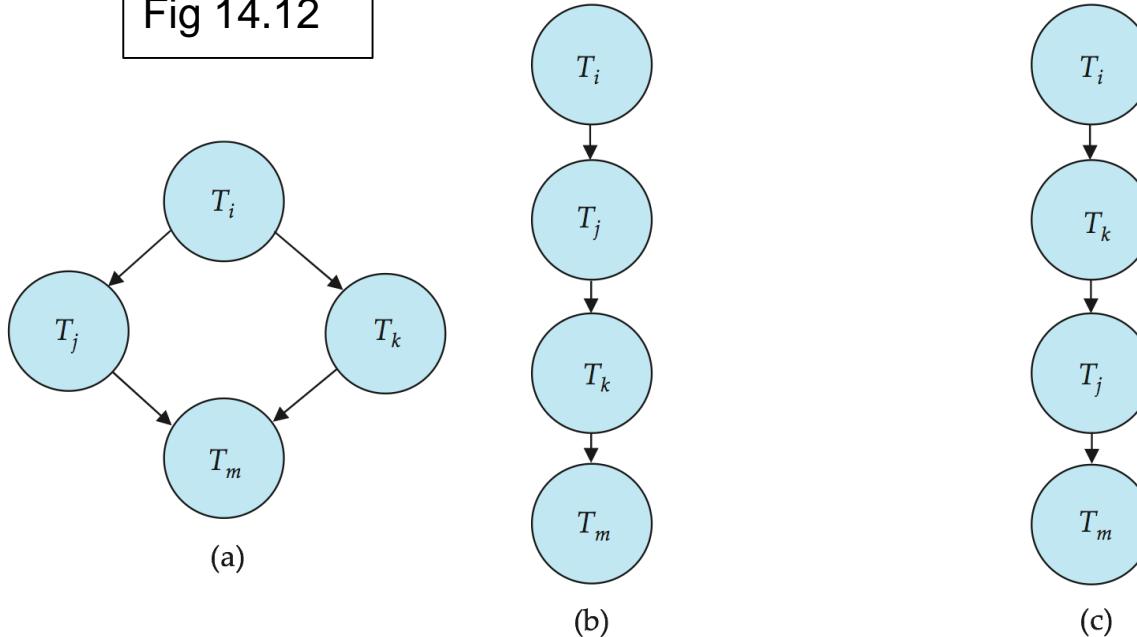
- Depth-First Search 기반 알고리즘:  $O(n^2)$
- Kahn's 알고리즘 :  $O(n+e)$
- Shortest path 방식:  $O(n+e)$  (단, 노드에 weight 필요)



# Test for Conflict Serializability [2/2]

- If precedence graph is acyclic, the serializability order can be obtained by a *topological sorting* of the graph
  - This is a linear order consistent with the partial order of the graph
  - For example, a serializability order for Schedule A would be  $T_5 \rightarrow T_1 \rightarrow T_3 \rightarrow T_2 \rightarrow T_4$ 
    - Are there others?

Fig 14.12



만약 여러 개의 Topological Sorting Order 가 존재한다면 그만큼의 serial 스케줄과 동등하다는 것을 나타냄



# Test for View Serializability

- The precedence graph test for conflict serializability cannot be used directly to test for view serializability
  - Extension to test for view serializability has **cost exponential in the size of the precedence graph**
- The problem of checking if a schedule is view serializable falls in the class of **NP-complete problems**.
  - Thus existence of an efficient algorithm is *extremely* unlikely
- However practical algorithms that just check some **sufficient conditions** for view serializability can still be used



# Chapter 14: Transactions

- Transaction Concept
- A Simple Transaction Model
- Storage Structure
- Transaction Atomicity and Durability
- Transaction Isolation
- Serializability
- **Transaction Isolation and Atomicity**
- Transaction Isolation Level
- Implementation of Isolation Levels
- Transaction Definition in SQL



# Recoverable Schedules

Need to address the effect of transaction failures on concurrently running transactions

- **Recoverable schedule** — if a transaction  $T_j$  reads a data item previously written by a transaction  $T_i$ , then the commit operation of  $T_i$  appears before the commit operation of  $T_j$
- The following schedule (Schedule 11) is not recoverable if  $T_9$  commits immediately after the read

Fig 14.14: Schedule 11

$T_8$	$T_9$
read (A)	
write (A)	read (A)
read (B)	commit



$T_8$	$T_9$
read (A) write (A)	
	read (A)
	read (B)
	commit

- If  $T_8$  should abort,  $T_9$  would have read (and possibly shown to the user) an inconsistent database state
  - Hence, DBMS must ensure that schedules are recoverable



# Cascading Rollbacks

- **Cascading rollback:** a single transaction failure leads to a series of transaction rollbacks
- Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

$T_{10}$	$T_{11}$	$T_{12}$
read (A) read (B) write (A)		
	read (A) write (A)	
abort		read (A)

Fig 14.15

불행의 쌍오:  $T_{10}$ 이 commit하기 전에 unlock을 했기에

- \*\* Suppose the lock of A is released without executing the commit instruction in  $T_{10}$ . If  $T_{10}$  fails,  $T_{10}$  is supposed to be rolled back, then  $T_{11}$  and  $T_{12}$  must also be rolled back
- **Cascading rollback** can lead to the **undoing** of a significant amount of work



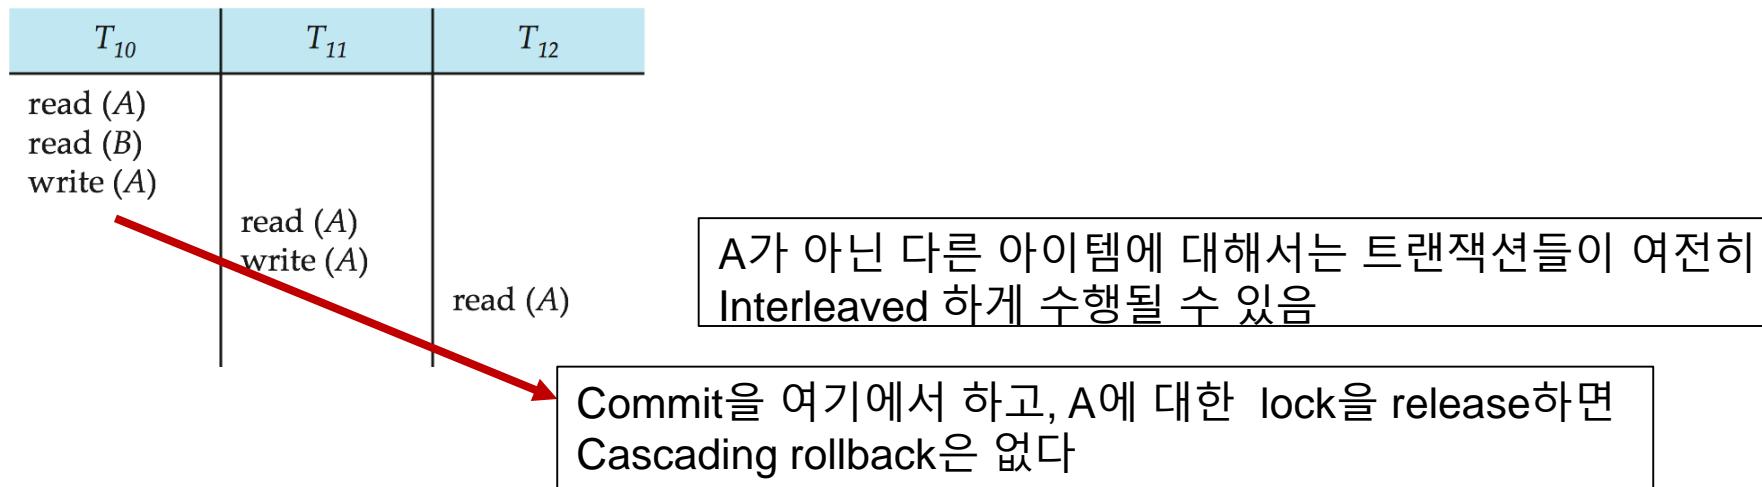
# Cascadeless Schedules

## ■ Cascadeless schedules: cascading rollbacks cannot occur

- for each pair of transactions  $T_i$  and  $T_j$  such that  $T_j$  reads a data item previously written by  $T_i$ , the commit operation of  $T_i$  appears before the read operation of  $T_j$
- Every cascadeless schedule is also recoverable

## ■ It is desirable to restrict the schedules to those that are cascadeless

## ■ Idea: hold locks until executing the commit instruction!



## ■ Trade Off

- More concurrency → More cascading rollback // early lock-release
- Less concurrency → Less cascading rollback // late lock-release



# Chapter 14: Transactions

- Transaction Concept
- A Simple Transaction Model
- Storage Structure
- Transaction Atomicity and Durability
- Transaction Isolation
- Serializability
- Transaction Isolation and Atomicity
- **Transaction Isolation Level**
- Implementation of Isolation Levels
- Transaction Definition in SQL



# Weak Levels of DB Consistency

- Serializability is rather strong
- Some applications are willing to live with weak levels of consistency, allowing schedules that are not serializable
  - Read-only transactions need not be serializable with respect to other transactions
    - ▶ A read-only transaction that wants to get an approximate total balance of all accounts
    - ▶ Database statistics computed for query optimization can be approximate
- Tradeoff accuracy for CC performance
  - Weak consistency → More serializable schedules → High performance
  - Strong consistency → Less serializable schedules → Low performance

Transaction Isolation Level == DB Consistency Level



# Transaction Isolation Levels in SQL-92

- **Serializable**: used to be default (conflict serializable) before 1992
- **Repeatable read**: only committed records to be read, repeated reads of same record must return same value
  - However, a transaction may not be serializable – it may find some records inserted by a transaction but not find others
- **Read committed**: only committed records can be read, but successive reads of record may return different (but committed) values
- **Read uncommitted**: even uncommitted records may be read
- Lower degrees of consistency are useful for gathering approximate information about the database
- Warning: some database systems do not ensure **serializable schedules** by default
  - MSSQL server는 default로 **read committed** 지원
  - Oracle and PostgreSQL by default support a level of consistency called **snapshot isolation** (not part of the SQL standard)
  - DBA may change (or reset) the transaction isolation level



# Chapter 14: Transactions

- Transaction Concept
- A Simple Transaction Model
- Storage Structure
- Transaction Atomicity and Durability
- Transaction Isolation
- Serializability
- Transaction Isolation and Atomicity
- Transaction Isolation Level
- **Implementation of Isolation Levels**
- Transaction Definition in SQL



# Implementation of Isolation Levels by Concurrency Control Schemes

- A database must provide a CC scheme ensuring all possible schedules are
  - either **conflict** or **view serializable**, and
  - are **recoverable** and **preferably cascadeless**
  - We study such protocols in **Chapter 16**
    - ▶ Locking based CC Schemes
    - ▶ Timestamps based CC Schemes
    - ▶ Multiple Version CC Schemes
    - ▶ Snapshot Isolation
- **Contradictory Goal of CC Schemes:** High Concurrency and High Isolation Level
- Concurrency-control schemes tradeoff between **the amount of concurrency** they allow and **the amount of overhead** that they incur
  - A policy in which only one transaction can execute at a time generates serial schedules in a simple way, but provides a poor degree of concurrency
  - Some schemes allow only **conflict-serializable schedules** to be generated, while others allow **view-serializable schedules** that are not conflict-serializable



# Concurrency Control vs. Serializability Tests

- Concurrency control protocols generally do not examine **the precedence graph** as it is being created
  - Testing a schedule for serializability *after* it has executed is a little too late!
  - Instead a protocol imposes a discipline that avoids nonserializable schedules
  - Tests for serializability help us understand **why a concurrency control protocol is correct**
  
- Serializable schedule만을 만들어내는 CC Scheme이 없이 그냥 concurrent하게 transaction들을 수행하고 precedence graph로 serializability를 판단한다면 어떨까?
  - Precedence Graph로 Serializability를 판단하는 것은 run-time에 어떤 트랜잭션들이 들어올지 모두 알고 있어야 하므로 비현실적
  - 만약 동시에 들어오는 트랜잭션들을 실시간으로 precedence 그래프를 그려서 serializability 판단한다고 해도 unserializable schedule들은 다 abort시켜야 하므로 비용이나 성능측면에서 비현실적임



# Chapter 14: Transactions

- Transaction Concept
- A Simple Transaction Model
- Storage Structure
- Transaction Atomicity and Durability
- Transaction Isolation
- Serializability
- Transaction Isolation and Atomicity
- Transaction Isolation Level
- Implementation of Isolation Levels
- **Transaction Definition in SQL**



# Transaction Definition in SQL

- Data manipulation language (DML) must include a construct for specifying the set of actions that comprise a transaction
- In SQL, a transaction begins implicitly and ends by:
  - **Commit work** commits current transaction and begins a new one
  - **Rollback work** causes current transaction to abort
- In almost all database systems, by default, every SQL statement also commits implicitly if it executes successfully
  - Implicit commit can be turned off by a database directive
  - In DBMS Schell, `setAutoCommit(false)`, 개발자가 원하는 시점에 `commit()` 가능
  - In JDBC, `connection.setAutoCommit(false); connection.commit()`

```
select name  
from instructor  
where salary < 90000
```

```
Transaction_begin  
S-Lock(name)  
Read(name)  
.....  
Commit-work
```

```
update instructor  
    set salary = salary * 1.03  
    where salary > 100000;
```

```
Transaction_begin  
S-lock (salary)  
Read(salary)  
Salary = salary * 1.03  
W-Lock(salary)  
Write(salary)  
Abort-work
```



# Transaction as SQL Statements [1/2]

- The current transaction model is based on `read(Q)` and `write(Q)` operations, not on SQL statements
- Concurrent execution of SQL statements may cause some strange situation
  - T1: `select count(ID) as num, name, sum(salary) as sum from instructor where salary > 90000;`  
 $\text{avg\_salary} = \text{sum}/\text{num};$
  - T2: `insert into instructor values ('11111', 'James', 'Marketing', 100000)`
  - T3: `update instructor set salary = salary * 0.9 where name = 'Wu'`
- Clearly there is a conflict in T1 and T2, but our current transaction model cannot capture the conflict
  - { T1; T2} 혹은 { T2; T1} 처럼 serial schedule의 결과를 가져야 하지만 our simple transaction model이 T2가 T1의 두단계 중간에 실행하는것을 허락하게 된다
  - 그러면 `avg_salary`의 값은 이상한 값을 return하게 된다
  - 새로 입력이 되는 tuple에 lock을 걸수 있는 방법이 없다!
  - Phantom Phenomenon: A conflict exists on phantom data



# Transaction as SQL Statements [2/2]

- T1과 T3도 interleaving하면 안되는데, simple CC scheme들은 interleaving 을 허락하여 strange 결과를 발생시킨다
- 결론: our simple transaction model은 read(Q), write(Q) 수준에서 serializability를 준수하는것을 기본으로 하지만, 실제 SQL 수준의 transaction에서는 고려해야 추가적인 issue들이 있고, serializability보다 weak한 기준으로 concurrency control을 해야 상황도 있다 (Ch15에 소개)

Semantic Gap between the read-write level transaction and the SQL level transaction!

- Phantom Phenomenon을 극복하는 방법 Index Locking Protocol: relation에 update 를 할때에 해당 record 들에 관련된 index bucket 에 locking
  - Predicate Locking Protocol: query의 predicate (where clause)와 conflict 이 나는 update 들에 locking
- Serializability보다 약한 기준으로 concurrency control 을 하는 방법
  - Snapshot Isolation
- Data record 들간에 concurrency control 을 하면서 Index 구조에도 동시에 concurrency control 를 해야 하는 상황을 위한 방법
  - Crabbing locking protocol