

Advanced DB

CHAPTER 15

TRANSACTION MNGMNT

Chapter 15: Transactions

- Transaction Concept
- Transaction State
- Implementation of Atomicity and Durability
- Concurrent Executions
- Serializability
- Recoverability
- Implementation of Isolation
- Testing for Serializability

Transactions

- What is a transaction?
 - Logical unit of work (user view)
 - A program unit that accesses and possibly updates various data items (system view)
 - DBMS must guarantee certain properties (ACID properties) for units of works declared as a DB transaction

Examples of Transactions

- Banks
 - Withdraw \$100 to account A.
 - Transfer \$50 from account A to B.
- Schools
 - Register course #409.433 for student #4321.
- Airlines
 - Check if two seats are available on flight #453.
 - Reserve the two seats on flight #453.
- Companies
 - Increase every employee's salary by 5%.

Dangers for Transactions

- Various types of failure
 - system crash
 - disk failure
 - system error
- Delayed disk write
 - disk access is performed in chunks: page (block)
 - i.e., write operation performed after the right amount of data has been gathered
 - buffer manager may pin a page

Properties of a Transaction

ACID properties

- **Atomicity**

“all or nothing”

- **Consistency (Correctness)**

Move from a consistent state to another consistent state

- **Isolation**

Should not be interfered by other transactions (concurrency)

- **Durability**

The effect of a completed transaction should be durable & public

Atomicity

- All or nothing

“Transfer \$50 from account A to account B”

Begin transaction

read(A,a)

a = a-50


write(A,a)

read(B,b)

b = b+50

write(B,b)

End Transaction



Consistency

- Move from a consistent state to another consistent state

“Withdraw \$100 from account A”

Begin transaction

read(A,a)

a = a-100

write(A,a)

End Transaction

What if A only had \$20?

- Responsibility of programmer

Isolation

- Should not be interfered by other transactions (concurrency)

“Transaction T1”

Begin transaction

read(A,a1)

a1 = a1-50

write(A,a1)

read(B,b1)

b1 = b1+50

write(B,b1)

End Transaction

“Transaction T2”


Begin transaction

read(A,a2)

a2 = a2-100

write(A,a2)

End Transaction



- Serial Execution VS Concurrent Execution

Durability

- The effect of a completed transaction should be durable & public

“Withdraw \$100 from account A”

Begin transaction

read (A,a)


a = a-100

write(A,a)

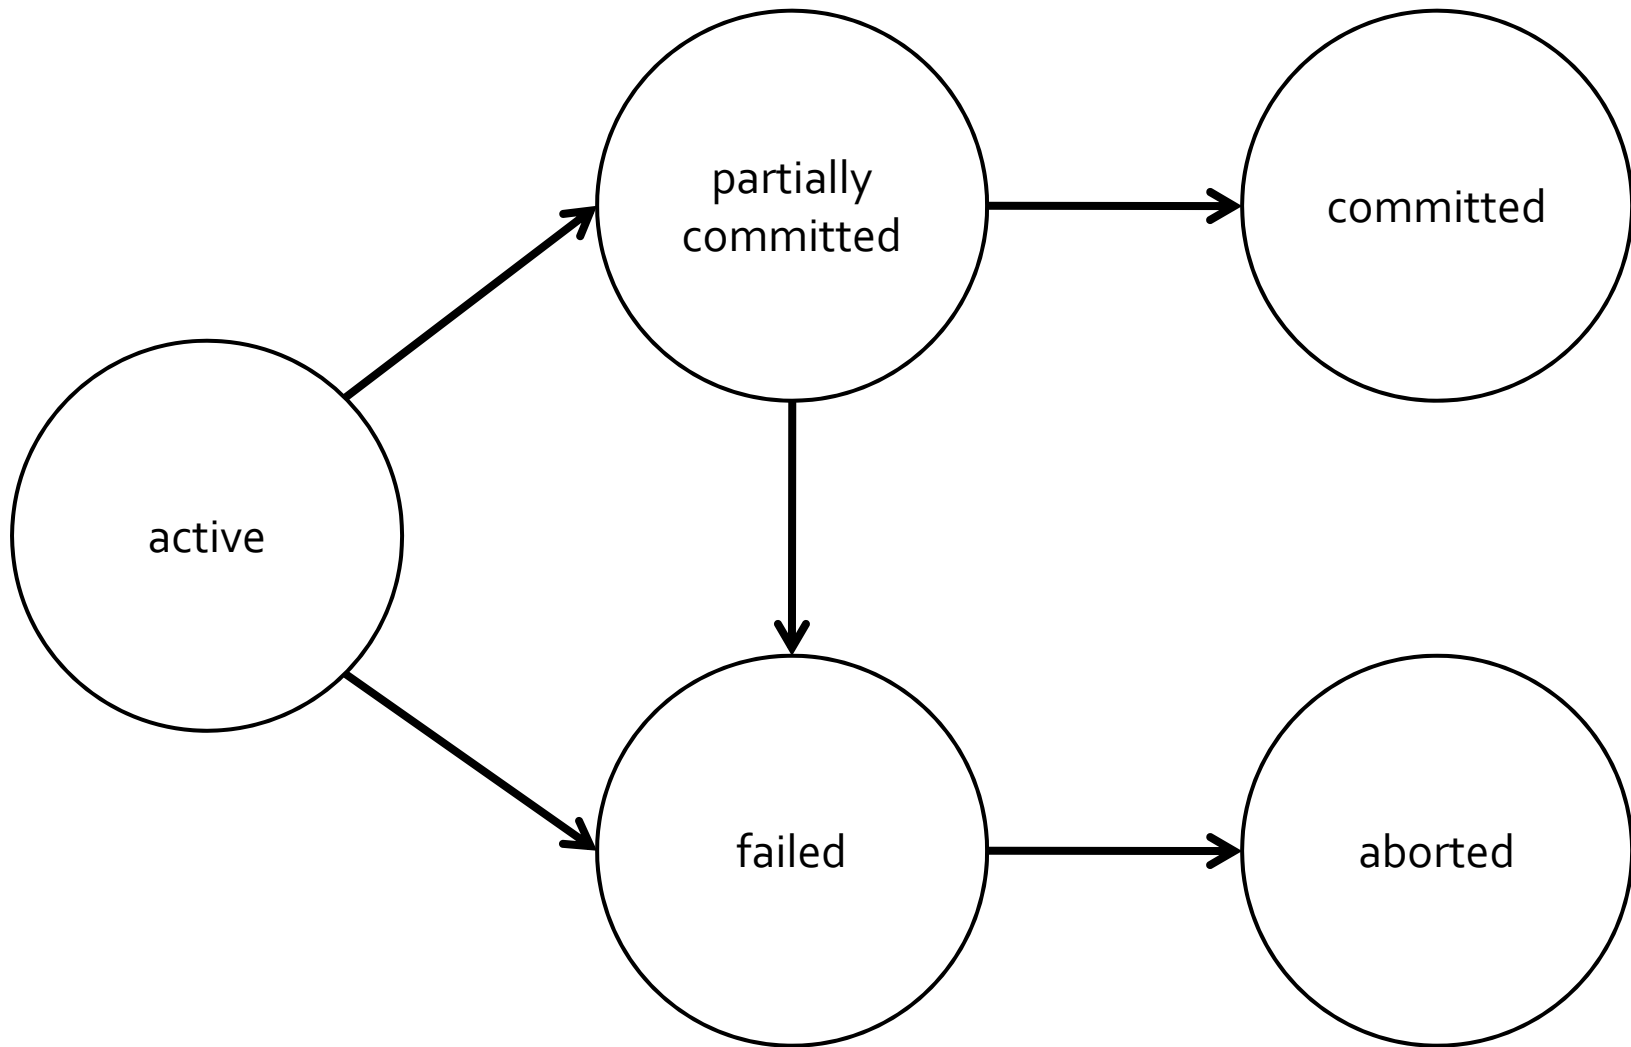
End Transaction

...

buffer flush



Transaction States



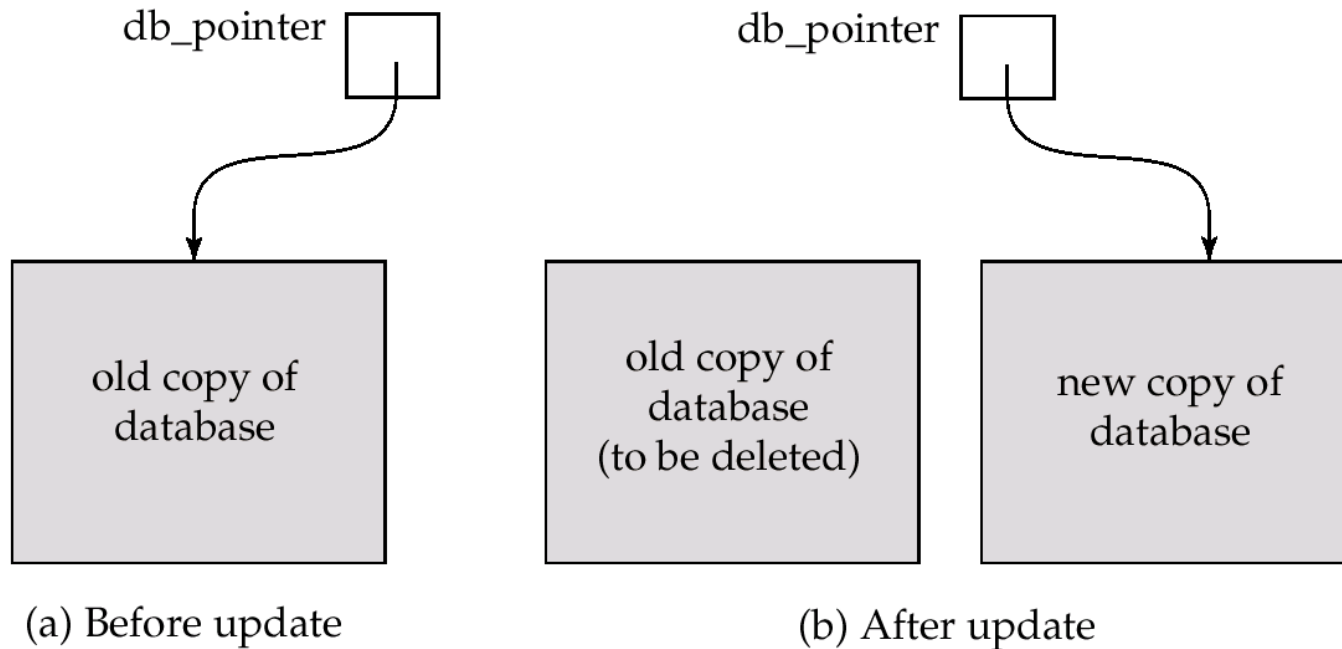
Transaction States (Cont.)

- *Active*: the initial state, transaction stays in this state while it is executing
- *Partially committed*: the final statement has been executed
- *Failed*: normal execution can no longer proceed
- *Aborted*: transaction has been rolled back and the database restored to its state prior to the start of the transaction.
 - Two options after it has been aborted:
 - restart the transaction – only if no internal logical error
 - kill the transaction
- *Committed*: after *successful completion*.

Implementation of Atomicity and Durability

- The *recovery-management* component of a database system implements the support for atomicity and durability.
- The *shadow-database* scheme:
 - assume that only one transaction is active at a time.
 - a pointer called `db_pointer` always points to the current consistent copy of the database.
 - all updates are made on a *shadow copy* of the database, and **db_pointer** is made to point to the updated shadow copy only after the transaction reaches partial commit and all updated pages have been flushed to disk.
 - in case transaction fails, old consistent copy pointed to by **db_pointer** can be used, and the shadow copy can be deleted.

Shadow Database Scheme



- Assumes disks to not fail
- Useful for text editors
- But extremely inefficient for large databases:
 - executing a single transaction requires copying the *entire* database.
 - Will see better schemes in Chapter 17.

Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system
 - increased processor and disk utilization
 - reduced average response time
- *Concurrency control schemes* – mechanisms to achieve isolation
 - to control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database
 - (Chapter 16)

Schedules

■ *Schedules*

- sequences that indicate the chronological order in which instructions of concurrent transactions are executed
- a schedule for a set of transactions must consist of all instructions of those transactions
- must preserve the order in which the instructions appear in each individual transaction.

Example Schedules

- Let T_1 transfer \$50 from A to B , and T_2 transfer 10% of the balance from A to B .
- The following is a serial schedule in which T_1 is followed by T_2 .

T_1	T_2
read(A) $A := A - 50$ write (A) read(B) $B := B + 50$ write(B)	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B)

Schedule 1 (fig 15.3)

Example Schedule (Cont.)

- Let T_1 and T_2 be the transactions defined previously.
- Schedule 3 is not a serial schedule, but it is **equivalent** to Sch.1.
- In both Schedule 1 and 3, the sum $A+B$ is preserved.*

T_1	T_2
$\text{read}(A)$ $A := A - 50$ $\text{write}(A)$	$\text{read}(A)$ $\text{temp} := A * 0.1$ $A := A - \text{temp}$ $\text{write}(A)$
$\text{read}(B)$ $B := B + 50$ $\text{write}(B)$	$\text{read}(B)$ $B := B + \text{temp}$ $\text{write}(B)$

Schedule 3 (fig 15.5)

Example Schedules (Cont.)

- Schedule 4 does not preserve the value of the the sum $A + B$.

T_1	T_2
$\text{read}(A)$ $A := A - 50$	$\text{read}(A)$ $\text{temp} := A * 0.1$ $A := A - \text{temp}$ $\text{write}(A)$ $\text{read}(B)$
$\text{write}(A)$ $\text{read}(B)$ $B := B + 50$ $\text{write}(B)$	$B := B + \text{temp}$ $\text{write}(B)$

Schedule 4 (fig 15.6)

Serializability

- Basic Assumption – Each transaction preserves database consistency.
 - Thus serial execution of a set of transactions preserves database consistency.
- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule.
 1. conflict serializability
 2. view serializability (not covered in this semester)
- We ignore operations other than *read* and *write* instructions.

Conflict Serializability

- Instructions I_i and I_j of transactions T_i and T_j respectively, conflict if and only if
 - there exists some item Q accessed by both I_i and I_j , and
 - at least one of these instructions is a write(Q)
 1. $I_i = \text{read}(Q)$, $I_j = \text{read}(Q)$. not conflict.
 2. $I_i = \text{read}(Q)$, $I_j = \text{write}(Q)$. conflict.
 3. $I_i = \text{write}(Q)$, $I_j = \text{read}(Q)$. conflict
 4. $I_i = \text{write}(Q)$, $I_j = \text{write}(Q)$. conflict
- Intuitively, a conflict between I_i and I_j forces a (logical) temporal order between them.

Conflict Serializability (Cont.)

- Schedules S and S' are *conflict equivalent* if
 - S can be transformed into a schedule S'
 - by a series of swaps of non-conflicting instructions
- A schedule S is *conflict serializable* if it is conflict equivalent to a serial schedule
- Example of a schedule that is not conflict serializable:

T_3	T_4
read(Q)	
	write(Q)
write(Q)	

Conflict Serializability (Cont.)

- A conflict serializable schedule

T_1	T_2
read(A) write(A)	
	read(A) write(A)
read(B) write(B)	
	read(B) write(B)

Recoverability

- Need to address the effect of transaction failures on concurrently running transactions
- *Recoverable schedule*
 - if a transaction T_j reads a data item previously written by a transaction T_i ,
 - the commit of T_i appears before the commit of T_j .
- The following schedule (Schedule 11) is not recoverable if T_9 commits immediately after the read

T_8	T_9
read(A)	
write(A)	
	read(A)
read(B)	

- DBMS must ensure that schedules are recoverable.

Cascading Rollback

- A single transaction failure can lead to a series of transaction rollbacks.
- Example: If T_{10} fails, T_{11} and T_{12} must also be rolled back.

T_{10}	T_{11}	T_{12}
read(A) read(B) write(A)	read(A) write(A)	read(A)

Can lead to the undoing of a significant amount of work


Cascadeless Schedules

- Cascading rollbacks cannot occur if
 - for each pair of transactions T_i and T_j
 - such that T_j reads a data item previously written by T_i ,
 - the commit of T_i appears before the read of T_j
- Every cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless

Testing for Serializability


- Consider some schedule of a set of transactions T_1, T_2, \dots, T_n
- **Precedence graph** — a direct graph where the vertices are the transactions (names).
- Draw an arc from T_i to T_j if
 - the two transaction conflict, and
 - T_i accessed the data item on which the conflict arose
- A schedule is conflict serializable if and only if its precedence graph is acyclic

example



Example

T_1	T_2	T_3	T_4	T_5
read(Y) read(Z)	read(X)			read(V) read(W) read(W)
	read(Y) write(Y)	write(Z)		
read(U)			read(Y) write(Y) read(Z) write(Z)	
read(U) write(U)				



*Precedence graph
for the schedule*

Serializability Tests vs. Concurrency Control

- Testing a schedule for serializability *after* it has executed is a little too late!
- Goal
 - to develop concurrency control protocols that will assure serializability
 - protocol will impose a discipline that avoids nonserializable schedules (Chapter 16)
- Tests for serializability help understand why a concurrency control protocol is correct.

END OF CHAPTER 15