



Chapter 14: Transactions

Database System Concepts, 6th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Transaction Concept

- A **transaction** is a *unit* of program execution that accesses and possibly updates various data items
- E.g. transaction to transfer \$50 from account A to account B:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)
- Two main issues to deal with:
 - **Failures of various kinds**, such as hardware failures and system crashes
 - **Concurrent execution** of multiple transactions



ACID Properties

To preserve the integrity of data the database system must ensure:

- **Atomicity.** Either **all** operations of the transaction are properly reflected in the database or **none** are.
- **Consistency.** Execution of a transaction in isolation preserves the consistency of the database.
- **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions.
- **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.



Atomicity

- “All or nothing”
- The system should ensure that updates of a partially executed transaction are not reflected in the database
- Transaction to transfer \$50 from account A to account B:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B) ← System crash
 5. $B := B + 50$
 6. **write**(B)
 - Money will be “lost” leading to an inconsistent database state



Consistency

- When the transaction completes successfully, the database must be consistent
 - During transaction execution, the database may be temporarily inconsistent
 - A transaction must see a consistent database

- Transaction to transfer \$50 from account A to account B:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)
 - The sum of A and B is unchanged by the execution of the transaction

- In general, consistency requirements include
 - Explicitly specified integrity constraints such as primary keys and foreign keys
 - Implicit integrity constraints
 - ▶ e.g. sum of balances of all accounts, minus sum of loan amounts



Isolation

- Each transaction must be unaware of other concurrently executing transactions
 - Intermediate transaction results must be hidden from other concurrently executed transactions
 - Isolation can be ensured trivially by running transactions **serially**

▶ That is, one after the other.

- Transaction to transfer \$50 from account A to account B:

T1

1. **read**(A)
2. $A := A - 50$
3. **write**(A)
4. **read**(B)
5. $B := B + 50$
6. **write**(B)

T2

read(A), read(B), print(A+B)

If T2 accesses the partially updated database, it will see an inconsistent database (the sum $A + B$ will be less than it should be).

- However, executing multiple transactions concurrently has significant benefits, as we will see later.



Durability

- The updates to the database by the transaction must persist even if there are software or hardware failures
- Transaction to transfer \$50 from account A to account B:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)
 - Once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), it must persist




Transaction State

- **Active** – the initial state; the transaction stays in this state while it is executing
- **Partially committed** – after the final statement has been executed
- **Failed** – after the discovery that normal execution can no longer proceed
- **Aborted** – after the transaction has been rolled back and the database restored to its state prior to the start of the transaction
Two options after it has been aborted:
 - restart the transaction
 - ▶ can be done only if no internal logical error
 - kill the transaction
- **Committed** – after successful completion



Transaction State (Cont.)





Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system.
Advantages are:
 - **Increased processor and disk utilization**, leading to better transaction *throughput*
 - **Reduced average response time** for transactions: short transactions need not wait behind long ones
- **Concurrency control schemes** – mechanisms to achieve isolation
 - To control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database



Schedules

- **Schedule** – a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
 - a schedule for a set of transactions must consist of all instructions of those transactions
 - must preserve the order in which the instructions appear in each individual transaction
- **Serial schedule** – instruction sequences from one by one transactions
- Simplified view of transactions
 - Our simplified schedules consist of only **read** and **write** instructions
 - We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes



Schedule 1

- Let T_1 transfer \$50 from A to B , and T_2 transfer 10% of the balance from A to B
- A serial schedule in which T_1 is followed by T_2 :

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit



Schedule 2

- A serial schedule in which T_2 is followed by T_1 :

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit



Schedule 3

- The following schedule is not a serial schedule, but it is *equivalent* to Schedule 1
 - We call it a *serializable schedule*

T_1	T_2
read (A) $A := A - 50$ write (A)	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	
	read (B) $B := B + temp$ write (B) commit

In Schedules 1, 2 and 3, the sum $A + B$ is preserved.



Schedule 4

- The following concurrent schedule does not preserve the value of $(A + B)$. The following schedule is **not serializable**.

T_1	T_2
read (A) $A := A - 50$	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
write (A) read (B) $B := B + 50$ write (B) commit	 $B := B + temp$ write (B) commit



Serializability

- **Basic Assumption** – Each transaction preserves database consistency
- Thus serial execution of a set of transactions preserves database consistency
- A (possibly concurrent) schedule is **serializable** if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notions of:
 1. **conflict serializability**
 2. **view serializability**



Conflicting Instructions

- Instructions I_i and I_j of transactions T_i and T_j respectively, **conflict** if and only if
 - There exists some item Q accessed by both I_i and I_j ,
 - and at least one of these instructions wrote Q .
 1. $I_i = \text{read}(Q)$, $I_j = \text{read}(Q)$. I_i and I_j don't conflict.
 2. $I_i = \text{read}(Q)$, $I_j = \text{write}(Q)$. They conflict.
 3. $I_i = \text{write}(Q)$, $I_j = \text{read}(Q)$. They conflict
 4. $I_i = \text{write}(Q)$, $I_j = \text{write}(Q)$. They conflict
- Intuitively, a conflict between I_i and I_j forces a (logical) temporal order between them
 - If I_i and I_j do not conflict, their results would remain the same even if they had been interchanged in the schedule



Conflict Serializability

- Schedules S and S' are **conflict equivalent** if S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions
- A schedule S is **conflict serializable** if It is conflict equivalent to a serial schedule
- Example of a schedule that is **not conflict serializable**:

T_3	T_4
read(Q)	write(Q)
write(Q)	

We are unable to swap instructions in the above schedule to obtain either the serial schedule $\langle T_3, T_4 \rangle$, or the serial schedule $\langle T_4, T_3 \rangle$.

T_3	T_4	T_3	T_4
read(Q)	write(Q)	read(Q)	write(Q)
write(Q)			
	write(Q)	write(Q)	



Conflict Serializability (Cont.)

- Schedule 3 can be transformed into Schedule 6, a serial schedule where T_2 follows T_1 , by series of **swaps of non-conflicting instructions**. Therefore Schedule 3 is conflict serializable.

Schedule3

T_1	T_2
read(A)	
write(A)	
	read(A)
	write(A)
read(B)	
write(B)	
	read(B)
	write(B)

Schedule 5 –
After Swapping a Pair of non
conflicting Instructions in
schedule 3

T_1	T_2
read(A)	
write(A)	
<u>read(B)</u>	read(A)
	<u>write(A)</u>
write(B)	
	read(B)
	write(B)

Schedule 6 –
A Serial Schedule That is
Equivalent to Schedule 3

T_1	T_2
read(A)	
write(A)	
read(B)	
write(B)	
	read(A)
	write(A)
	read(B)
	write(B)



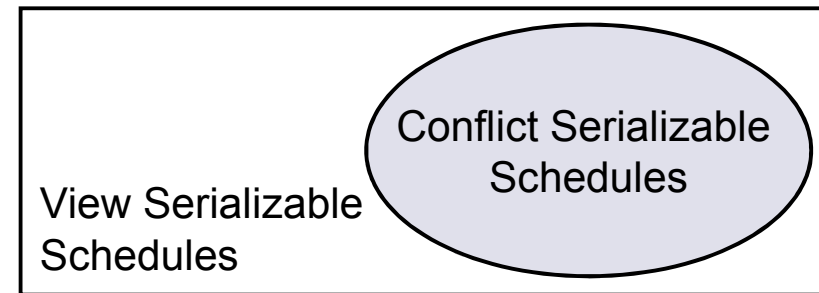
View Serializability

- Schedules S and S' are **view equivalent** if the following three conditions are met, for each data item Q :
 1. If in S , transaction T_i reads the initial value of Q , then in S' also transaction T_i must read the initial value of Q .
 2. If in S , T_i executes **read**(Q), and that value was produced by T_j (if any), then in S' also T_i must read the value of Q that was produced by the same **write**(Q) operation of T_j .
 3. The transaction (if any) that performs the final **write**(Q) operation in S must also perform the final **write**(Q) operation in S' .
- A schedule S is **view serializable** if it is view equivalent to a serial schedule



View Serializability (Cont.)

- Every conflict serializable schedule is also view serializable



- Below is a schedule which is **view-serializable** but **not conflict serializable**.

T_{27}	T_{28}	T_{29}
read (Q)	write (Q)	
write (Q)		write (Q)

A view equivalent serial schedule

T_{27}	T_{28}	T_{29}
read (Q)	write (Q)	
write (Q)		write (Q)


- Every view serializable schedule that is not conflict serializable has **blind writes**



Testing for Serializability

- Consider some schedule of a set of transactions T_1, T_2, \dots, T_n
- **Precedence graph** – a direct graph where the vertices are transactions(names)
 - draw an arc from T_i to T_j if the two transaction conflict, and T_i accessed the data item on which the conflict arose earlier
 - may label the arc by the item that was accessed
- Example




T_1	T_2
read(A) $A := A - 50$	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B)
write(A) read(B) $B := B + 50$ write(B)	 $B := B + temp$ write(B)






Test for Conflict Serializability

- A schedule is **conflict serializable** if and only if its **precedence graph** is **acyclic**
- Cycle-detection algorithms exist which take order n^2 time, where n is the number of vertices in the graph
 - (Better algorithms take order $n + e$ where e is the number of edges.)
- If precedence graph is acyclic, the serializability order can be obtained by a **topological sorting** of the graph
 - A linear order consistent with the partial order of the graph






Precedence Graph for Serial Schedules



(a) Schedule 1



(b) Schedule 2

T_1	T_2
$\text{read}(A)$ $A := A - 50$ $\text{write}(A)$ $\text{read}(B)$ $B := B + 50$ $\text{write}(B)$	$\text{read}(A)$ $\text{temp} := A * 0.1$ $A := A - \text{temp}$ $\text{write}(A)$ $\text{read}(B)$ $B := B + \text{temp}$ $\text{write}(B)$

T_1	T_2
$\text{read}(A)$ $A := A - 50$ $\text{write}(A)$ $\text{read}(B)$ $B := B + 50$ $\text{write}(B)$	$\text{read}(A)$ $\text{temp} := A * 0.1$ $A := A - \text{temp}$ $\text{write}(A)$ $\text{read}(B)$ $B := B + \text{temp}$ $\text{write}(B)$



Test for View Serializability

- The precedence graph test for conflict serializability cannot be used directly to test for view serializability
 - Extension to test for view serializability has cost exponential in the size of the precedence graph
- The problem of checking if a schedule is view serializable falls in the class of *NP-complete problems*
 - Thus existence of an efficient algorithm is *extremely* unlikely
- However, practical algorithms that just check some **sufficient conditions** for view serializability can still be used



Recoverable Schedules

- Need to address the effect of transaction failures on concurrently running transactions
- **Recoverable schedule**
 - If a transaction T_j reads a data item previously written by a transaction T_i ,
 - then the commit operation of T_i appears before the commit operation of T_j
- The following schedule is not recoverable if T_9 commits immediately after the read

T_8	T_9
read (A) write (A)	
	read (A) commit
read (B)	

- DBMS must ensure that schedules are recoverable



Cascading Rollbacks

- A single transaction failure leads to a series of transaction rollbacks
- Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

T_{10}	T_{11}	T_{12}
read (A) read (B) write (A) abort	read (A) write (A)	read (A)

- Can lead to **the undoing** of a significant amount of work



Cascadeless Schedules

- **Cascadeless schedules** — cascading rollbacks cannot occur;
 - For each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i ,
 - The commit operation of T_i appears before the read operation of T_j
- Every cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless

- Idea: block other transactions until executing the commit instruction
 - More concurrency → More cascading rollback
 - Less cascading rollback → Less concurrency



Concurrency Control

- A database must provide a mechanism that will ensure that all possible schedules are
 - either conflict or view serializable, and
 - are recoverable and preferably cascadeless
- A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency
- Testing a schedule for serializability *after* it has executed is a little too late
- **Goal** – to develop concurrency control protocols that will assure serializability
 - Concurrency-control schemes tradeoff between the amount of concurrency they allow and the amount of overhead that they incur



Weak Levels of Consistency

- Some applications are willing to live with weak levels of consistency, allowing schedules that are not serializable
 - Some transactions need not be serializable with respect to other transactions
 - ▶ E.g. a read-only transaction that wants to get an approximate total balance of all accounts
 - ▶ E.g. database statistics computed for query optimization can be approximate (why?)
- Tradeoff accuracy for performance



Levels of Consistency in SQL-92

- **Serializable** — default
- **Repeatable read** — only committed records to be read, repeated reads of same record must return same value
 - However, a transaction may not be serializable – it may find some records inserted by a transaction but not find others
- **Read committed** — only committed records can be read, but successive reads of record may return different (but committed) values
- **Read uncommitted** — even uncommitted records may be read

- Warning: some database systems do not ensure serializable schedules by default
 - E.g. Oracle and PostgreSQL by default support a level of consistency called snapshot isolation (not part of the SQL standard)



Transaction Definition in SQL

- Data manipulation language must include a construct for specifying the set of actions that comprise a transaction
- In SQL, **a transaction begins implicitly**
- A transaction in SQL **ends by:**
 - **Commit work** commits current transaction and begins a new one.
 - **Rollback work** causes current transaction to abort.
- In almost all database systems, by default, every SQL statement also commits implicitly if it executes successfully
 - Implicit commit can be turned off by a database directive
 - ▶ E.g. in JDBC, `connection.setAutoCommit(false);`



End of Chapter 14

Database System Concepts, 6th Ed.

©Silberschatz, Korth and Sudarshan
See www.db-book.com for conditions on re-use