



Chapter 15 : Concurrency Control

Database System Concepts, 6th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Concurrency Control

- A database must provide a mechanism that will ensure that all possible schedules are
 - either conflict or view serializable, and
 - are recoverable and preferably cascadeless
- A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency
- Testing a schedule for serializability *after* it has executed is a little too late
- **Goal** – to develop concurrency control protocols that will assure serializability
 - Concurrency-control schemes tradeoff between the amount of concurrency they allow and the amount of overhead that they incur



Lock-Based Protocols

- A lock is a mechanism to control concurrent access to a data item
- Data items can be locked in two modes:
 1. **exclusive** (X) *mode*. Data item can be both read as well as written. X-lock is requested using **lock-X** instruction.
 2. **shared** (S) *mode*. Data item can only be read. S-lock is requested using **lock-S** instruction.
- Lock requests are made to concurrency-control manager
- Transaction can proceed **only after request is granted**



Granting of Locks

■ Lock-compatibility matrix

	S	X
S	true	false
X	false	false

- A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions
- Any number of transactions can hold shared locks on an item
- If any transaction holds an exclusive lock on the item no other transaction may hold any lock on the item
- If a lock cannot be granted, the requesting transaction is made to wait till all incompatible locks held by other transactions have been released



Transaction with Locking

- Example of a transaction performing locking:

```
 $T_2$ : lock-S(A);  
      read (A);  
      unlock(A);  
      lock-S(B);  
      read (B);  
      unlock(B);  
      display(A+B)
```

- A **locking protocol** is a set of rules followed by **all transactions** while requesting and releasing locks
- Locking protocols restrict **the set of possible schedules**



Pitfalls of Lock-Based Protocols

- Consider the partial schedule

T_3	T_4
lock-x (B)	
read (B)	
$B := B - 50$	
write (B)	
	lock-s (A)
	read (A)
	lock-s (B)
lock-x (A)	

- Neither T_3 nor T_4 can make progress
 - Executing **lock-S(B)** causes T_4 to wait for T_3 to release its lock on B
 - Executing **lock-X(A)** causes T_3 to wait for T_4 to release its lock on A
- Such a situation is called a **deadlock**
 - To handle a deadlock one of T_3 or T_4 must be rolled back and its locks released



Pitfalls of Lock-Based Protocols (Cont.)

- The potential for deadlock exists in most locking protocols
 - Deadlocks are a necessary evil
- **Starvation** is also possible if concurrency control manager is badly designed
- Example
 - A transaction may be waiting for an X-lock on an item,
 - while a sequence of other transactions request and are granted an S-lock on the same item.
 - The same transaction is repeatedly rolled back due to deadlocks
- Concurrency control manager can be designed to prevent **starvation**



Two-Phase Locking Protocol (2PL)

- This is a protocol which ensures conflict-serializable schedules
- Phase 1: Growing Phase
 - transaction may obtain locks
 - transaction may not release locks
- Phase 2: Shrinking Phase
 - transaction may release locks
 - transaction may not obtain locks
- The protocol assures (conflict) serializability
 - The transactions can be serialized in the order of their **lock points** (i.e., the point where a transaction acquired its final lock)
 - There can be conflict serializable schedules that cannot be obtained if two-phase locking is used



Partial Schedule Under 2PL

T_5	T_6	T_7
<u>lock-X(A)</u> read(A) <u>lock-S(B)</u> read(B) write(A) <u>unlock(A)</u>	<u>lock-X(A)</u> read(A) write(A) <u>unlock(A)</u>	<u>lock-S(A)</u> read(A)



Strict / Rigorous 2PL

- Two-phase locking *does not* ensure freedom from deadlocks
- Cascading roll-back is possible under two-phase locking

- **Strict two-phase locking**
 - Here a transaction must hold **all its exclusive locks** till it commits/aborts
 - No cascading rollback
- **Rigorous two-phase locking**
 - Here ***all* locks(shared and exclusive)** are held till commit/abort
 - No cascading rollback (of course)
 - In this protocol transactions can be serialized in the order in which they commit



Lock Conversions

- The original lock mode with (lock-X, lock-S)
 - assign lock-X on a data D when D is both read and written
- Two-phase locking with lock conversions:
 - First Phase:
 - can acquire a lock-S on item
 - can acquire a lock-X on item
 - can convert a lock-S to a lock-X (upgrade)
 - Second Phase:
 - can release a lock-S
 - can release a lock-X
 - can convert a lock-X to a lock-S (downgrade)
- This protocol assures serializability
- The refined 2PL gets more concurrency than the original 2PL




Deadlock Detection


- Deadlocks can be described as a *wait-for graph*, which consists of a pair $G = (V, E)$,
 - V is a set of vertices (all the transactions in the system)
 - E is a set of edges; each element is an ordered pair $T_i \rightarrow T_j$.
- If $T_i \rightarrow T_j$ is in E , then there is a directed edge from T_i to T_j , implying that T_i is waiting for T_j to release a data item
- When T_i requests a data item currently being held by T_j , then the edge $T_i \rightarrow T_j$ is inserted in the wait-for graph
 - This edge is removed only when T_j is no longer holding a data item needed by T_i
- The system is in a deadlock state if and only if the wait-for graph has a cycle
 - Must invoke a deadlock-detection algorithm periodically to look for cycles



Deadlock Detection (Cont.)



Wait-for graph without a cycle



Wait-for graph with a cycle



Multiversion Two-Phase Locking

- Motivation: Decision support queries that read large amounts of data have concurrency conflicts with OLTP transactions that update a few rows
 - Poor performance results
- Multiversion schemes keep old versions of data item to increase concurrency
 - Differentiates between read-only transactions and update transactions
 - **Ts-counter** is a global time-stamp clock
 - ▶ This is incremented during commit processing
- **Update transactions** acquire read and write locks, and hold all locks up to the end of the transaction
 - Each successful **write** results in the creation of a new version of the data item written
 - Each version of a data item has a single timestamp whose value is obtained from **ts-counter**
- **Read-only transactions** are assigned a timestamp by reading the current value of **ts-counter** before they start execution



Multiversion Two-Phase Locking (Cont.)

- Creation of multiple versions increases storage overhead
 - Extra tuples
 - Extra space in each tuple for storing version information
- Versions can, however, be garbage collected
 - E.g., if Q has two versions Q5 and Q9, and the oldest active transaction has timestamp > 9 , then Q5 will never be required again
- Problem: works well, but how does system know a transaction is read only?



Snapshot Isolation

■ A transaction T2 executing with Snapshot Isolation

- takes snapshot of committed data at start
- always reads/modifies data in its own snapshot
- updates of concurrent transactions are not visible to T2
- writes of T2 complete when it commits
- **First-committer-wins rule:**
 - ▶ Commits only if no other concurrent transaction has already written data that T2 intends to write

T1	T2	T3
W(Y := 1) Commit		
	Start R(X) → 0 R(Y) → 1	
		W(X:=2) W(Z:=3) Commit
	R(Z) → 0 R(Y) → 1 W(X:=3) Commit-Req Abort	

Concurrent updates not visible

Own updates are visible

Not first-committer of X

Serialization error, T2 is rolled back



Benefits of Snapshot Isolation

- Reading is *never* blocked
 - and also doesn't block other txs activities
- Performance similar to Read Committed
- Avoids the usual anomalies
 - No dirty read
 - No lost update
 - No non-repeatable read
 - Predicate based selects are repeatable
- Problems with SI
 - SI does not always give serializable executions
 - ▶ Serializable: among two concurrent txs, one sees the effects of the other
 - ▶ SI: neither sees the effects of the other
 - Result: integrity constraints can be violated
- Variants implemented in many database systems
 - E.g., Oracle, PostgreSQL, SQL Server 2005



Weak Levels of Consistency in SQL

- SQL allows non-serializable executions
 - **Serializable**: is the default
 - **Repeatable read**: allows only committed records to be read, and repeating a read should return the same value (so read locks should be retained)
 - ▶ However, the phantom phenomenon need not be prevented
 - T1 may see some records inserted by T2, but may not see others inserted by T2
 - **Read committed**: only committed records can be read, but successive reads of record may return different (but committed) values
 - **Read uncommitted**: even uncommitted records may be read
- In many database systems, read committed is the default consistency level
 - has to be explicitly changed to serializable when required
 - ▶ **set isolation level serializable**



End of Chapter 15

Database System Concepts, 6th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use