



Chapter 11: Indexing and Hashing

Database System Concepts, 6th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Basic Concepts

- Indexing mechanisms used to speed up access to desired data
 - e.g., author catalog in library, term index at the end of a book
- **Search key** - attribute to set of attributes used to look up records in a file
- An **index file** consists of records (called **index entries**) of the form

search-key	pointer
------------	---------

- Index files are typically much smaller than the original file
- Two basic kinds of indices:
 - **Ordered indices:** search keys are stored in sorted order
 - **Hash indices:** search keys are distributed uniformly across “buckets” using a “hash function”



Index Evaluation Metrics

- Access types supported efficiently
 - **Point query:** records with a specified value in the attribute
 - **Range query:** records with an attribute value falling in a specified range of values

- Access time
- Insertion time
- Deletion time

- Space overhead



Ordered Indices

- **Ordered index:** index entries are stored sorted on the search key value
 - e.g., author catalog in library
- **Primary index:** an index whose search key specifies the sequential order of the file
 - Also called **clustering index**
 - The search key of a primary index is usually (but not necessarily) the primary key
- **Secondary index:** an index whose search key specifies an order different from the sequential order of the file
 - Also called **non-clustering index**
- **Index-sequential file:** ordered sequential file with a primary index



Dense Index Files

■ Dense index

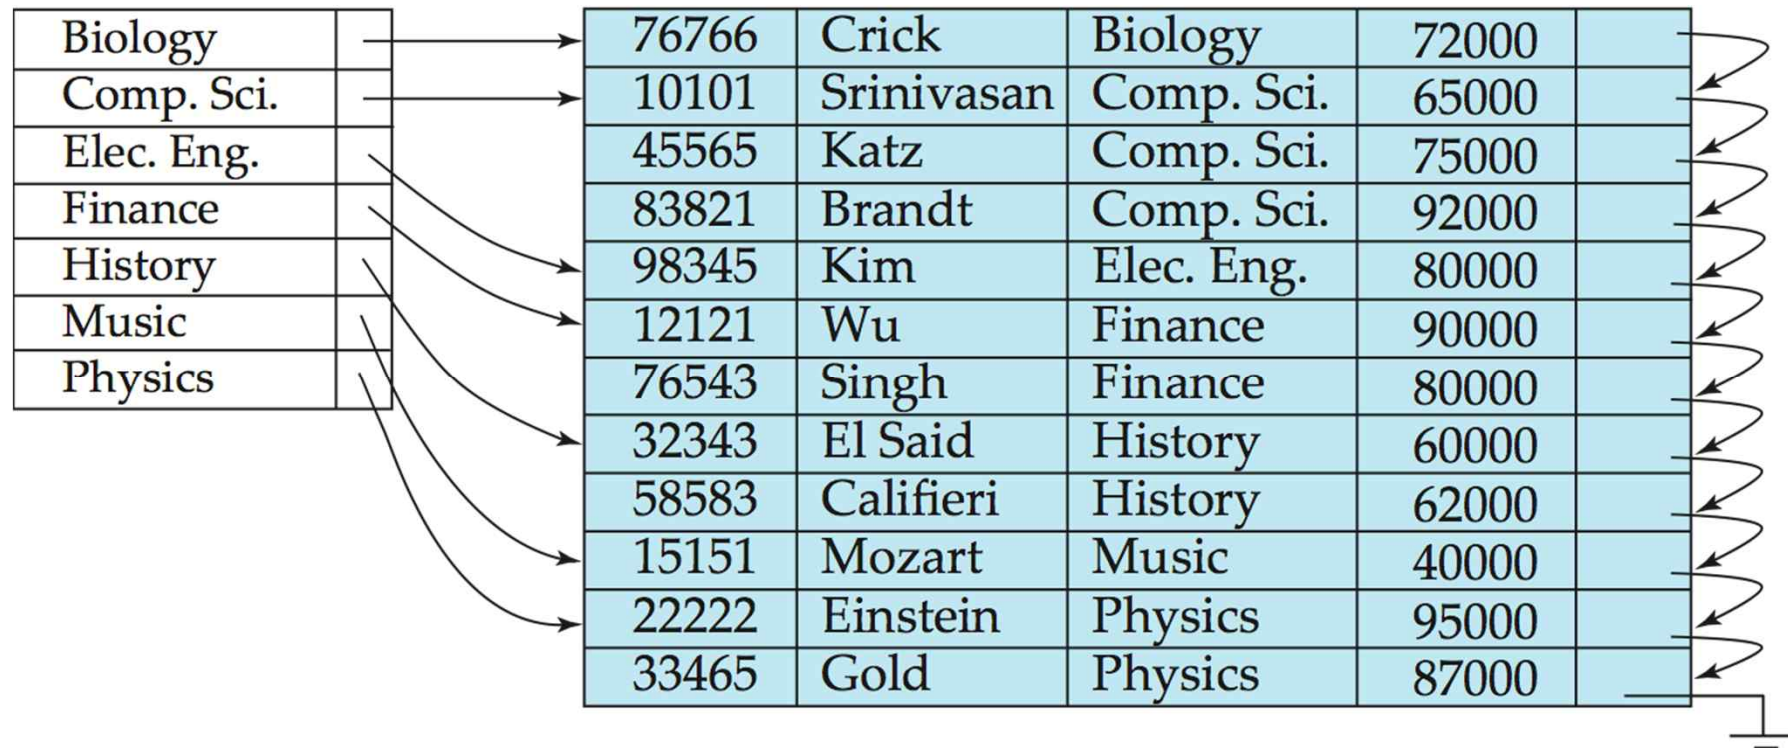
- Index record appears for every search-key value in the file

10101		→	10101	Srinivasan	Comp. Sci.	65000	
12121		→	12121	Wu	Finance	90000	
15151		→	15151	Mozart	Music	40000	
22222		→	22222	Einstein	Physics	95000	
32343		→	32343	El Said	History	60000	
33456		→	33456	Gold	Physics	87000	
45565		→	45565	Katz	Comp. Sci.	75000	
58583		→	58583	Califieri	History	62000	
76543		→	76543	Singh	Finance	80000	
76766		→	76766	Crick	Biology	72000	
83821		→	83821	Brandt	Comp. Sci.	92000	
98345		→	98345	Kim	Elec. Eng.	80000	



Primary Indices Example

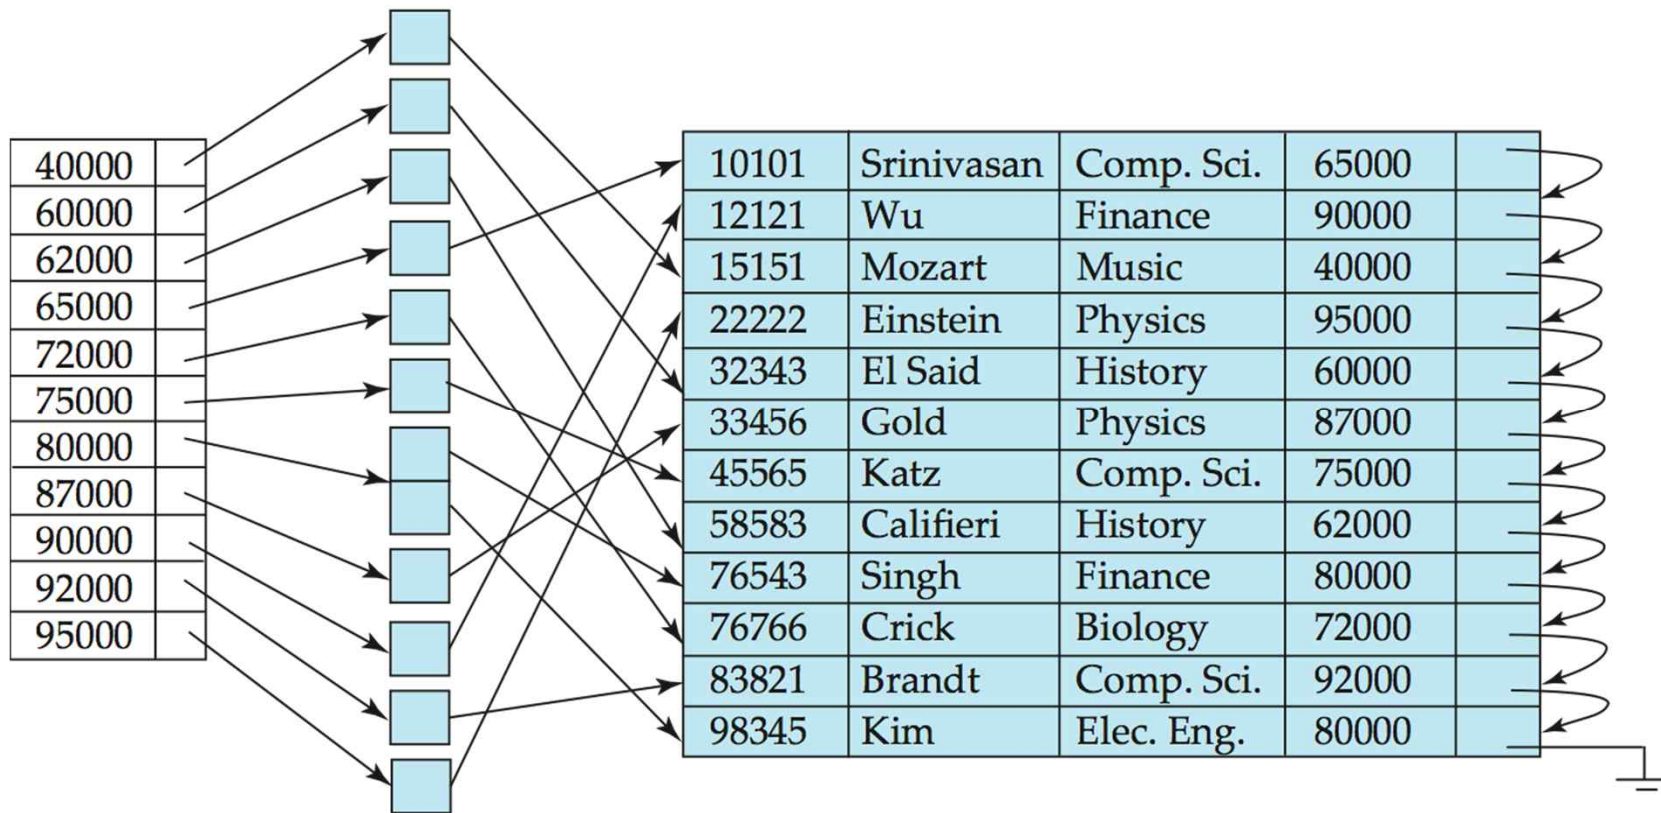
- Example: dense index on *dept_name*, with *instructor* file sorted on *dept_name*





Secondary Indices Example

- Index record points to a bucket that contains pointers to all the actual records with that particular search-key value.
- Secondary indices have to be dense

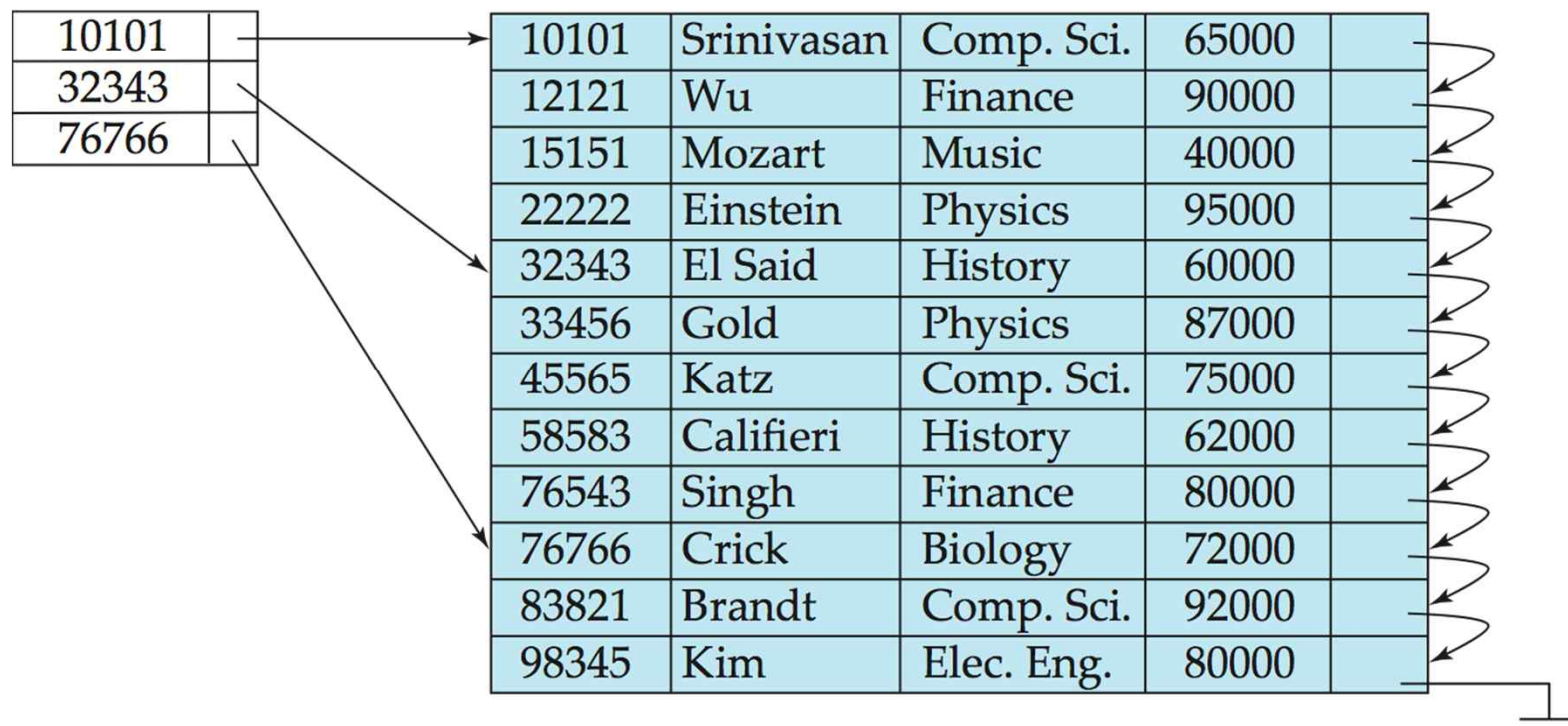


Secondary index on *salary* field of *instructor*



Sparse Index Files

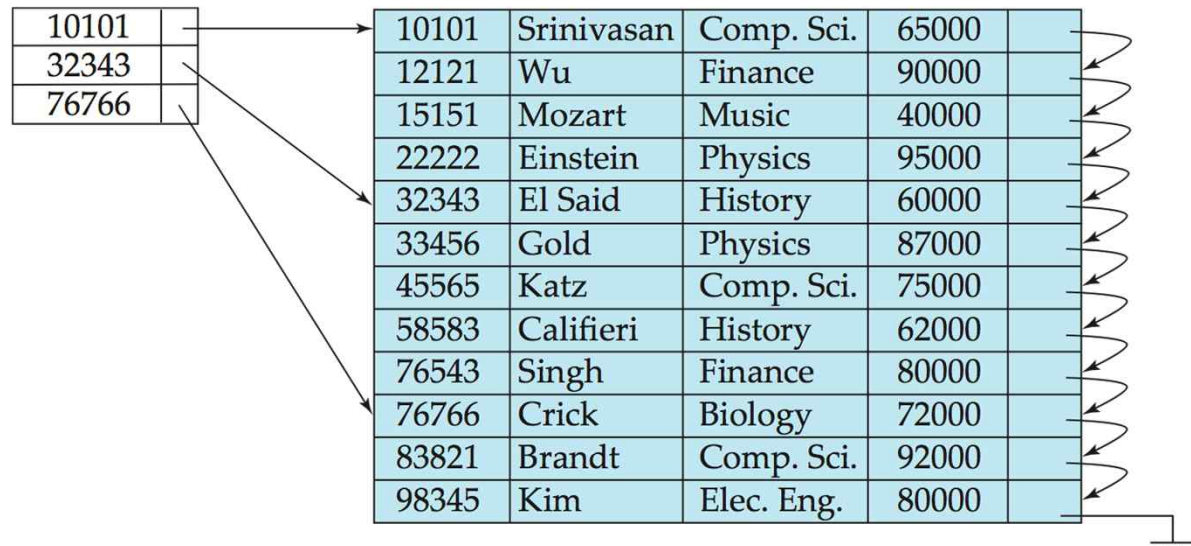
- **Sparse Index:** contains index records for **only some** search-key values.
 - Applicable when records are sequentially ordered on search-key
- Compared to dense indices:
 - Less space and less maintenance overhead for insertions and deletions
 - Generally slower than dense index for locating records





Index Update: Deletion

- If deleted record was the only record in the file with its particular search-key value, the search-key is deleted from the index also.
- **Dense indices** – deletion of search-key is similar to file record deletion
- **Sparse indices**
 - If an entry for the search key exists in the index,
 - ▶ Replace the entry in the index with the next search-key value in the file (in search-key order)
 - If the next search-key value already has an index entry
 - ▶ Delete the entry





Index Update: Insertion

- Perform a lookup using the search-key value appearing in the record to be inserted
- **Dense indices**
 - If the search-key value does not appear in the index, insert it.
- **Sparse indices**
 - If index stores an entry for each block of the file
 - ▶ No change needs to be made to the index, unless a new block is created
 - ▶ If a new block is created, the first search-key value appearing in the new block is inserted into the index




Primary and Secondary Indices

- Indices offer substantial benefits when searching for records
- When a file is modified, every index on the file must be updated
 - Updating indices imposes overhead on database modification
- Sequential scan using primary index is efficient, but a sequential scan using a secondary index is expensive
 - Each record access may fetch a new block from disk
 - Block fetch requires about 5 to 10 milliseconds, versus about 100 nanoseconds for memory access



Multilevel Index

- If primary index does not fit in memory, access becomes expensive
- Solution: treat primary index **kept on disk** as a sequential file and construct a sparse index on it
 - outer index – a sparse index of primary index
 - inner index – the primary index file
- If even outer index is too large to fit in main memory, yet another level of index can be created, and so on
- Indices at all levels must be updated on insertion or deletion from the file



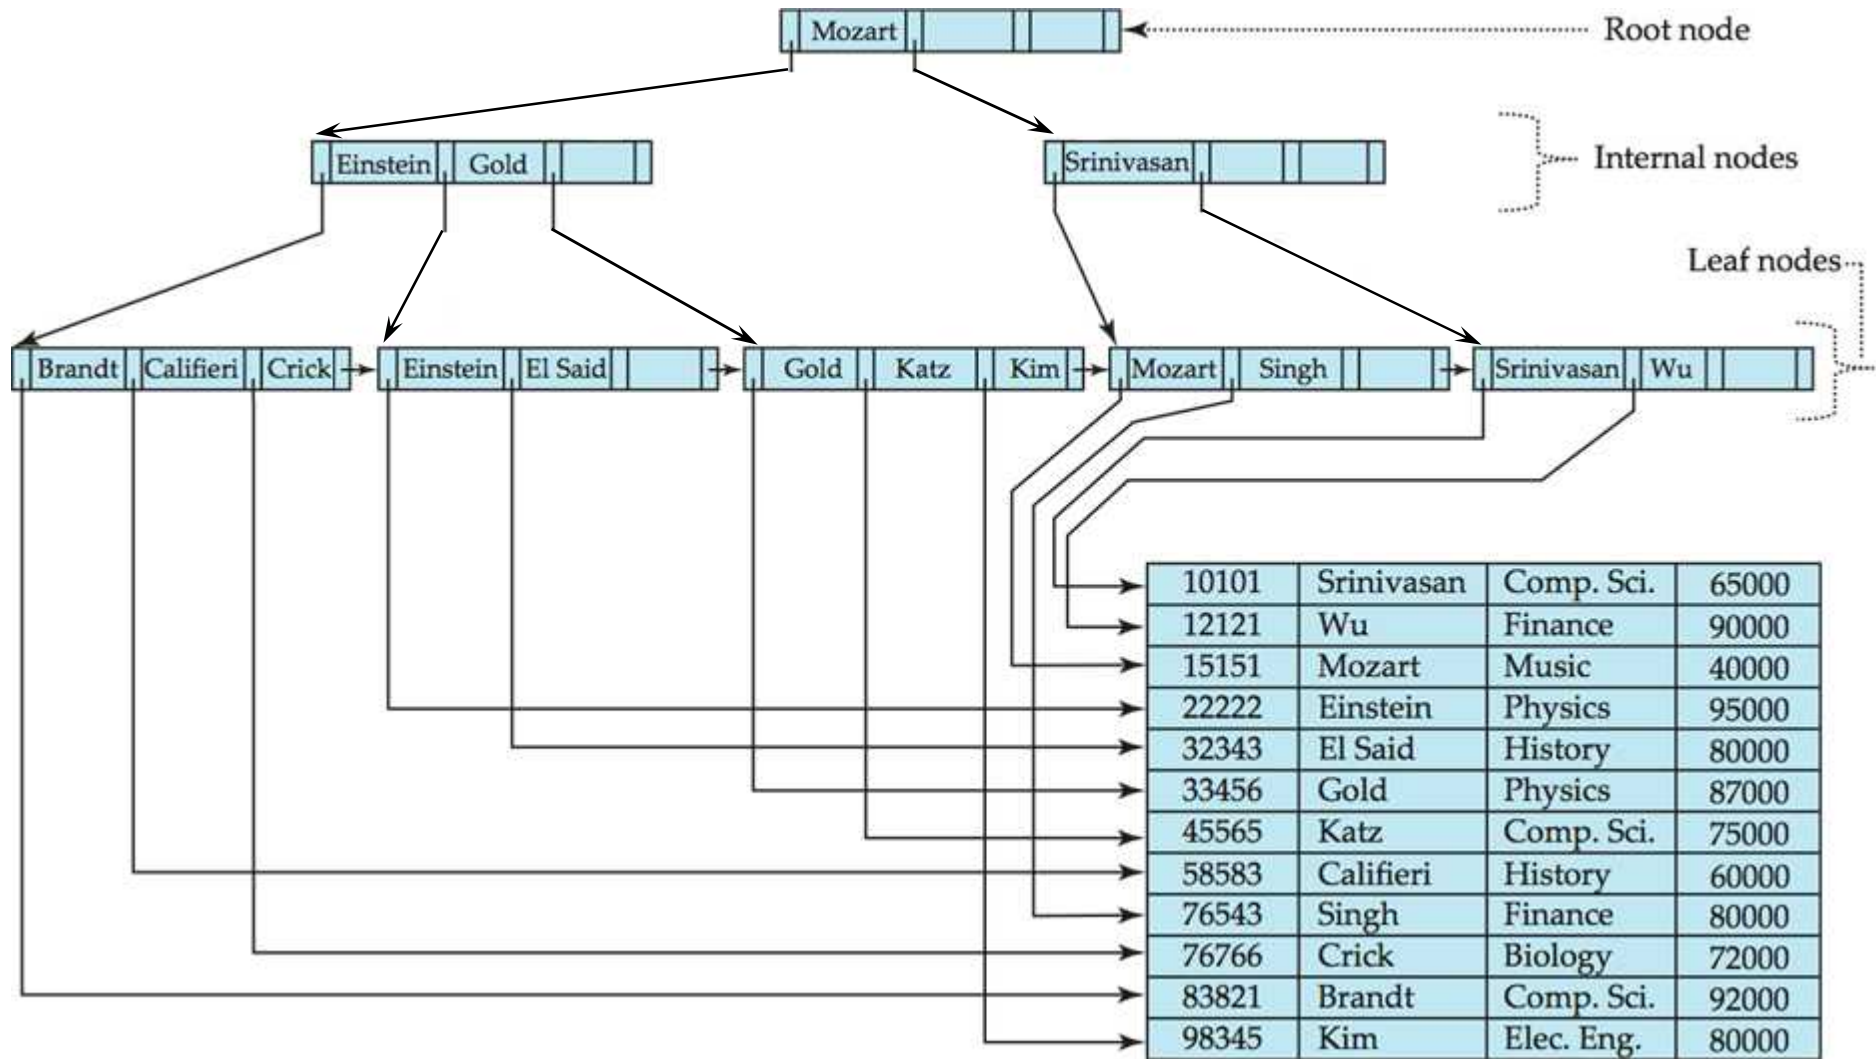


B⁺-Tree Index Files

- Advantage of B⁺-tree index files:
 - Automatically reorganizes itself with **small local changes**, in the face of insertions and deletions
 - Reorganization of entire file is not required to maintain performance
- (Minor) disadvantage of B⁺-trees:
 - Extra insertion and deletion overhead, space overhead
- B⁺-trees are used extensively
 - Advantages of B⁺-trees outweigh disadvantages



Example of B⁺-Tree





B⁺-Tree Index Files (Cont.)


A B⁺-tree is a rooted tree satisfying the following properties:

- All paths from root to leaf are of the same length
- Each node that is not a root or a leaf has between $\lceil n/2 \rceil$ and n children
- A leaf node has between $\lceil (n-1)/2 \rceil$ and $n-1$ values
- Special cases: root node
 - If the root is not a leaf, it has at least 2 children
 - If the root is a leaf (that is, there are no other nodes in the tree), it can have between 0 and $(n-1)$ values



B⁺-Tree Node Structure

- Typical node (non-leaf node)

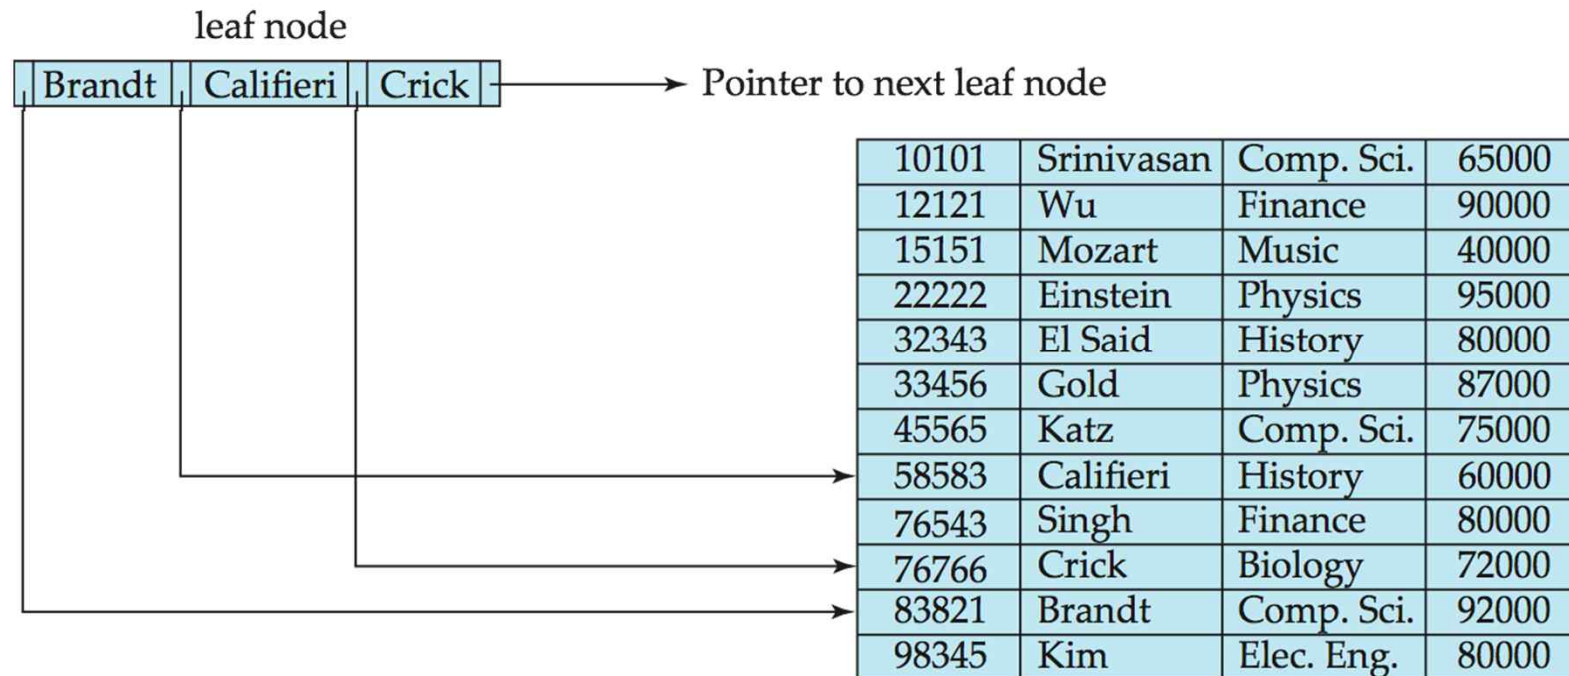


- K_i are the search-key values
 - P_i are pointers to children (for non-leaf nodes)
or pointers to records (for leaf nodes)
-
- The search-keys in a node are ordered
$$K_1 < K_2 < K_3 < \dots < K_{n-1}$$
 - We assume no duplicate keys



Leaf Nodes in B⁺-Trees

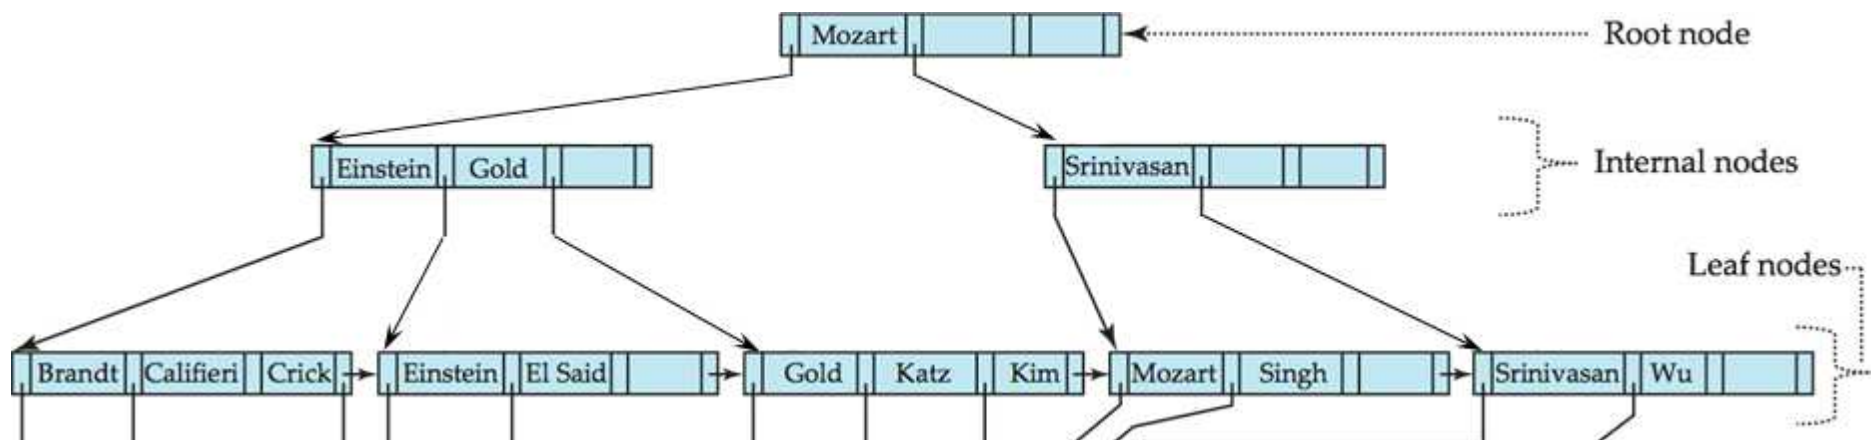
- For $i = 1, 2, \dots, n-1$, pointer P_i points to a file record with search-key value K_i
- P_n points to next leaf node in search-key order
- If L_i, L_j are leaf nodes and $i < j$, L_i 's search-key values are less than or equal to L_j 's search-key values






Non-Leaf Nodes in B⁺-Trees

- Non leaf nodes form a multi-level sparse index on the leaf nodes
- For a non-leaf node with m pointers ($m \leq n$):
 - All the search-keys in the subtree to which P_1 points are less than K_1
 - For $2 \leq i \leq m - 1$, all the search-keys in the subtree to which P_i points have values greater than or equal to K_{i-1} and less than K_i
 - All the search-keys in the subtree to which P_m points have values greater than or equal to K_{m-1}





Example of B⁺-tree



B⁺-tree for *instructor* file ($n = 6$)

- Leaf nodes must have between 3 and 5 values
 - $\lceil (n-1)/2 \rceil$ and $n-1$, with $n = 6$
- Non-leaf nodes other than root must have between 3 and 6 children
 - $\lceil (n/2) \rceil$ and n , with $n = 6$
- Root must have at least 2 children



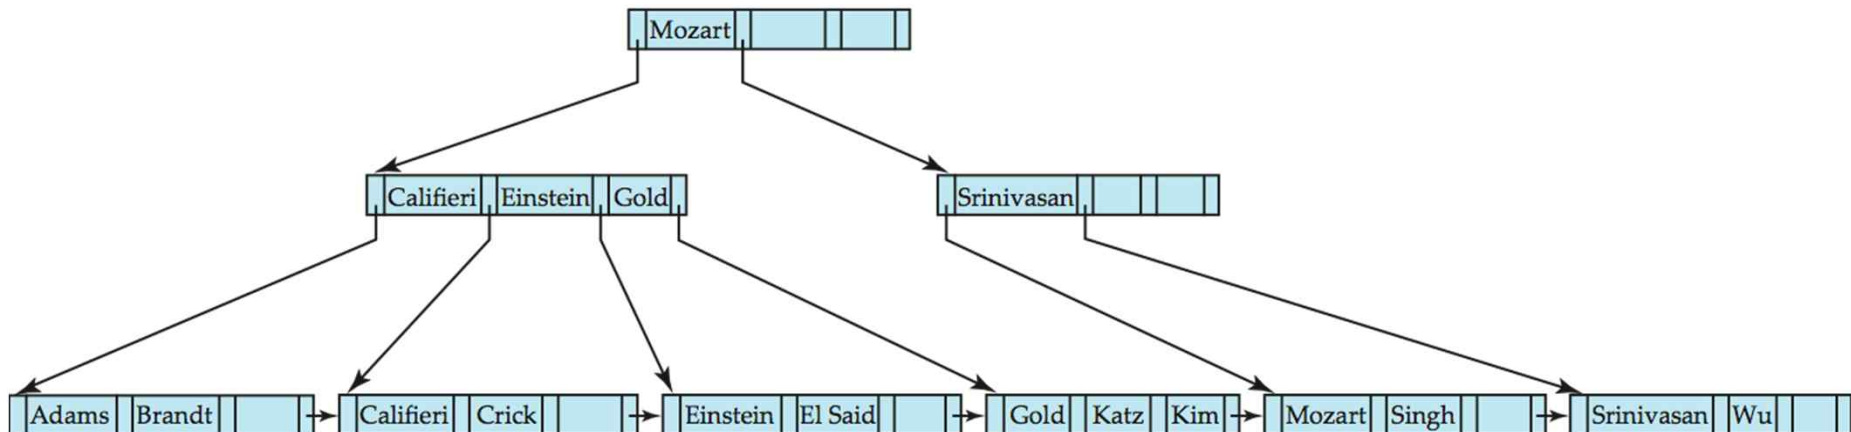
Observations about B⁺-trees

- Since the inter-node connections are done by pointers, “logically” close blocks need not be “physically” close
- The non-leaf levels of the B⁺-tree form a hierarchy of sparse indices
- The B⁺-tree contains a relatively small number of levels
 - ▶ Level below root has at least $2 * \lceil n/2 \rceil$ values
 - ▶ Next level has at least $2 * \lceil n/2 \rceil * \lceil n/2 \rceil$ values
 - ▶ .. etc.
- If there are K search-key values in the file, the tree height is no more than $\lceil \log_{\lceil n/2 \rceil}(K) \rceil$
- Thus searches can be conducted efficiently
- Insertions and deletions to the main file can be handled efficiently, as the index can be restructured in logarithmic time (as we shall see).



Queries on B⁺-Trees

- Find record with search-key value V
 1. $C = \text{root}$
 2. While C is not a leaf node {
 1. Let i be least value s.t. $V \leq K_i$.
 2. If no such exists, set $C = \text{last non-null pointer in } C$
 3. Else { if ($V = K_i$) set $C = P_{i+1}$
else set $C = P_i$ }
 3. Let i be least value s.t. $K_i = V$
 4. If there is such a value i , follow pointer P_i to the desired record
 5. Else no record with search-key value k exists





Queries on B⁺-Trees (Cont.)

- If there are K search-key values in the file, the height of the tree is no more than $\lceil \log_{\lceil n/2 \rceil}(K) \rceil$
- A node is generally the same size as a disk block, typically 4 kilobytes
 - n is typically around 100 (40 bytes per index entry)
- With 1 million search key values and $n = 100$
 - At most $\log_{50}(1,000,000) = 4$ nodes are accessed in a lookup
- Contrast this with a balanced binary tree with 1 million search key values
 - Around 20 nodes are accessed in a lookup
 - Difference is significant since every node access may need a disk I/O, costing around 20 milliseconds



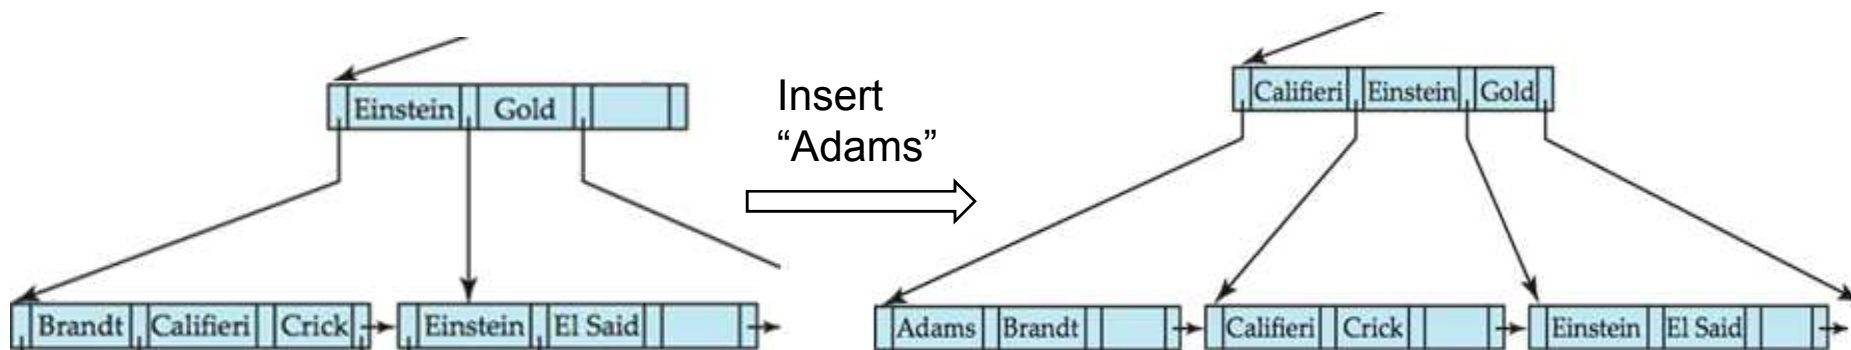
Updates on B⁺-Trees: Insertion

1. Find the leaf node in which the search-key value would appear
2. Add the record to the file
3. If there is room in the leaf node, insert (key-value, pointer) pair in the leaf node
4. Otherwise, split the node (along with the new (key-value, pointer) entry)



Leaf Node Split in B⁺-Trees

- Splitting a leaf node:
 - Take the n (search-key value, pointer) pairs (including the one being inserted) in sorted order
 - ▶ Place the first $\lceil n/2 \rceil$ in the original node, and the rest in a new node
 - ▶ Make the original node point to the new node
 - Let the new node be p , and let k be the least key value in p
 - ▶ Insert (k, p) in the parent of the node being split
 - If the parent is full, split it and **propagate** the split further up
- Splitting of nodes proceeds upwards till a node that is not full is found



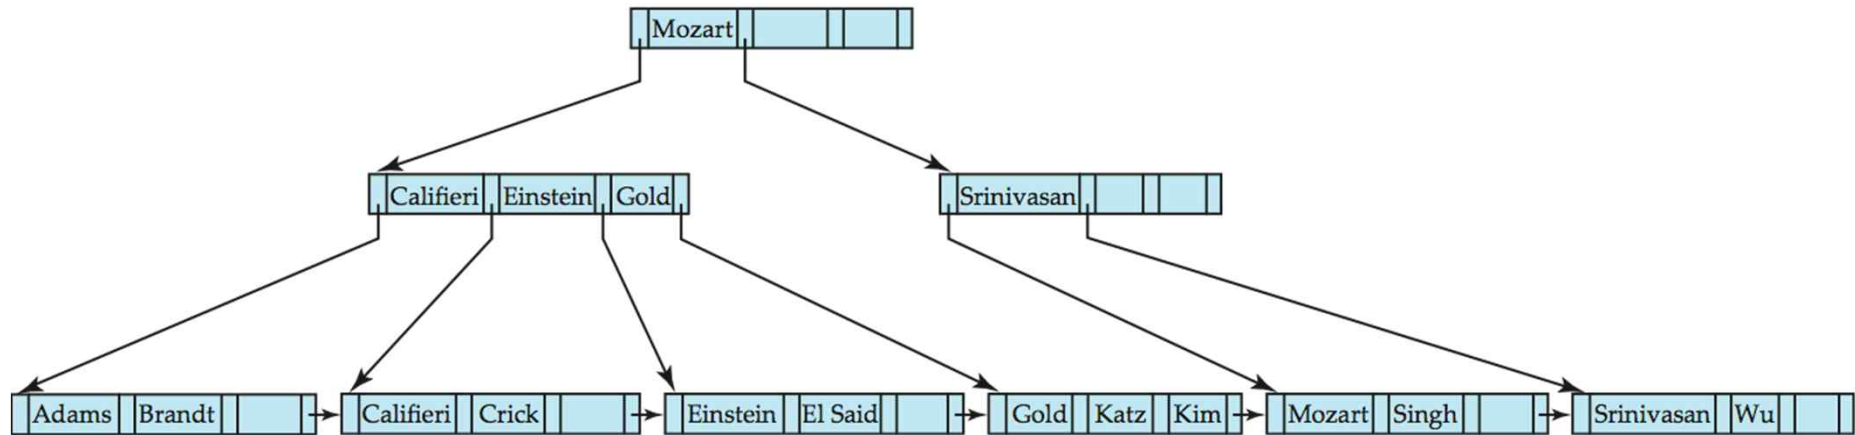


Non-Leaf Node Split in B⁺-Trees

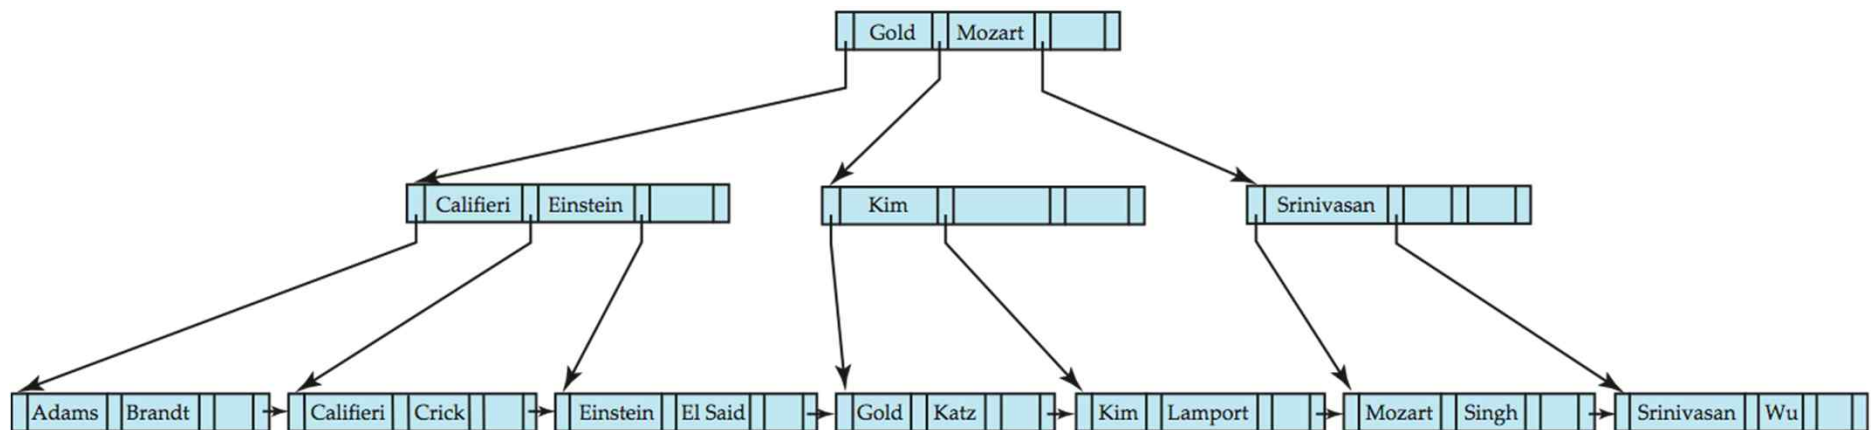
- Splitting a non-leaf node: when inserting (k,p) into an already full internal node N
 - Copy N to an in-memory area M that can hold N and (k,p)
 - Insert (k,p) into M
 - Copy $P_1, K_1, \dots, K_{\lceil n/2 \rceil - 1}, P_{\lceil n/2 \rceil}$ from M back into node N
 - Copy $P_{\lceil n/2 \rceil + 1}, K_{\lceil n/2 \rceil + 1}, \dots, K_{n-1}, P_n$ from M into newly allocated node N'
 - Insert $(K_{\lceil n/2 \rceil}, N')$ into parent N



B⁺-Tree Non-Leaf Node Split Example



After insertion of “Lamport”





Updates on B⁺-Trees: Deletion

1. Find the record to be deleted, and remove it from the main file
2. Remove (search-key value, pointer) from the leaf node
3. If the node has too few entries due to the removal,
merge siblings or **redistributed pointers** (next slide)
4. The node deletions may cascade upwards till a node which has $\lceil n/2 \rceil$ or more pointers is found
5. If the root node has only one pointer after deletion, it is deleted and the sole child becomes the root




Updates on B⁺-Trees: Deletion (Cont.)

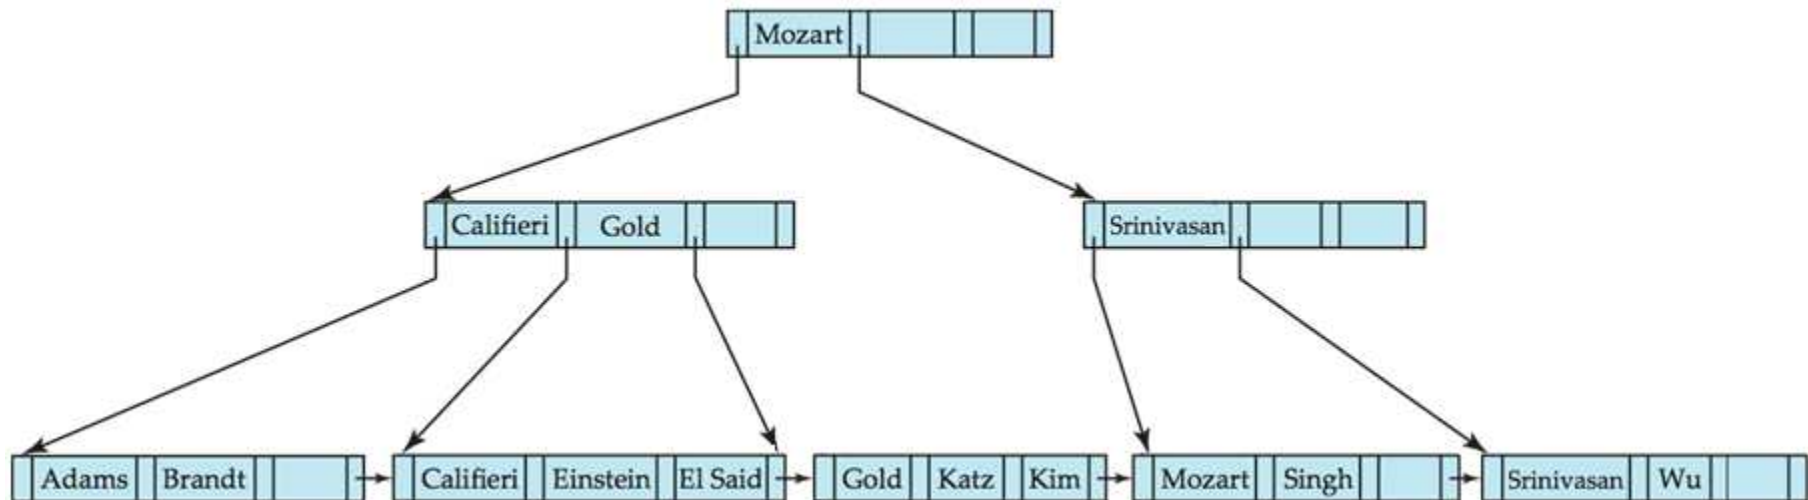
3. If the node has too few entries due to the removal,
 1. The entries in **the node and a sibling** fit into a single node – **merge siblings**:
 - ▶ Insert all the search-key values in the two nodes into a single node (the one on the left), and delete the other node
 - ▶ Delete the pair (K_{i-1}, P_i) , where P_i is the pointer to the deleted node, from its parent, recursively using the above procedure
 2. The entries in **the node and a sibling** do **not** fit into a single node – **redistribute pointers**:
 - ▶ **Redistribute** the pointers between the node and a sibling such that both have more than the minimum number of entries
 - ▶ Update the corresponding search-key value in the parent of the node
- Order to check an adjacent sibling
 - **Right sibling first**
 - If not possible, then, left sibling



Example: B⁺-Tree Deletion – Merge Siblings



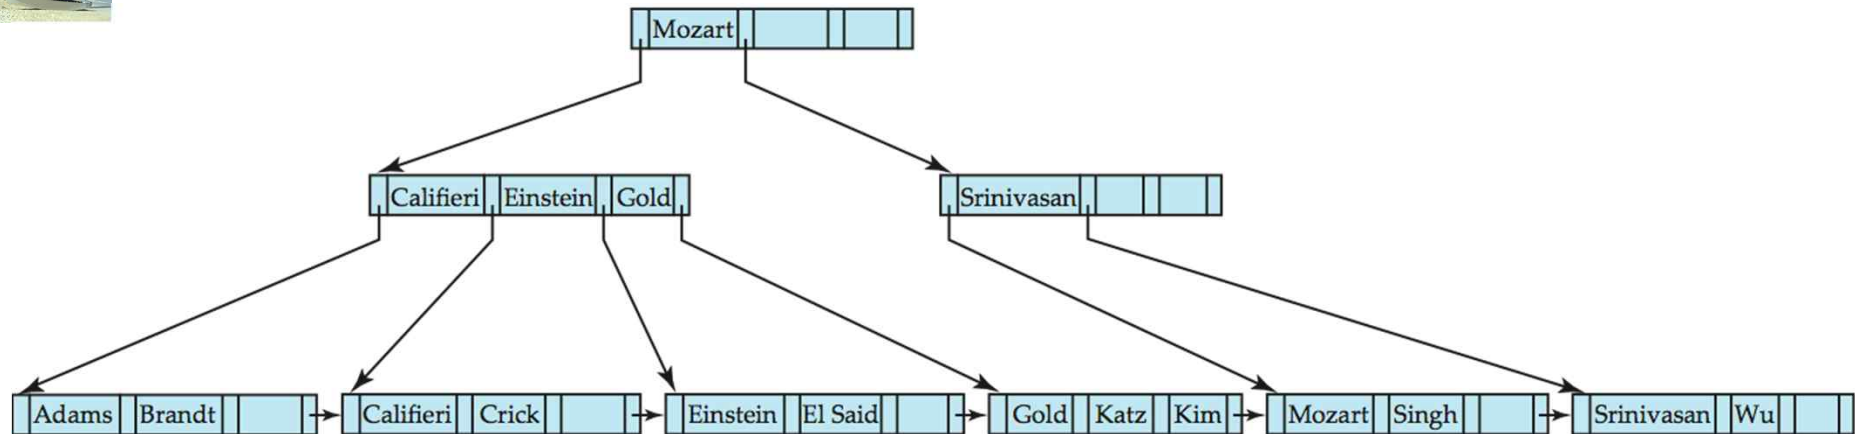
After deletion of “Crick”



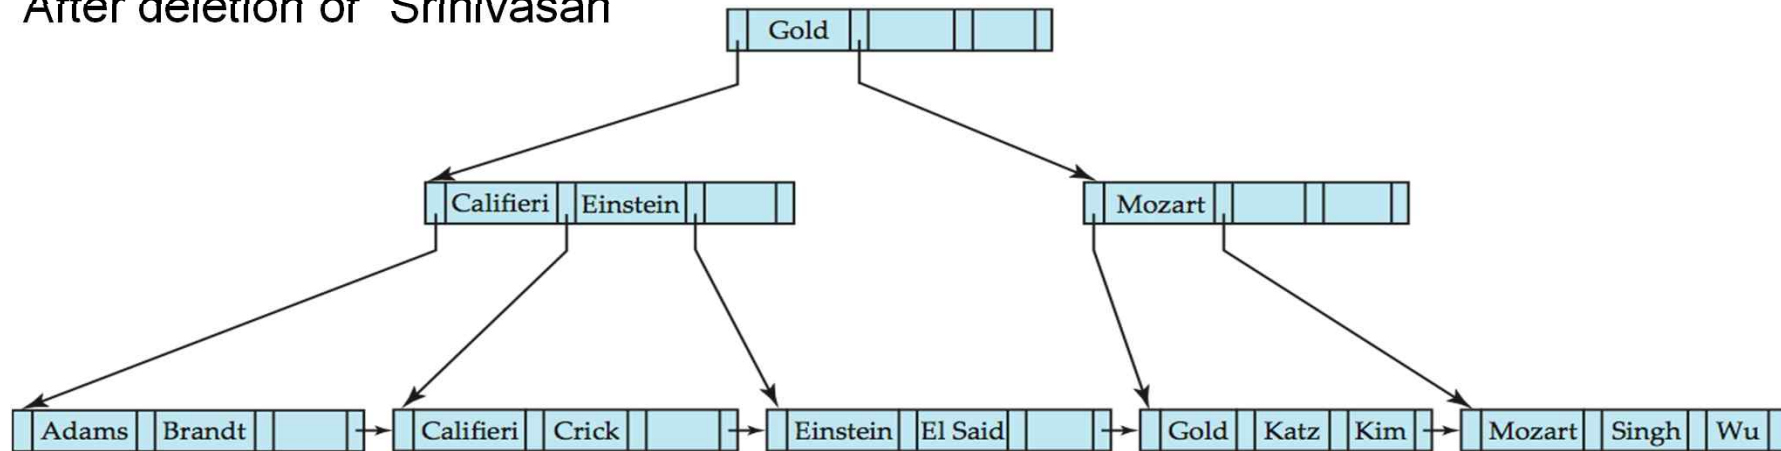
- Deleting “Crick” causes merging of under-full leaves



Example: B⁺-Tree Deletion – Merge Siblings (*)



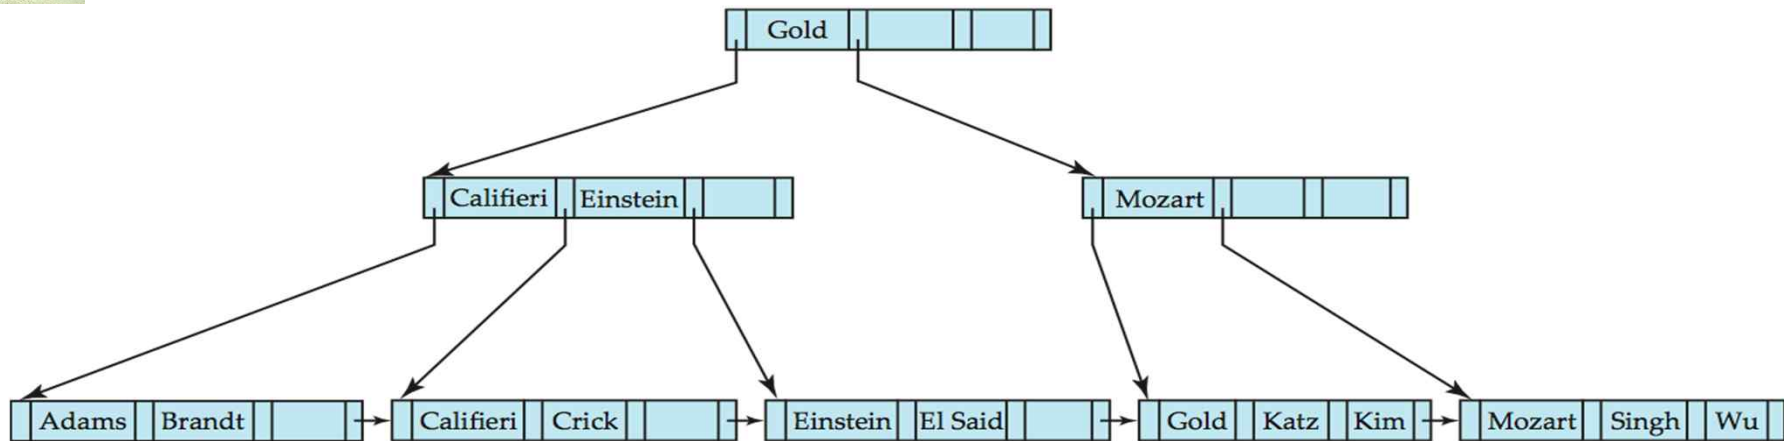
After deletion of “Srinivasan”



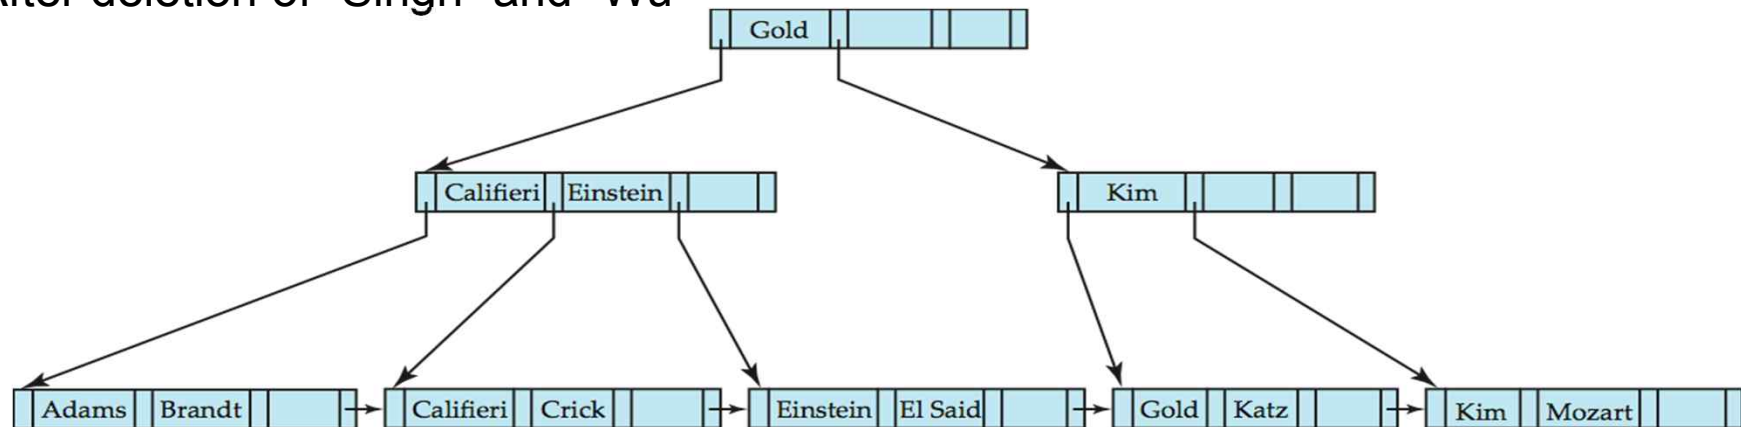
- Deleting “Srinivasan” causes merging of under-full leaves
- Parent node becomes **underfull**, and is merged with its sibling
 - Value separating two nodes (at the parent) is pulled down when merging
- Root node then has only one child, and is deleted, but re-created after merging



Example: B⁺-tree Deletion – Redistribute Pointers



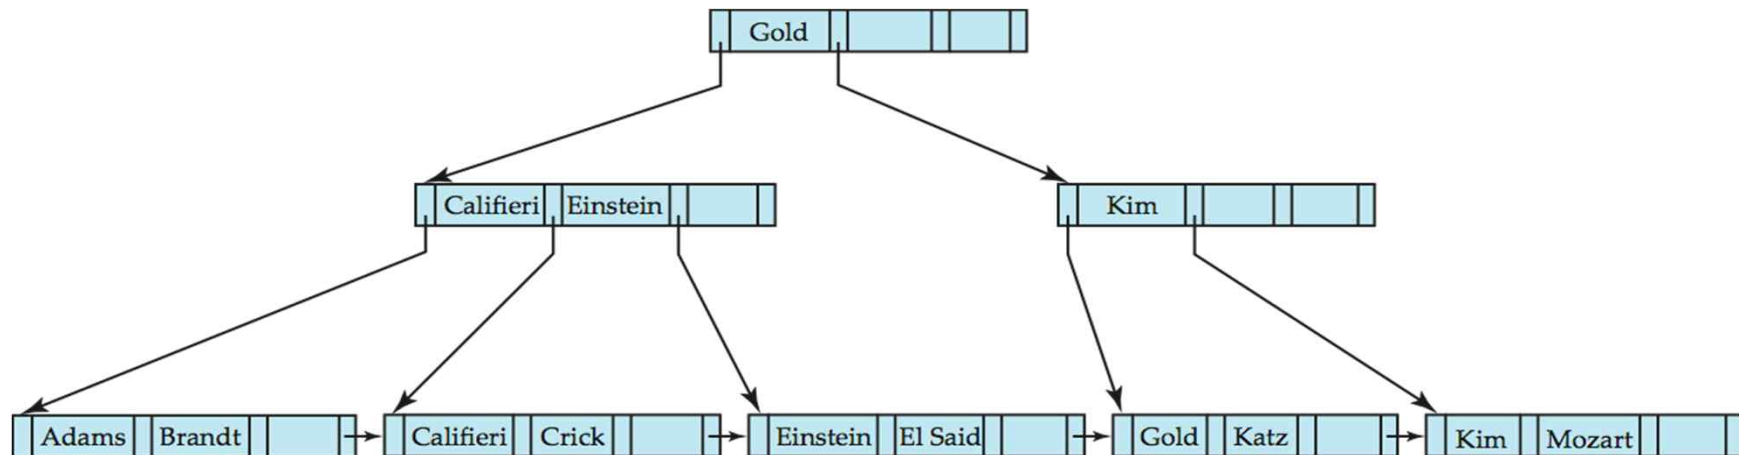
After deletion of “Singh” and “Wu”



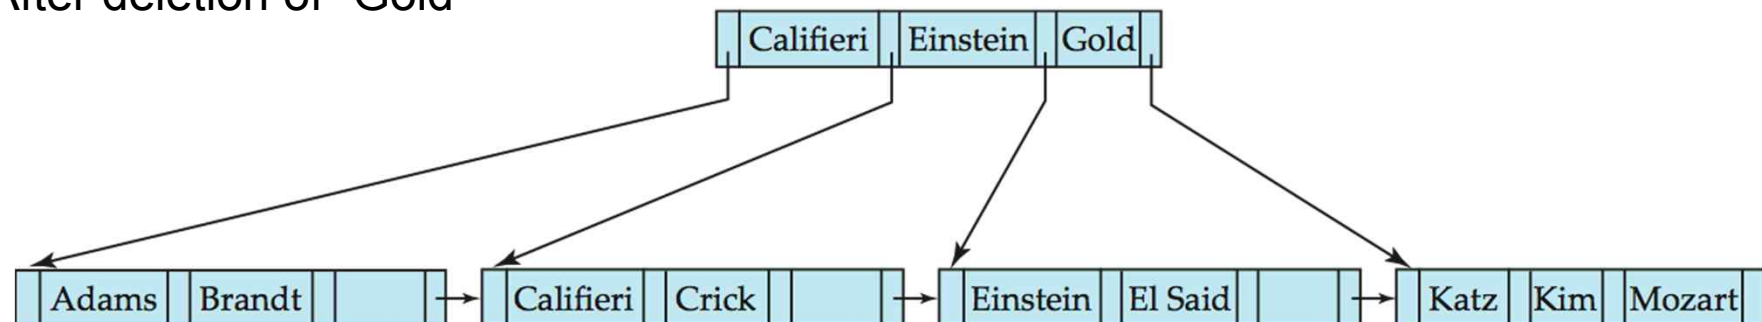
- Leaf containing Singh and Wu became **underfull**, and borrowed a value Kim from its left sibling
- Search-key value in the parent changes as a result



Example: B⁺-tree Deletion – Root Node Deletion



After deletion of “Gold”



- Node with Gold and Katz became **underfull**, and was merged with its sibling
- Parent node becomes **underfull**, and is merged with its sibling
 - Value separating two nodes (at the parent) is pulled down when merging
- Root node then has only one child, and is deleted



Handling Non-Unique Search Keys

- Buckets on separate block (bad idea)
- List of tuple pointers with each key
 - Extra code to handle long lists
 - Deletion of a tuple can be expensive if there are many duplicates on search key
 - Low space overhead, no extra cost for queries
- Make search key unique by adding a record-identifier
 - Extra storage overhead for keys
 - Simpler code for insertion/deletion
 - Widely used



Bulk Loading and Bottom-Up Build

- Inserting entries one-at-a-time into a B⁺-tree requires ≥ 1 IO per entry
 - assuming leaf level does not fit in memory
 - can be very inefficient for loading a large number of entries at a time (**bulk loading**)
- Efficient alternative 1:
 - Sort entries first (using efficient external-memory sort algorithms discussed later in Section 12.4)
 - Insert in sorted order
 - ▶ insertion will go to existing page (or cause a split)
 - ▶ much improved IO performance, but most leaf nodes half full
- Efficient alternative 2: **Bottom-up B⁺-tree construction**
 - As before sort entries
 - And then create tree layer-by-layer, starting with leaf level
 - ▶ details as an exercise
 - Implemented as part of bulk-load utility by most database systems



Multiple-Key Access

- Use **multiple indices** for certain types of queries.

- Example:

select *ID*

from *instructor*

where *dept_name* = "Finance" **and** *salary* = 80000

- Possible strategies for processing query using indices on single attributes:
 1. Use index on *dept_name* to find instructors with department name Finance; test *salary* = 80000
 2. Use index on *salary* to find instructors with a salary of \$80000; test *dept_name* = "Finance".
 3. Use *dept_name* index to find pointers to all records pertaining to the "Finance" department. Similarly use index on *salary*. Take **intersection** of both sets of pointers obtained.



Indices on Multiple Keys

- **Composite search keys** are search keys containing more than one attribute
 - E.g. (*dept_name*, *salary*)
- Lexicographic ordering: $(a_1, a_2) < (b_1, b_2)$ if either
 - $a_1 < b_1$, or
 - $a_1 = b_1$ and $a_2 < b_2$
- Suppose we have an **index** on combined search-key (*dept_name*, *salary*).
- Can efficiently handle
 - where *dept_name* = "Finance" and *salary* = 80000**
 - Fetch only records that satisfy both conditions
- Can also efficiently handle
 - where *dept_name* = "Finance" and *salary* < 80000**
- But cannot efficiently handle
 - where *dept_name* < "Finance" and *balance* = 80000**
 - May fetch many records that satisfy the first but not the second condition



Static Hashing

■ Bucket

- A unit of storage containing one or more records
- Typically a disk block

■ Hash file organization

- The bucket of a record is directly obtained from its search-key value using a **hash function**

■ Hash function h

- A function **from** the set of all search-key values K **to** the set of all bucket addresses B
- is used to locate records for access, insertion as well as deletion.

■ Records with different search-key values may be mapped to the same bucket

- Thus entire bucket has to be searched sequentially to locate a record



Example of Hash File Organization

Hash file organization of *instructor* file, using *dept_name* as key

- # of buckets = 10
- The binary representation of the *i*th character is assumed to be the integer *i*
- The hash function returns the sum of the binary representations of the characters modulo 10
 - E.g. $h(\text{Music}) = 1$
 $h(\text{History}) = 2$
 $h(\text{Physics}) = 3$
 $h(\text{Elec. Eng.}) = 3$

bucket 0

bucket 1

15151	Mozart	Music	40000

bucket 2

32343	El Said	History	80000
58583	Califieri	History	60000

bucket 3

22222	Einstein	Physics	95000
33456	Gold	Physics	87000
98345	Kim	Elec. Eng.	80000

bucket 4

12121	Wu	Finance	90000
76543	Singh	Finance	80000

bucket 5

76766	Crick	Biology	72000

bucket 6

10101	Srinivasan	Comp. Sci.	65000
45565	Katz	Comp. Sci.	75000
83821	Brandt	Comp. Sci.	92000

bucket 7




Hash Functions

- Worst hash function
 - All search-key values are mapped to **the same bucket**
 - Access time is proportional to the number of search-key values in the file
- Ideal hash function
 - **Uniform**: each bucket is assigned **the same number of search-key values** from the set of *all* possible values
 - **Random**: each bucket will have **the same number of records assigned to it** irrespective of the *actual distribution* of search-key values in the file
- Typical hash functions perform computation on the internal binary representation of the search-key



Handling of Bucket Overflows

- **Bucket overflow** can occur because of
 - Insufficient buckets
 - Skew in distribution of records. This can occur due to two reasons:
 - ▶ multiple records have same search-key value
 - ▶ chosen hash function produces non-uniform distribution of key values
- Although the probability of bucket overflow can be reduced, it cannot be eliminated; it is handled by using **overflow buckets**.
- **Overflow chaining**
 - The overflow buckets of a given bucket are chained together in a linked list



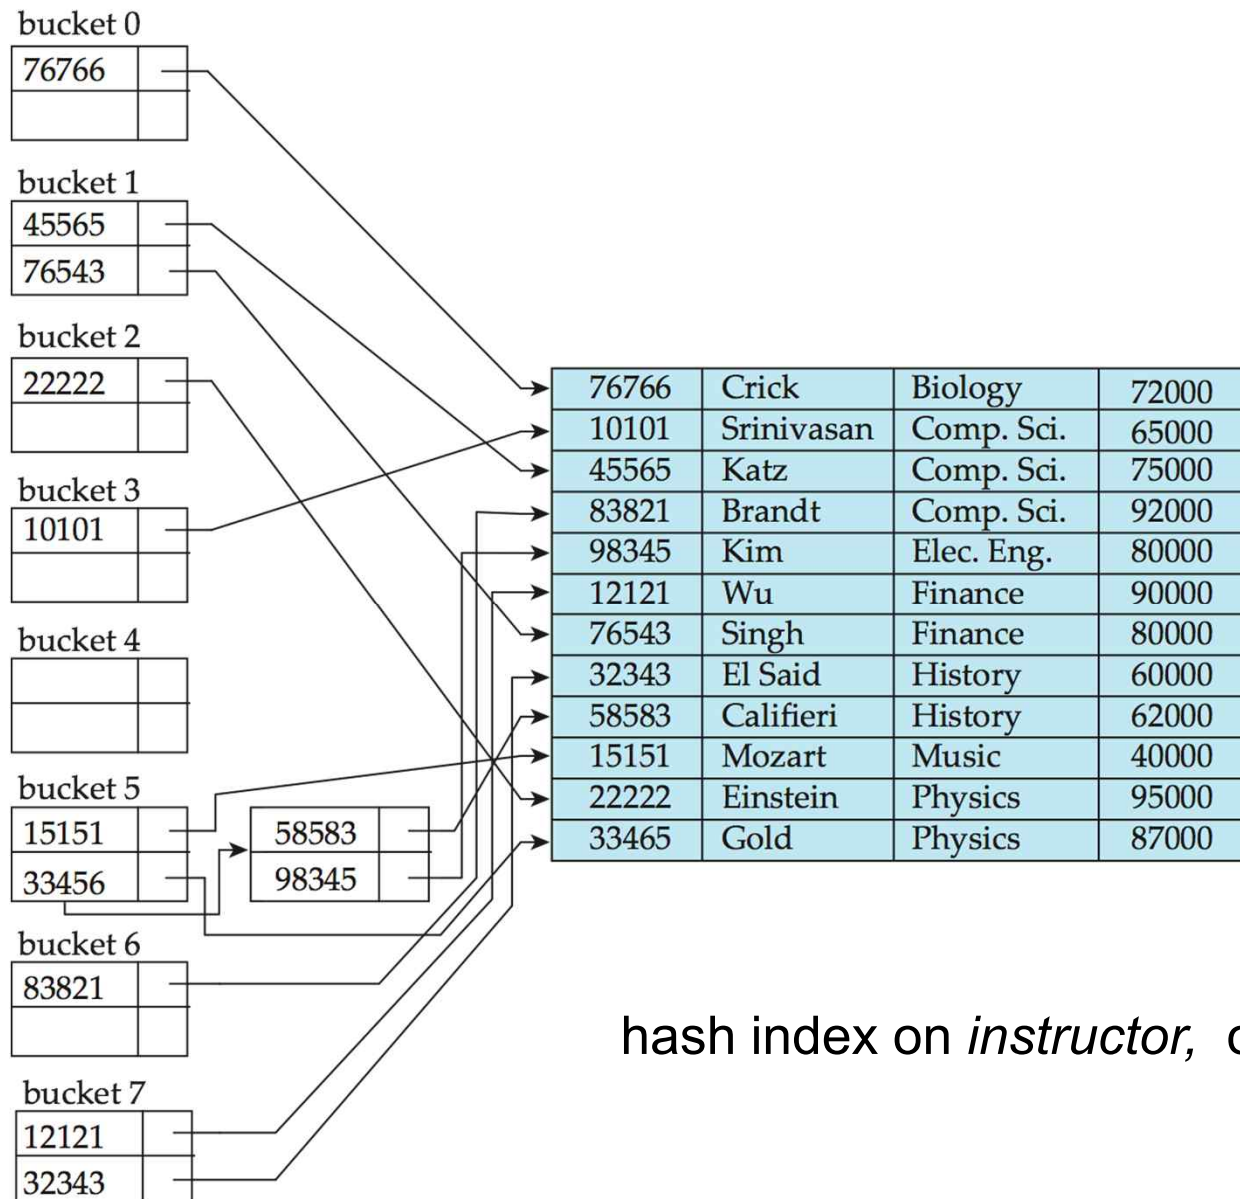


Hash Indices

- A **hash index** organizes the search keys, with their associated record pointers, into a hash file structure.
- Strictly speaking, hash indices are always secondary indices
 - If the file itself is organized using hashing, a separate primary hash index on it using the same search-key is unnecessary
 - However, we use the term **hash index** to refer to both secondary index structures and hash organized files



Example of Hash Index



hash index on *instructor*, on attribute *ID*



Deficiencies of Static Hashing

- In static hashing, function h maps search-key values to a fixed set of B of bucket addresses
 - Databases grow or shrink with time
 - If initial number of buckets is too small, and file grows, performance will degrade due to too much overflows
 - If space is allocated for anticipated growth, a significant amount of space will be wasted initially (and buckets will be underfull)
 - If database shrinks, again space will be wasted
- One solution: periodic re-organization of the file with a new hash function
 - Expensive, disrupts normal operations
- Better solution: allow the number of buckets to be modified dynamically
 - Dynamic hashing is not covered in this class



Comparison of Ordered Indexing and Hashing

- Cost of periodic re-organization
- Relative frequency of insertions and deletions
- Is it desirable to optimize average access time at the expense of worst-case access time?
- Expected type of queries:
 - Hashing is generally better at retrieving records having a specified value of the key (**point query**)
 - If **range queries** are common, ordered indices are to be preferred
- In practice:
 - PostgreSQL supports hash indices, but discourages use due to poor performance
 - Oracle supports static hash organization, but not hash indices
 - SQLServer supports only B⁺-trees



Bitmap Indices

- **Bitmap**: simply an array of bits
- **Bitmap index**: a specialized type of index designed for **efficient querying on multiple keys**
- In the simplest form, a bitmap index on an attribute has **a bitmap for each value of the attribute**
 - Bitmap has as many bits as records
 - In a bitmap for value v , the bit for a record is 1 if the record has the value v for the attribute, and is 0 otherwise

record number	<i>ID</i>	<i>gender</i>	<i>income_level</i>	Bitmaps for <i>gender</i>		Bitmaps for <i>income_level</i>	
				m	10010		
				f	01101	L1	10100
0	76766	m	L1			L2	01000
1	22222	f	L2			L3	00001
2	12121	f	L1			L4	00010
3	15151	m	L4			L5	00000
4	58583	f	L3				



Bitmap Indices (Cont.)

- Applicable on attributes that take on a relatively small number of distinct values
 - E.g. gender, country, state, ...
 - E.g. income-level (income broken up into a small number of levels such as 0-9999, 10000-19999, 20000-50000, 50000- infinity)
- Bitmap indices generally very small compared with relation size
 - E.g. if record is 100 bytes, space for a single bitmap is 1/800 of space used by relation
 - ▶ If # of distinct attribute values is 8, bitmap is only 1% of relation size



Bitmap Indices (Cont.)

- Bitmap indices are useful for **queries on multiple attributes**
 - not particularly useful for single attribute queries
- Queries are answered using bitmap operations
 - Intersection (and): e.g., $100110 \text{ AND } 110011 = 100010$
 - Union (or): e.g., $100110 \text{ OR } 110011 = 110111$
 - Complementation (not): e.g., $100110 \text{ NOT } 100110 = 011001$
- **Each operation takes two bitmaps of the same size** and applies the operation on corresponding bits to get the result bitmap
 - E.g. males with income level L1: $10010 \text{ AND } 10100 = 10000$
 - ▶ Can then retrieve required tuples
 - ▶ Counting number of matching tuples is even faster



Index Definition in SQL

- Create an index

create index <index-name> **on** <relation-name>
(<attribute-list>)

E.g.: **create index** *dept_index* **on** *instructor* (*dept_name*)

- Use **create unique index** to indirectly specify and enforce the condition that the search key is a candidate key
 - Not really required if SQL **unique** integrity constraint is supported

- To drop an index

drop index <index-name>

- Most database systems allow specification of type of index, and clustering



End of Chapter 11

Database System Concepts, 6th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use