



Chapter 22: Object-Based Databases

Database System Concepts, 6th Ed.

- Chapter 1: Introduction
- **Part 1: Relational databases**
 - Chapter 2: Introduction to the Relational Model
 - Chapter 3: Introduction to SQL
 - Chapter 4: Intermediate SQL
 - Chapter 5: Advanced SQL
 - Chapter 6: Formal Relational Query Languages
- **Part 2: Database Design**
 - Chapter 7: Database Design: The E-R Approach
 - Chapter 8: Relational Database Design
 - Chapter 9: Application Design
- **Part 3: Data storage and querying**
 - Chapter 10: Storage and File Structure
 - Chapter 11: Indexing and Hashing
 - Chapter 12: Query Processing
 - Chapter 13: Query Optimization
- **Part 4: Transaction management**
 - Chapter 14: Transactions
 - Chapter 15: Concurrency control
 - Chapter 16: Recovery System
- **Part 5: System Architecture**
 - Chapter 17: Database System Architectures
 - Chapter 18: Parallel Databases
 - Chapter 19: Distributed Databases
- **Part 6: Data Warehousing, Mining, and IR**
 - Chapter 20: Data Mining
 - Chapter 21: Information Retrieval
- **Part 7: Specialty Databases**
 - [Chapter 22: Object-Based Databases](#)
 - Chapter 23: XML
- **Part 8: Advanced Topics**
 - Chapter 24: Advanced Application Development
 - Chapter 25: Advanced Data Types
 - Chapter 26: Advanced Transaction Processing
- **Part 9: Case studies**
 - Chapter 27: PostgreSQL
 - Chapter 28: Oracle
 - Chapter 29: IBM DB2 Universal Database
 - Chapter 30: Microsoft SQL Server
- **Online Appendices**
 - Appendix A: Detailed University Schema
 - Appendix B: Advanced Relational Database Model
 - Appendix C: Other Relational Query Languages
 - Appendix D: Network Model
 - Appendix E: Hierarchical Model



Chapter 22: Object-Based Databases

- 22.1 Overview
- 22.2 Complex Data Types
- 22.3 Structured Data Types and Inheritance in SQL
- 22.4 Table Inheritance
- 22.5 Array and Multiset Types in SQL
- 22.6 Object Identity and Reference Types in SQL
- 22.7 Implementing O-R Features
- 22.8 Persistent Programming Languages
- 22.9 Object-Relational Mapping
- 22.10 Object-Oriented versus Object-Relational



Object-Relational Data Models

- Extend the relational data model by including **object orientation and constructs** to deal with added data types
- Allow attributes of tuples to have **complex types**, including non-atomic values such as nested relations
- Preserve **relational foundations**, in particular the declarative access to data, while extending modeling power
- **Upward compatibility** with existing relational languages



Chapter 22: Object-Based Databases

- 22.1 Overview
- 22.2 Complex Data Types
- 22.3 Structured Data Types and Inheritance in SQL
- 22.4 Table Inheritance
- 22.5 Array and Multiset Types in SQL
- 22.6 Object Identity and Reference Types in SQL
- 22.7 Implementing O-R Features
- 22.8 Persistent Programming Languages
- 22.9 Object-Relational Mapping
- 22.10 Object-Oriented versus Object-Relational



Complex Data Types

- Motivation:
 - Permit non-atomic domains (atomic \equiv indivisible)
 - Example of non-atomic domain: [set of integers](#) or [set of tuples](#)
 - Allows more intuitive modeling for applications with complex data

- Intuitive definition:
 - allow relations whenever we allow atomic (scalar) values
 - ▶ [relations within relations](#)
 - Retains mathematical foundation of relational model
 - Violates first normal form



Example of a Nested Relation

- Example: library information system
- Each book has
 - title,
 - a list (array) of authors,
 - Publisher, with subfields *name* and *branch*, and
 - a set of keywords
- Non-1NF relation *books*

<i>title</i>	<i>author_array</i>	<i>publisher</i>	<i>keyword_set</i>
		(<i>name, branch</i>)	
Compilers	[Smith, Jones]	(McGraw-Hill, New York)	{parsing, analysis}
Networks	[Jones, Frick]	(Oxford, London)	{Internet, Web}



1NF Version of Nested Relation

- 1NF version of books

<i>title</i>	<i>author</i>	<i>pub-name</i>	<i>pub-branch</i>	<i>keyword</i>
Compilers	Smith	McGraw-Hill	New York	parsing
Compilers	Jones	McGraw-Hill	New York	parsing
Compilers	Smith	McGraw-Hill	New York	analysis
Compilers	Jones	McGraw-Hill	New York	analysis
Networks	Jones	Oxford	London	Internet
Networks	Frick	Oxford	London	Internet
Networks	Jones	Oxford	London	Web
Networks	Frick	Oxford	London	Web

flat-books



4NF Decomposition of Nested Relation

- Suppose for simplicity that title uniquely identifies a book
 - In real world ISBN is a unique identifier
- Decompose *books* into 4NF using the schemas:
 - $(title, author, position)$
 - $(title, keyword)$
 - $(title, pub_name, pub_branch)$
- 4NF design requires users to include joins in their queries.

<i>title</i>	<i>author</i>	<i>position</i>
Compilers	Smith	1
Compilers	Jones	2
Networks	Jones	1
Networks	Frick	2

authors

<i>title</i>	<i>keyword</i>
Compilers	parsing
Compilers	analysis
Networks	Internet
Networks	Web

keywords

<i>title</i>	<i>pub_name</i>	<i>pub_branch</i>
Compilers	McGraw-Hill	New York
Networks	Oxford	London

books4



4NF Decomposition of Nested Relation

- Dependencies in *doc*

$title \twoheadrightarrow author$ (MVD)

$title \twoheadrightarrow keyword$

$title \rightarrow pub_name, pub_branch$

MVD: multi-valued dependency;

$X \twoheadrightarrow Y$ means that a set of Y values is associated with each X value

- Decomposed version

- 4NF (BCNF extended to include MVD)

- Loose 1-to-1 correspondence between a tuple and a doc

<i>title</i>	<i>author</i>	<i>position</i>
Compilers	Smith	1
Compilers	Jones	2
Networks	Jones	1
Networks	Frick	2

authors

<i>title</i>	<i>keyword</i>
Compilers	parsing
Compilers	analysis
Networks	Internet
Networks	Web

keywords

<i>title</i>	<i>pub-name</i>	<i>pub-branch</i>
Compilers	McGraw-Hill	New York
Networks	Oxford	London

books4



Complex Types and SQL

- Extensions introduced in SQL:1999 to support complex types:
 - Collection and large object types
 - ▶ Nested relations are an example of collection types
 - Structured types
 - ▶ Nested record structures like composite attributes
 - Inheritance
 - Other object orientation features
 - ▶ Including object identifiers and references
- Not fully implemented in any database system currently
 - But some features are present in each of the major commercial database systems
 - ▶ Read the manual of your database system to see what it supports



Chapter 22: Object-Based Databases

- 22.1 Overview
- 22.2 Complex Data Types
- 22.3 Structured Data Types and Inheritance in SQL
- 22.4 Table Inheritance
- 22.5 Array and Multiset Types in SQL
- 22.6 Object Identity and Reference Types in SQL
- 22.7 Implementing O-R Features
- 22.8 Persistent Programming Languages
- 22.9 Object-Relational Mapping
- 22.10 Object-Oriented versus Object-Relational



Structured Types and Inheritance in SQL [1/2]

- **Structured types** (a.k.a. **user-defined types**) can be declared and used in SQL

create type *Name* **as**

(*firstname* **varchar**(20),
 lastname **varchar**(20))

final

create type *Address* **as**

(*street* **varchar**(20),
 city **varchar**(20),
 zipcode **varchar**(20))

not final

- Note: **final** and **not final** indicate whether subtypes can be created

- Structured types can be used to create tables with composite attributes

create table *person* (

name *Name*,
 address *Address*,
 dateOfBirth **date**)

- Dot notation used to reference components: *name.firstname*



Structured Types and Inheritance in SQL [2/2]

■ User-defined row types

```
create type PersonType as (  
    name Name,  
    address Address,  
    dateOfBirth date)  
not final
```

- Can then create a table whose rows are a user-defined type
create table *customer* **of** *CustomerType*

- Alternative using **unnamed row types**.

```
create table person_r(  
    name    row(firstname varchar(20),  
                lastname varchar(20)),  
    address row(street    varchar(20),  
                city       varchar(20),  
                zipcode   varchar(20)),  
    dateOfBirth date)
```



Methods

- Can add a method declaration with a structured type
method *ageOnDate* (*onDate* **date**)
 returns **interval year**
- Method body is given separately
create instance method *ageOnDate* (*onDate* **date**)
 returns **interval year**
 for *CustomerType*
 begin
 return *onDate* - **self.dateOfBirth**;
 end
- We can now find the age of each customer:
 select *name.lastname*, *ageOnDate* (**current_date**)
 from *customer*



Constructor Functions

- **Constructor functions** are used to create values of structured types
- E.g.
create function *Name*(*firstname* **varchar**(20), *lastname* **varchar**(20))
returns *Name*
begin
 set **self**.*firstname* = *firstname*;
 set **self**.*lastname* = *lastname*;
end
- To create a value of type *Name*, we use
 new *Name*('John', 'Smith')
- Normally used in insert statements
insert into *Person* **values**
 (**new** *Name*('John', 'Smith'),
 new *Address*('20 Main St', 'New York', '11001'),
 date '1960-8-22');



Type Inheritance

- Suppose that we have the following type definition for people:

```
create type Person  
  (name varchar(20),  
   address varchar(20))
```

- Using **inheritance** to define the student and teacher types

```
create type Student under Person  
  (degree varchar(20),  
   department varchar(20))  
create type Teacher under Person  
  (salary integer,  
   department varchar(20))
```

- Subtypes can redefine methods by using **overriding method** in place of **method** in the method declaration



Multiple Type Inheritance

- SQL:1999 and SQL:2003 **do not support multiple inheritance**
- If our type system supports multiple inheritance, we can define a type for teaching assistant as follows:
create type *Teaching Assistant*
under *Student, Teacher*
- To avoid a conflict between the two occurrences of *department* we can rename them
create type *Teaching Assistant*
under
Student **with** (*department as student_dept*),
Teacher **with** (*department as teacher_dept*)
- Each value must have a **most-specific type**



Chapter 22: Object-Based Databases

- 22.1 Overview
- 22.2 Complex Data Types
- 22.3 Structured Data Types and Inheritance in SQL
- 22.4 Table Inheritance
- 22.5 Array and Multiset Types in SQL
- 22.6 Object Identity and Reference Types in SQL
- 22.7 Implementing O-R Features
- 22.8 Persistent Programming Languages
- 22.9 Object-Relational Mapping
- 22.10 Object-Oriented versus Object-Relational



Table Inheritance

- Tables created from subtypes can further be specified as **subtables**
- E.g. **create table** *people* **of** *Person*;
 create table *students* **of** *Student* **under** *people*;
 create table *teachers* **of** *Teacher* **under** *people*;
- Tuples added to a subtable are automatically visible to queries on the supertable
 - E.g. query on *people* also sees *students* and *teachers*.
 - Similarly updates/deletes on *people* also result in updates/deletes on subtables
 - To override this behaviour, use “**only people**” in query
- Conceptually, **multiple inheritance is possible with tables**
 - e.g. *teaching_assistants* under *students* and *teachers*
 - *But is not supported in SQL currently*
 - ▶ So we cannot create a person (tuple in *people*) who is both a student and a teacher



Consistency Requirements for Subtables

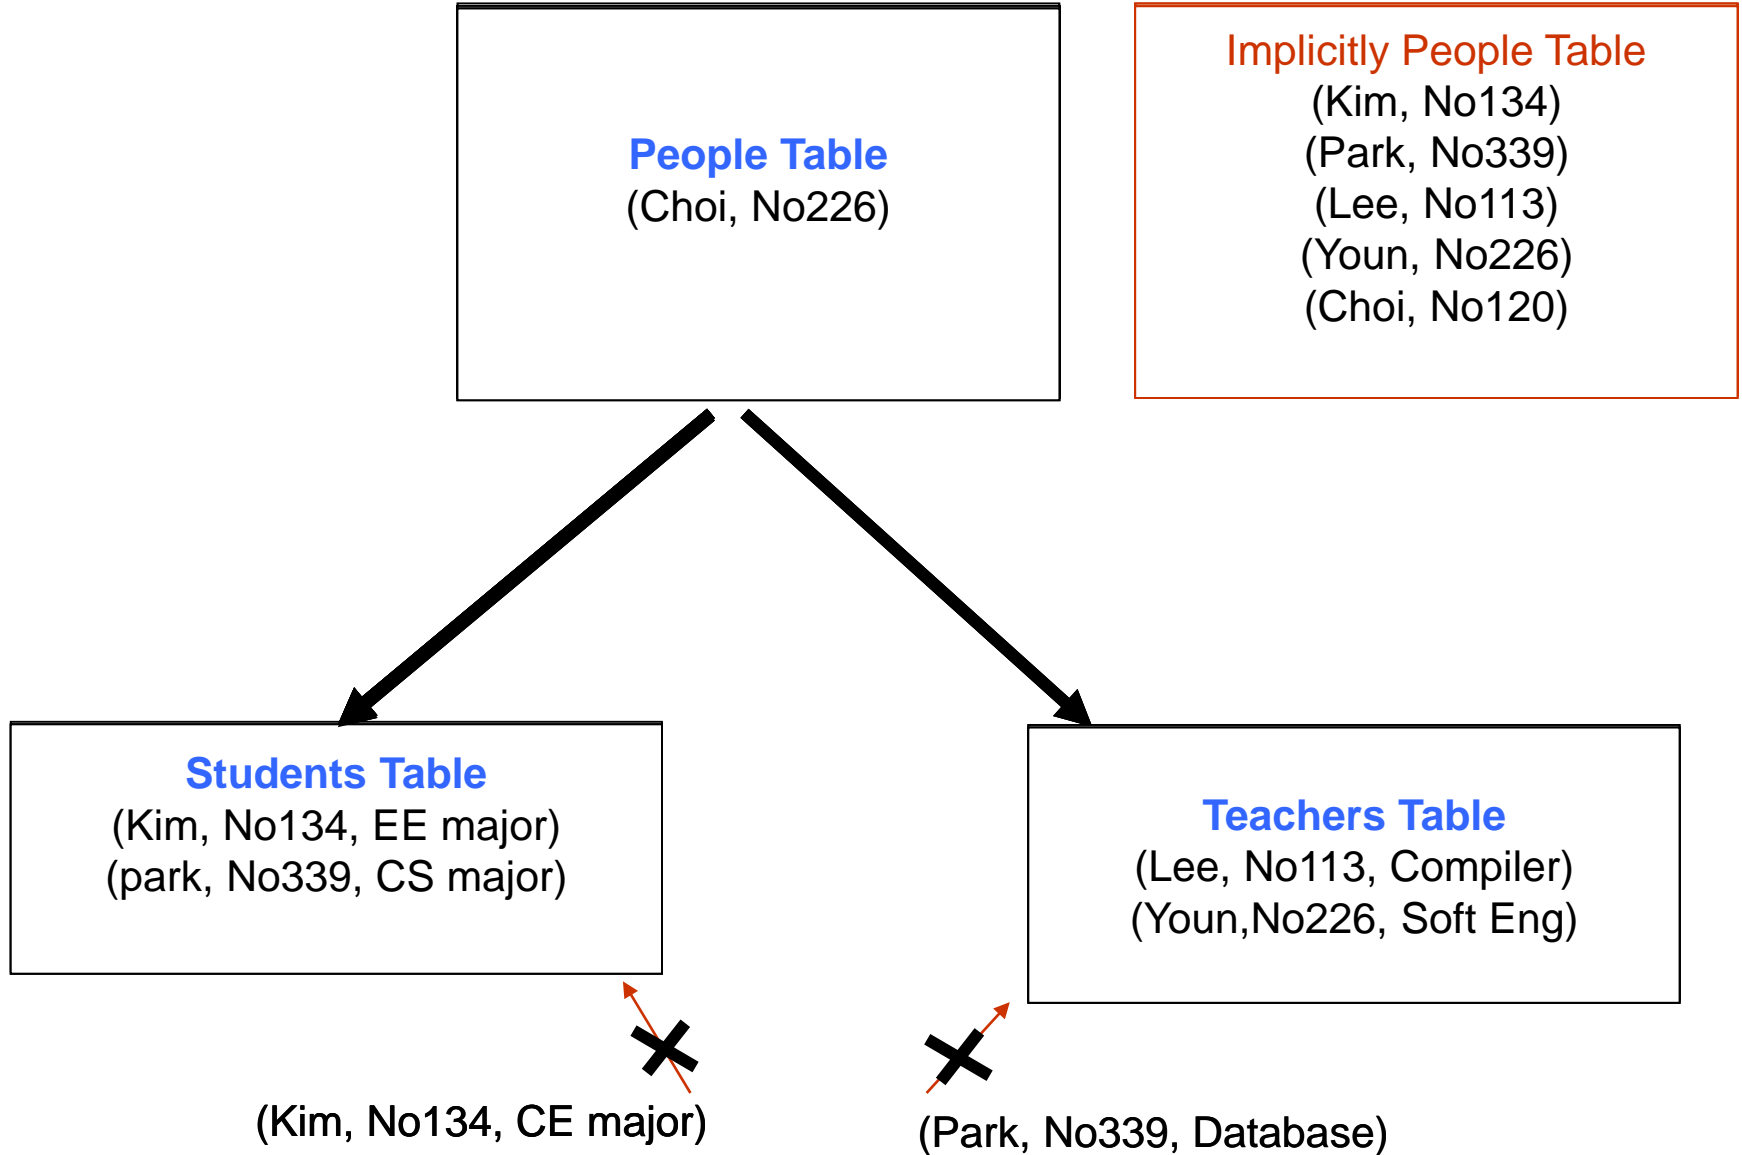
- Consistency requirements on subtables and supertables.
 - Each tuple of the supertable (e.g. *people*) can correspond to **at most one tuple** in each of the subtables (e.g. *students* and *teachers*)
 - Additional constraint in SQL:1999:

All tuples corresponding to each other (that is, with the same values for inherited attributes) must be derived from one tuple (inserted into one table).

 - ▶ That is, each entity must have **a most specific type**
 - ▶ We cannot have a tuple in *people* corresponding to a tuple each in *students* and *teachers*



Subtable Consistency





Chapter 22: Object-Based Databases

- 22.1 Overview
- 22.2 Complex Data Types
- 22.3 Structured Data Types and Inheritance in SQL
- 22.4 Table Inheritance
- 22.5 Array and Multiset Types in SQL
- 22.6 Object Identity and Reference Types in SQL
- 22.7 Implementing O-R Features
- 22.8 Persistent Programming Languages
- 22.9 Object-Relational Mapping
- 22.10 Object-Oriented versus Object-Relational



Array and Multiset Types in SQL

- Example of array and multiset declaration:

create type *Publisher* **as**

 (*name* **varchar**(20),
 branch **varchar**(20));

create type *Book* **as**

 (*title* **varchar**(20),
 author_array **varchar**(20) **array** [10],
 pub_date **date**,
 publisher *Publisher*,
 keyword-set **varchar**(20) **multiset**);

create table *books* **of** *Book*;



Creation of Collection Values

- Array construction

array ['Silberschatz', 'Korth', 'Sudarshan']

- Multisets

multiset ['computer', 'database', 'SQL']

- To create a tuple of the type defined by the books relation:

('Compilers', **array**['Smith', 'Jones'],
 new *Publisher* ('McGraw-Hill', 'New York'),
 multiset ['parsing', 'analysis'])

- To insert the preceding tuple into the relation books

insert into *books*
values

('Compilers', **array**['Smith', 'Jones'],
 new *Publisher* ('McGraw-Hill', 'New York'),
 multiset ['parsing', 'analysis']);



Unnesting

- The transformation of a nested relation into a form with fewer (or no) relation-valued attributes is called **unnesting**
- E.g.

```
select title, A as author, publisher.name as pub_name,  
       publisher.branch as pub_branch, K.keyword  
from books as B, unnest(B.author_array) as A (author),  
       unnest (B.keyword_set) as K (keyword)
```
- Result relation *flat_books*

<i>title</i>	<i>author</i>	<i>pub_name</i>	<i>pub_branch</i>	<i>keyword</i>
Compilers	Smith	McGraw-Hill	New York	parsing
Compilers	Jones	McGraw-Hill	New York	parsing
Compilers	Smith	McGraw-Hill	New York	analysis
Compilers	Jones	McGraw-Hill	New York	analysis
Networks	Jones	Oxford	London	Internet
Networks	Frick	Oxford	London	Internet
Networks	Jones	Oxford	London	Web
Networks	Frick	Oxford	London	Web



Querying Collection-Valued Attributes

- To find all books that have the word “database” as a keyword,

```
select title  
from books  
where 'database' in (unnest(keyword-set))
```

- We can access individual elements of an array by using indices

E.g.: If we know that a particular book has three authors, we could write:

```
select author-array[1], author-array[2], author-array[3]  
from books  
where title = 'Database System Concepts'
```

- To get a relation containing pairs of the form “title, author-name” for each book and each author of the book

```
select B.title, A.author  
from books as B, unnest (B.author-array) as A (author)
```

- To retain ordering information we add a **with ordinality** clause

```
select B.title, A.author, A.position  
from books as B, unnest (B.author-array) with ordinality as A (author, position)
```



Nesting [1/3]

- **Nesting** is the opposite of unnesting, creating a collection-valued attribute
- Nesting can be done in a manner similar to aggregation, but using the function **collect()** in place of an aggregation operation, to create a multiset
- To nest the *flat_books* relation on the attribute *keyword*:

```
select title, author, Publisher (pub_name, pub_branch ) as publisher,  
       collect (keyword) as keyword_set  
from flat_books  
groupby title, author, publisher
```
- To nest on both authors and keywords:

```
select title, collect (author ) as author_set,  
       Publisher (pub_name, pub_branch) as publisher,  
       collect (keyword ) as keyword_set  
from flat_books  
group by title, publisher
```



Nesting [2/3]

<i>title</i>	<i>author</i>	<i>pub-name</i>	<i>pub-branch</i>	<i>keyword</i>
Compilers	Smith	McGraw-Hill	New York	parsing
Compilers	Jones	McGraw-Hill	New York	parsing
Compilers	Smith	McGraw-Hill	New York	analysis
Compilers	Jones	McGraw-Hill	New York	analysis
Networks	Jones	Oxford	London	Internet
Networks	Frick	Oxford	London	Internet
Networks	Jones	Oxford	London	Web
Networks	Frick	Oxford	London	Web



<i>title</i>	<i>author-set</i>	<i>publisher</i>	<i>keyword-set</i>
		(<i>name, branch</i>)	
Compilers	{Smith, Jones}	(McGraw-Hill, New York)	{parsing, analysis}
Networks	{Jones, Frick}	(Oxford, London)	{Internet, Web}

** note: group by title, publisher



Nesting [3/3]

- Another approach to creating nested relations is to use subqueries in the **select** clause, starting from the 4NF relation *books4*

```
select title,  
       array (select author  
              from authors as A  
              where A.title = B.title  
              order by A.position) as author_array,  
       Publisher (pub-name, pub-branch) as publisher,  
       multiset (select keyword  
                  from keywords as K  
                  where K.title = B.title) as keyword_set  
from books4 as B
```



Nesting & Unnesting

- Unnesting

```
select title, A as author,  
        publisher.name as pub_name,  
        publisher.branch as pub_branch,  
        K as keyword  
from books as B,  
        unnest(B.author_array) as A,  
        unnest(B.keyword_set) as K
```

- result is *flat_books*

- Nesting

```
select title, author,  
        Publisher(pubname, pubbranch)  
        as publisher, collectkeyword() as  
        keyword_list  
from flat_books  
group by title, author, publisher
```

- result is shown below

<i>title</i>	<i>author</i>	<i>publisher</i>	<i>keyword-set</i>
		(pub-name, pub-branch)	
Compilers	Smith	(McGraw-Hill, New York)	{parsing, analysis}
Compilers	Jones	(McGraw-Hill, New York)	{parsing, analysis}
Networks	Jones	(Oxford, London)	{Internet, Web}
Networks	Frick	(Oxford, London)	{Internet, Web}



Chapter 22: Object-Based Databases

- 22.1 Overview
- 22.2 Complex Data Types
- 22.3 Structured Data Types and Inheritance in SQL
- 22.4 Table Inheritance
- 22.5 Array and Multiset Types in SQL
- 22.6 Object Identity and Reference Types in SQL
- 22.7 Implementing O-R Features
- 22.8 Persistent Programming Languages
- 22.9 Object-Relational Mapping
- 22.10 Object-Oriented versus Object-Relational



Object-Identity and Reference Types

- Define a type *Department* with a field *name* and a field *head* which is a reference to the type *Person*, with table *people* as scope:

```
create type Department (  
    name varchar (20),  
    head ref (Person) scope people)
```

- We can then create a table *departments* as follows

```
create table departments of Department
```

- We can omit the declaration **scope** *people* from the type declaration and instead make an addition to the **create table** statement:

```
create table departments of Department  
    (head with options scope people)
```

- Referenced table must have an attribute that stores the identifier, called the **self-referential attribute**

```
create table people of Person  
ref is person_id system generated;
```




Initializing Reference-Typed Values

- To create a tuple with a reference value, we can first create the tuple with a null reference and then set the reference separately:

```
insert into departments
```

```
  values ('CS', null)
```

```
update departments
```

```
  set head = (select p.person_id  
                from people as p  
                where name = 'John')
```

```
where name = 'CS'
```



User Generated Identifiers [1/2]

- The type of the object-identifier must be specified as part of the type definition of the referenced table, and
- The table definition must specify that the reference is user generated

```
create type Person  
  (name varchar(20)  
   address varchar(20))  
  ref using varchar(20)  
create table people of Person  
  ref is person_id user generated
```

- When creating a tuple, we must provide a unique value for the identifier:

```
insert into people (person_id, name, address) values  
  ('01284567', 'John', '23 Coyote Run')
```

- We can then use the identifier value when inserting a tuple into *departments*
 - Avoids need for a separate query to retrieve the identifier:

```
insert into departments  
values('CS', '02184567')
```



User Generated Identifiers [2/2]

- Can use an existing primary key value as the identifier:

```
create type Person  
  (name varchar (20) primary key,  
   address varchar(20))  
ref from (name)  
create table people of Person  
  ref is person_id derived
```

- When inserting a tuple for *departments*, we can then use

```
insert into departments  
  values(`CS`,`John`)
```



Path Expressions [1/2]

```
create type Department (  
    name varchar (20),  
    head ref (Person) scope people)
```

```
create type Person  
    ( name          varchar(20)  
      address       varchar(20) )  
    ref using      varchar(20)
```

- Find the names and addresses of the heads of all departments:

```
select head → name, head → address  
from departments
```

- An expression such as “*head* → *name*” is called a **path expression**
- Path expressions help avoid explicit joins
 - If department head were not a reference, a join of *departments* with *people* would be required to get at the address
 - Makes expressing the query much easier for the user



Path Expressions [2/2]

- Dot (.) notation is used for **composite attributes**

```
select title, publisher.name  
from books
```

- Pointer (→) notation is used for **reference attributes**

```
select head→name, head→address  
from departments
```



Chapter 22: Object-Based Databases

- 22.1 Overview
- 22.2 Complex Data Types
- 22.3 Structured Data Types and Inheritance in SQL
- 22.4 Table Inheritance
- 22.5 Array and Multiset Types in SQL
- 22.6 Object Identity and Reference Types in SQL
- 22.7 Implementing O-R Features
- 22.8 Persistent Programming Languages
- 22.9 Object-Relational Mapping
- 22.10 Object-Oriented versus Object-Relational



Implementing O-R Features in RDB

- If we want to keep existing RDBMS and utilize O-R advantages
 - Structured Type, Array, Multiset, Nested relations, Inheritance, Subtable
- Convert tables with O-R tables into Relational Tables
 - Similar to how E-R features are mapped onto relation schemas
 - Multivalued attribute vs Multi-Set valued attribute
 - Composite attribute vs Structured Type
 - ISA vs Table Inheritance
- Subtable implementation
 - Each table stores primary key and those attributes locally defined in that table or,
 - Each table stores both locally defined and inherited attributes



Chapter 22: Object-Based Databases

- 22.1 Overview
- 22.2 Complex Data Types
- 22.3 Structured Data Types and Inheritance in SQL
- 22.4 Table Inheritance
- 22.5 Array and Multiset Types in SQL
- 22.6 Object Identity and Reference Types in SQL
- 22.7 Implementing O-R Features
- 22.8 Persistent Programming Languages
- 22.9 Object-Relational Mapping
- 22.10 Object-Oriented versus Object-Relational



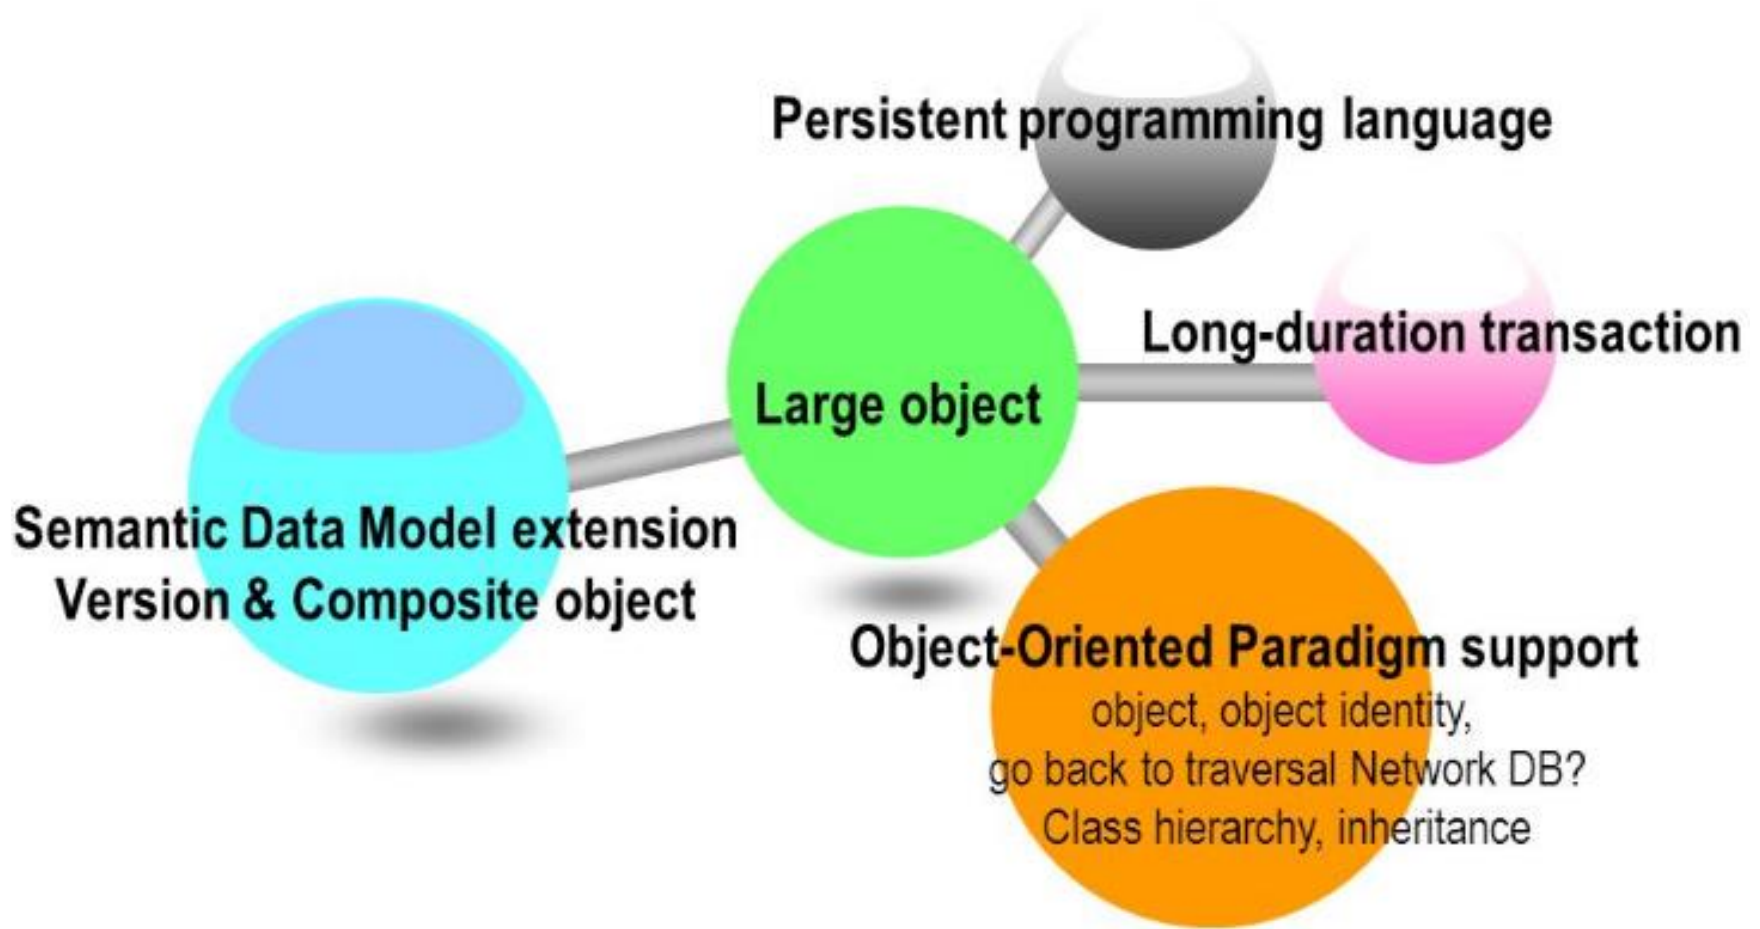
Persistent OO Programming Languages

- Languages extended with **constructs** to handle persistent data
- Programmer can manipulate persistent data directly
 - no need to fetch it into memory and store it back to disk (unlike embedded SQL)
- Supporting Persistent Objects inside Programming Language!

- Persistent objects:
 - **Persistence by class** - explicit declaration of persistence
 - **Persistence by creation** - special syntax to create persistent objects
 - **Persistence by marking** - make objects persistent after creation
 - **Persistence by reachability** - object is persistent if it is declared explicitly to be so or is reachable from a persistent object



Feature of OODBMS





Concerns in Persistent PL

- Object Identifiers
 - We need stronger version of in-memory pointers in Persistent PL
 - Degrees of permanence of object identity
 - ▶ **Intraprocedure**: only during execution of a single procedure
 - ▶ **Intraprogram**: only during execution of a single program or query
 - ▶ **Interprogram**: across program executions, but not if data-storage format on disk changes
 - ▶ **Persistent**: interprogram, plus persistent across data reorganizations
- How to represent class and its instances
- How to support Query
- How to support Transaction



Object Identity and Pointers

- Degrees of permanence of object identity
 - **Intraprocedure**: only during execution of a single procedure
 - **Intraprogram**: only during execution of a single program or query
 - **Interprogram**: across program executions, but not if data-storage format on disk changes
 - **Persistent**: interprogram, plus persistent across data reorganizations
- Persistent versions of C++ and Java have been implemented
 - C++
 - ▶ ODMG C++
 - ▶ ObjectStore
 - Java
 - ▶ Java Database Objects (JDO)



Persistent C++ Systems

- Extensions of C++ language to support persistent storage of objects
- Several proposals, ODMG standard proposed, but not much action of late
 - **persistent pointers:** e.g. `d_Ref<T>`
 - **creation of persistent objects:** e.g. `new (db) T()`
 - **Class extents:** access to all persistent objects of a particular class
 - **Relationships:** Represented by pointers stored in related objects
 - ▶ Issue: consistency of pointers
 - ▶ Solution: extension to type system to automatically maintain back-references
 - **Iterator interface**
 - **Transactions**
 - **Updates:** `mark_modified()` function to tell system that a persistent object that was fetched into memory has been updated
 - **Query language**



Persistent Java Systems

- Standard for adding persistence to Java : **Java Database Objects (JDO)**
 - Persistence by reachability
 - Byte code enhancement
 - ▶ Classes separately declared as persistent
 - ▶ Byte code modifier program modifies class byte code to support persistence
 - E.g. Fetch object on demand
 - Mark modified objects to be written back to database
 - Database mapping
 - ▶ Allows objects to be stored in a relational database
 - Class extents
 - Single reference type
 - ▶ no difference between in-memory pointer and persistent pointer
 - ▶ Implementation technique based on **hollow objects** (a.k.a. **pointer swizzling**)



Chapter 22: Object-Based Databases

- 22.1 Overview
- 22.2 Complex Data Types
- 22.3 Structured Data Types and Inheritance in SQL
- 22.4 Table Inheritance
- 22.5 Array and Multiset Types in SQL
- 22.6 Object Identity and Reference Types in SQL
- 22.7 Implementing O-R Features
- 22.8 Persistent Programming Languages
- 22.9 Object-Relational Mapping
- 22.10 Object-Oriented versus Object-Relational



Object-Relational Mapping

- **Object-Relational Mapping (ORM)** systems built on top of traditional relational databases
- Implementor provides a mapping from objects to relations
 - Objects are purely transient, no permanent object identity
- Objects can be retrieved from database
 - System uses mapping to fetch relevant data from relations and construct objects
 - Updated objects are stored back in database by generating corresponding update/insert/delete statements
- The **Hibernate** ORM system is widely used
 - described in Section 9.4.2
 - Provides API to start/end transactions, fetch objects, etc
 - Provides query language operating directly on object model
 - ▶ queries translated to SQL
- Limitations: overheads, especially for bulk updates



Chapter 22: Object-Based Databases

- 22.1 Overview
- 22.2 Complex Data Types
- 22.3 Structured Data Types and Inheritance in SQL
- 22.4 Table Inheritance
- 22.5 Array and Multiset Types in SQL
- 22.6 Object Identity and Reference Types in SQL
- 22.7 Implementing O-R Features
- 22.8 Persistent Programming Languages
- 22.9 Object-Relational Mapping
- 22.10 Object-Oriented versus Object-Relational



Comparison of O-O and O-R Databases

- **Relational systems**
 - simple data types, powerful query languages, high protection
- **Persistent-programming-language-based OODBs**
 - complex data types, integration with programming language, high performance
- **Object-relational systems**
 - complex data types, powerful query languages, high protection.
- **Object-relational mapping systems**
 - complex data types integrated with programming language, but built as a layer on top of a relational database system
- **Note: Many real systems blur these boundaries**
 - E.g. persistent programming language built as a wrapper on a relational database offers first two benefits, but may have poor performance.