

Matrix Factorization and Collaborative Filtering



Hyewon Lim

4 Jan 2017

Outline

- **Recommender System**
- Matrix Factorization
- Reference



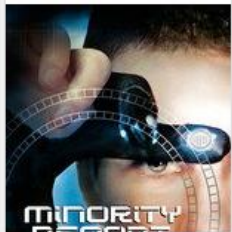









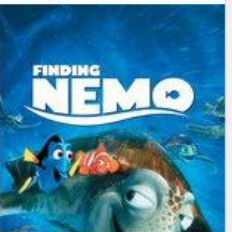




























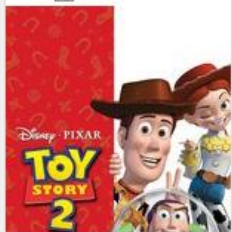


































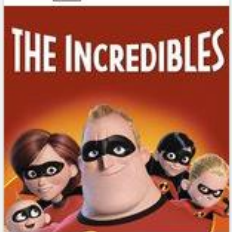

Recommender System

movielens  

MovieLens recommends these movies

top picks

found 48995 movies. sort by:

<p>Minority Report  </p> <p>2002 • 145 min</p>  <p></p>	<p>Beauty and the Beast  </p> <p>1991  84 min</p>  <p></p>	<p>Finding Nemo  </p> <p>2003  100 min</p>  <p></p>	<p>Edge of Tomorrow  </p> <p>2014  113 min</p>  <p></p>	<p>Iron Man  </p> <p>2008  126 min</p>  <p></p>	<p>The Theory of Everything  </p> <p>2014  123 min</p>  <p></p>	<p>The Man from U.N.C.L.E.  </p> <p>2015  116 min</p>  <p></p>	<p>Big Hero 6  </p> <p>2014 • 102 min</p>  <p></p>
<p>Toy Story 2  </p> <p>1999  92 min</p>  <p></p>	<p>District 9  </p> <p>2009  112 min</p>  <p></p>	<p>21 Jump Street  </p> <p>2012  109 min</p>  <p></p>	<p>John Wick  </p> <p>2014  101 min</p>  <p></p>	<p>Star Trek  </p> <p>2009  127 min</p>  <p></p>	<p>Watchmen  </p> <p>2009  162 min</p>  <p></p>	<p>The Hangover  </p> <p>2009  100 min</p>  <p></p>	<p>The Incredibles  </p> <p>2004  115 min</p>  <p></p>

Recommender System Strategies

1. Content Filtering

- Create a profile for each user or product to characterize its nature



Genres: Crime, Comedy, Action, Adventure

Directors: Matthew Vaughn

Cast: Taron Egerton, Colin Firth, Samuel L. Jackson, ...

Distributor: Fox

Box Office Popularity: ...



Gender

Region


...

Answers provided on a suitable questionnaire

Recommender System Strategies

2. Collaborative filtering

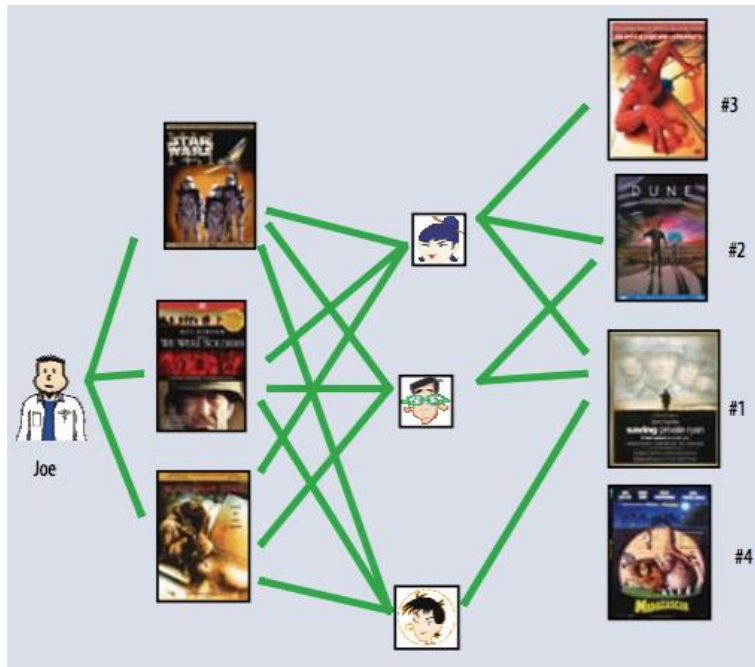
- Rely only on past user behavior
- Everyday examples
 - Bestseller lists
 - Top 40 music lists
 - Unmarked but well-used paths thru the woods
 - The “recent returns” shelf at the library
- Common insight: **personal tastes are correlated**



Types of Collaborative Filtering

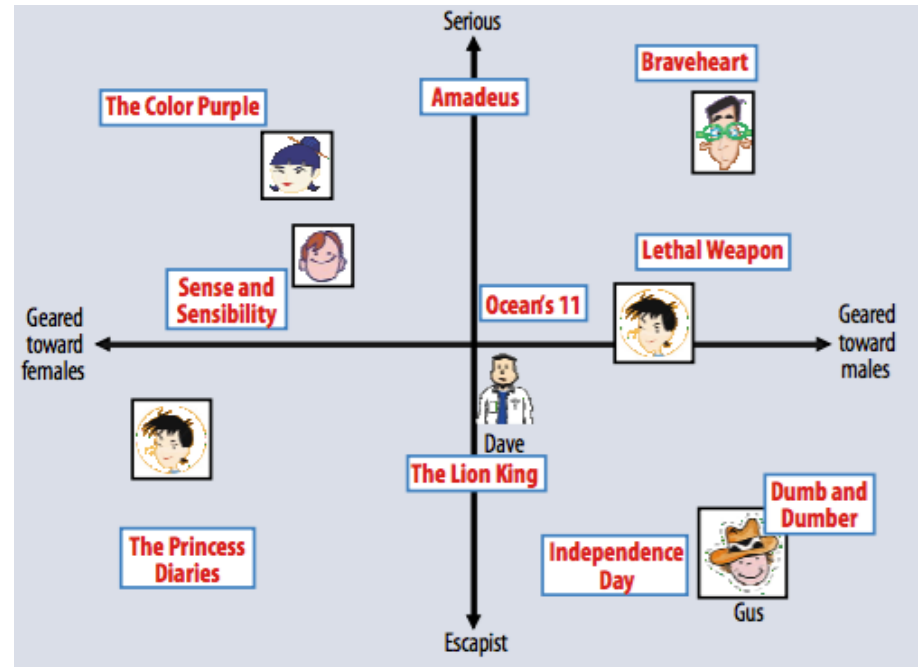
a. Neighborhood Methods

- Find neighbors based on similarity of movie preferences
- Recommend movies that those neighbors watched



b. Latent Factor Methods


- Characterize both items and users
- Recommend a movie based on its proximity to the user in the latent space



Outline

- Recommender System
- **Matrix Factorization**
- Reference

Netflix Prize



Netflix Prize

COMPLETED

[Home](#) [Rules](#) [Leaderboard](#) [Update](#)

Leaderboard

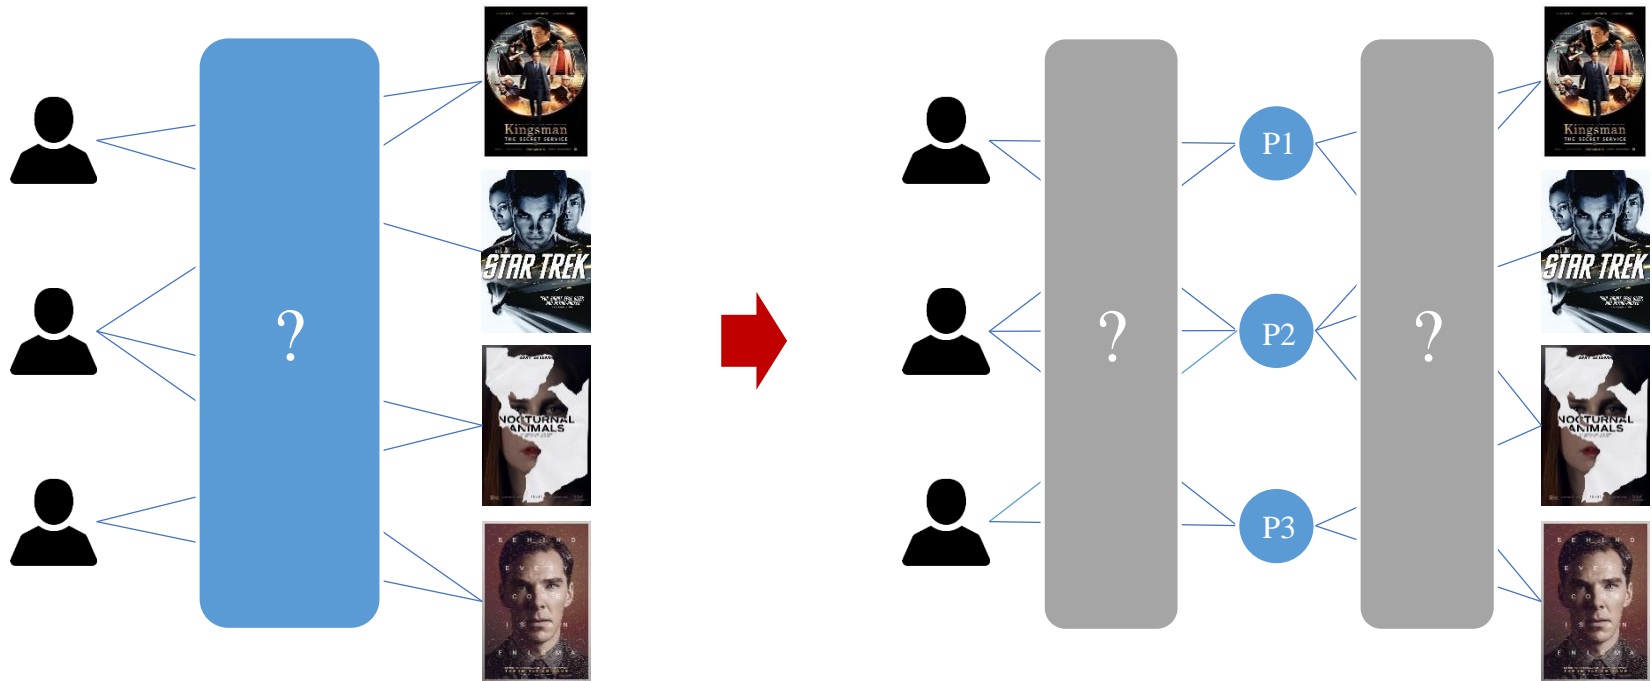
Showing Test Score. [Click here to show quiz score](#)

500,000 users
20,000 movies
100M ratings

Rank	Team Name	Best Test Score	% Improvement	Best Submit Time
Grand Prize - RMSE = 0.8567 - Winning Team: BellKor's Pragmatic Chaos				
1	BellKor's Pragmatic Chaos	0.8567	10.06	2009-07-26 18:18:28
2	The Ensemble	0.8567	10.06	2009-07-26 18:38:22
3	Grand Prize Team	0.8582	9.90	2009-07-10 21:24:40
4	Opera Solutions and Vandelay United	0.8588	9.84	2009-07-10 01:12:31
5	Vandelay Industries I	0.8591	9.81	2009-07-10 00:32:20
6	PragmaticTheory	0.8594	9.77	2009-06-24 12:06:56
7	BellKor in BigChaos	0.8601	9.70	2009-05-13 08:14:09
8	Dace	0.8612	9.59	2009-07-24 17:18:43
9	Feeds2	0.8622	9.48	2009-07-12 13:11:51
10	BigChaos	0.8623	9.47	2009-04-07 12:33:59
11	Opera Solutions	0.8623	9.47	2009-07-24 00:34:07
12	BellKor	0.8624	9.46	2009-07-26 17:19:11

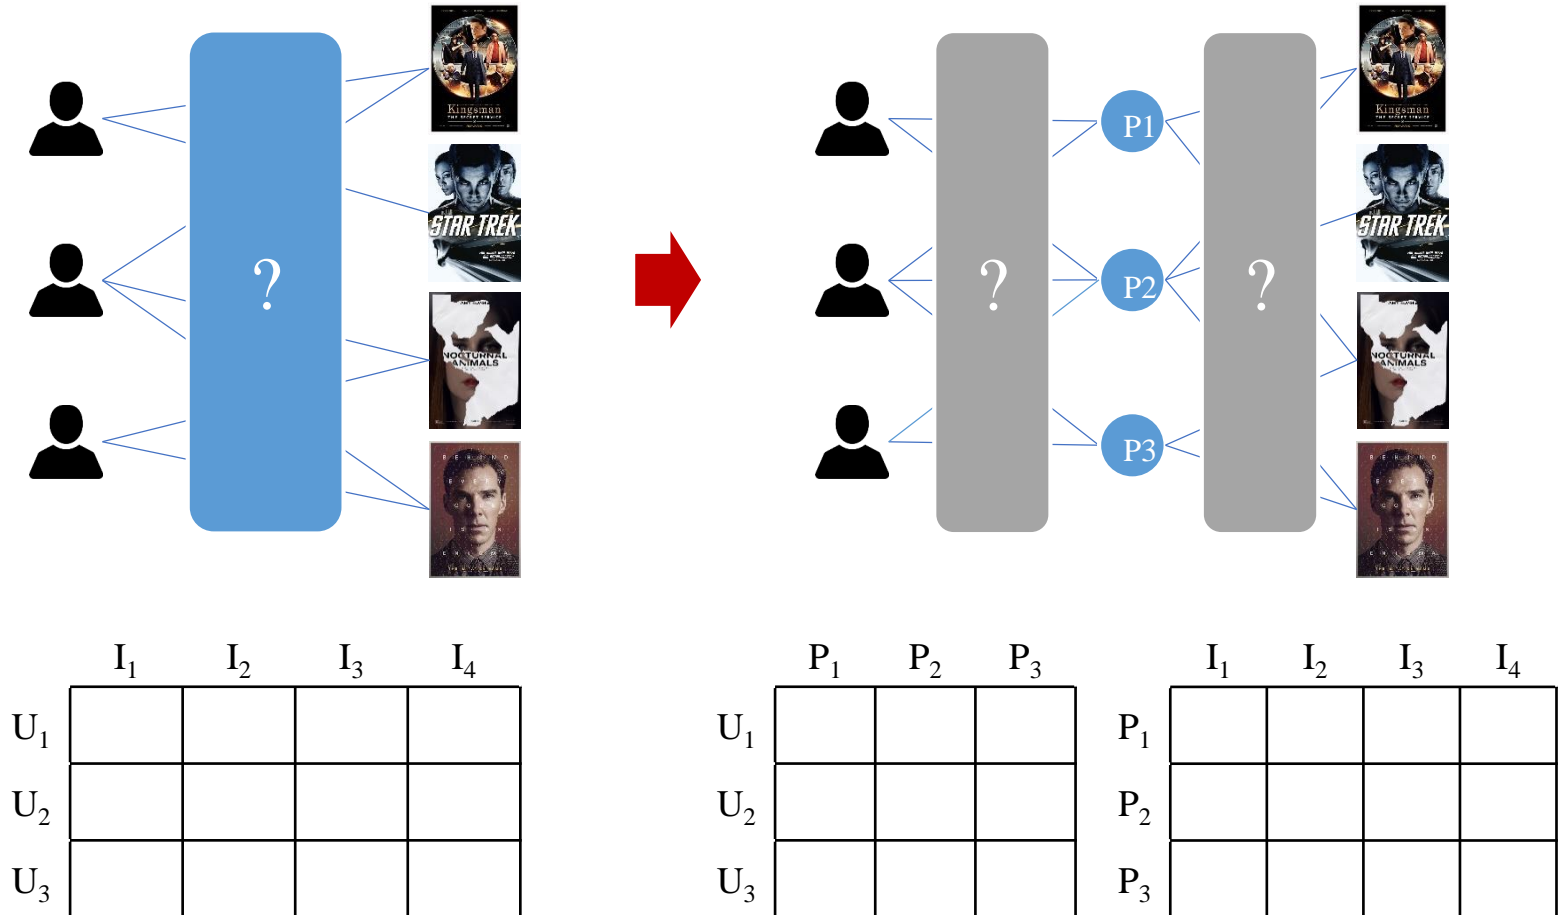
Matrix Factorization

- Assume latent factors in user preference



Matrix Factorization

- Assume latent factors in user preference



* Redraw figures in [2]

Singular Value Decomposition

$$A = U \times S \times V^T$$

$$\begin{bmatrix} 4. & 2. & 3. & 5. & 1. \\ 0. & 3. & 0. & 4. & 2. \\ 5. & 4. & 3. & 3. & 0. \\ 0. & 0. & 5. & 5. & 2. \\ 5. & 0. & 0. & 5. & 0. \end{bmatrix} =$$

$$\begin{bmatrix} -0.54 & 0.03 & -0.021 & 0.099 & -0.835 \\ -0.29 & -0.225 & 0.393 & -0.84 & 0.07 \\ -0.506 & 0.372 & 0.574 & 0.374 & 0.371 \\ -0.417 & -0.79 & -0.217 & 0.277 & 0.279 \\ -0.441 & 0.432 & -0.685 & -0.26 & 0.287 \end{bmatrix} \times \begin{bmatrix} 13.707 & 0. & 0. & 0. & 0. \\ 0. & 5.607 & 0. & 0. & 0. \\ 0. & 0. & 3.791 & 0. & 0. \\ 0. & 0. & 0. & 3.645 & 0. \\ 0. & 0. & 0. & 0. & 0.155 \end{bmatrix} \times \begin{bmatrix} -0.503 & -0.29 & -0.381 & -0.705 & -0.143 \\ 0.738 & 0.156 & -0.489 & -0.254 & -0.357 \\ -0.169 & 0.905 & 0.15 & -0.35 & 0.087 \\ 0.265 & -0.227 & 0.769 & -0.455 & -0.282 \\ -0.321 & 0.147 & 0.029 & 0.33 & -0.875 \end{bmatrix}$$



We can drop less important information

Singular Value Decomposition

$$A \approx U \times S \times V^T$$

$$\begin{bmatrix} -0.54 & 0.03 & -0.021 & 0.099 & -0.835 \\ -0.29 & -0.225 & 0.393 & -0.84 & 0.07 \\ -0.506 & 0.372 & 0.574 & 0.374 & 0.371 \\ -0.417 & -0.79 & -0.217 & 0.277 & 0.279 \\ -0.441 & 0.432 & -0.685 & -0.26 & 0.287 \end{bmatrix} \times \begin{bmatrix} 13.707 & 0. & 0. & 0. & 0. \\ 0. & 5.607 & 0. & 0. & 0. \\ 0. & 0. & 3.791 & 0. & 0. \\ 0. & 0. & 0. & 3.645 & 0. \\ 0. & 0. & 0. & 0. & 0.155 \end{bmatrix} \times \begin{bmatrix} -0.503 & -0.29 & -0.381 & -0.705 & -0.143 \\ 0.738 & 0.156 & -0.489 & -0.254 & -0.357 \\ -0.169 & 0.905 & 0.15 & -0.35 & 0.087 \\ 0.265 & -0.227 & 0.769 & -0.455 & -0.282 \\ -0.321 & 0.147 & 0.029 & 0.33 & -0.875 \end{bmatrix}$$




$$\begin{bmatrix} -0.54 & 0.03 & -0.021 & 0.099 & -0.835 \\ -0.29 & -0.225 & 0.393 & -0.84 & 0.07 \\ -0.506 & 0.372 & 0.574 & 0.374 & 0.371 \\ -0.417 & -0.79 & -0.217 & 0.277 & 0.279 \\ -0.441 & 0.432 & -0.685 & -0.26 & 0.287 \end{bmatrix} \times \begin{bmatrix} 13.707 & 0. & 0. & 0. & 0. \\ 0. & 5.607 & 0. & 0. & 0. \\ 0. & 0. & 3.791 & 0. & 0. \\ 0. & 0. & 0. & 3.645 & 0. \\ 0. & 0. & 0. & 0. & 0.155 \end{bmatrix} \times \begin{bmatrix} -0.503 & -0.29 & -0.381 & -0.705 & -0.143 \\ 0.738 & 0.156 & -0.489 & -0.254 & -0.357 \\ -0.169 & 0.905 & 0.15 & -0.35 & 0.087 \\ 0.265 & -0.227 & 0.769 & -0.455 & -0.282 \\ -0.321 & 0.147 & 0.029 & 0.33 & -0.875 \end{bmatrix}$$

Matrix Factorization

■ Matrices

- User vector
 - $(P_{u*})^T \in \mathbb{R}^f$
- Item vectors:
 - $(Q_{*i}) \in \mathbb{R}^f$
- Rating prediction
 - $R_{ui} = P_{u*} Q_{*i} = [PQ]_{ui}$



■ Vectors

- User vector
 - $p_u \in \mathbb{R}^r$
- Item vectors:
 - $q_i \in \mathbb{R}^r$
- Rating prediction
 - $\hat{r}_{ui} = q_i^T p_u$
- Set of non-zero entries
 - $\kappa = \{(u, i): r_{ui} \neq 0\}$
- Objective
 - $\min_{q^*, p^*} \sum_{(u, i) \in \kappa} (r_{ui} - q_i^T p_u)^2$

Matrix Factorization

$$R_{ui} \approx \hat{R}_{ui} = U \times I$$

- Minimize the error between R and \hat{R}

$$\min_{q^*, p^*} \sum_{(u,i) \in \kappa} (r_{ui} - q_i^T p_u)^2$$

$$\min_{q^*, p^*} \sum_{(u,i) \in \kappa} (r_{ui} - q_i^T p_u)^2 + \lambda(\|q_i\|^2 + \|p_u\|^2)$$

Regularization factor
- avoid overfitting
- make simple model

Matrix Factorization

- How to deal with empty cells in matrix

	I ₁	I ₂	I ₃	I ₄
U ₁		3	4	2
U ₂	5			
U ₃	3		2	

$$b_{ui} = \mu + b_i + b_u$$

μ : average of the whole users

b_i, b_u : the observed deviations of u and i



$$\hat{r}_{ui} = \mu + b_i + b_u + q_i^T p_u$$

$$\min_{q^*, p^*, b^*} \sum_{(u,i) \in \kappa} (r_{ui} - \mu - b_i - b_u - p_u^T q_i)^2 + \lambda(\|q_i\|^2 + \|p_u\|^2 + b_u^2 + b_i^2)$$

Matrix Factorization

- Approaches to minimizing $\min_{q^*, p^*} \sum_{(u,i) \in \kappa} (r_{ui} - q_i^T p_u)^2 + \lambda(\|q_i\|^2 + \|p_u\|^2)$

1. Stochastic gradient descent



- Associated prediction error e_{ui}
 - $q_i \leftarrow q_i + \gamma(e_{ui} \cdot p_u - \lambda \cdot q_i)$
 - $p_u \leftarrow p_u + \gamma(e_{ui} \cdot q_i - \lambda \cdot p_u)$

Matrix Factorization

- Approaches to minimizing $\min_{q^*, p^*} \sum_{(u,i) \in \kappa} (r_{ui} - q_i^T p_u)^2 + \lambda(\|q_i\|^2 + \|p_u\|^2)$
 - 2. Alternating least squares
 - Rotate between fixing the q_i 's and fixing the p_u 's
 - When all p_u 's are fixed, the system recomputes the q_i 's by solving a least-squares problems, and vice versa
 - Stochastic gradient descent is easier and faster than ALS in general, ALS is favorable in at least two cases
 - When the system can use parallelization
 - For systems centered on implicit data

HOSVD [4]

- SVD on each matrix

$$A_1 = U^{(1)} \cdot S_1 \cdot V_1^T$$

$$A_2 = U^{(2)} \cdot S_2 \cdot V_2^T$$



$$A_3 = U^{(3)} \cdot S_3 \cdot V_3^T$$

- Construction of core tensor

$$\mathcal{S} = \mathcal{A} \times_1 U_{c_1}^{(1)T} \times_2 U_{c_2}^{(2)T} \times_3 U_{c_3}^{(3)T}$$

- Construction of tensor $\hat{\mathcal{A}}$

$$\hat{\mathcal{A}} = \mathcal{S} \times_1 U_{c_1}^{(1)} \times_2 U_{c_2}^{(2)} \times_3 U_{c_3}^{(3)}$$



Example

■ MF in Python

```
1 #!/usr/bin/python
2
3 import numpy as np
4
5 np.set_printoptions(precision = 3) # set decimal display
6
7 # Matrix
8 A = np.zeros((5, 5))
9
10 A[0, 4] = 1
11 A[0, 1] = A[1, 4] = A[3, 4] = 2
12 A[0, 2] = A[1, 1] = A[2, 2] = A[2, 3] = 3
13 A[0, 0] = A[1, 3] = A[2, 1] = 4
14 A[0, 3] = A[2, 0] = A[3, 2] = A[3, 3] = A[4, 0] = A[4, 3] = 5
15
16 # SVD
17 U, s, V = np.linalg.svd(A, full_matrices = True)
18
19 # Reconstruction
20 S = np.diag(s)
21
22 P = np.dot(U, np.dot(S, V))
23
```

```
[ 4.  2.  3.  5.  1.]
[ 0.  3.  0.  4.  2.]
[ 5.  4.  3.  3.  0.]
[ 0.  0.  5.  5.  2.]
[ 5.  0.  0.  5.  0.]
```

```
[ 4e+00  2e+00  3e+00  5e+00  1e+00]
[-4e-16  3e+00  3e-15  4e+00  2e+00]
[ 5e+00  4e+00  3e+00  3e+00  7e-16]
[-4e-15  4e-15  5e+00  5e+00  2e+00]
[ 5e+00  1e-15 -2e-16  5e+00 -5e-16]
```

Example

■ MF in Python with r

```
1 #!/usr/bin/python
2
3 import rpy2.robjects as robjects
4
5 r = robjects.r
6
7 r('''
8     rsvd <- function() {
9         # MATRIX
10        A <- matrix(c(4, 2, 3, 5, 1, 0, 3, 0, 4, 2, 5, 4, 3
11          , 3, 0, 0, 0, 5, 5, 2, 5, 0, 0, 5, 0), nrow = 5, ncol = 5,
12          byrow = TRUE)
13
14        # SVD
15        result <- svd(A)
16
17        # RECONSTRUCTION
18        U <- result$u
19        s <- result$d
20        V <- result$v
21
22        ApproxA <- U %*% diag(s) %*% t(V)
23    }
24    ''')
25
26 svd = r['rsvd']
27
28 result = svd()
29
30 print result
```

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	4	2	3	5	1
[2,]	0	3	0	4	2
[3,]	5	4	3	3	0
[4,]	0	0	5	5	2
[5,]	5	0	0	5	0

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	4.000000e+00	2.000000e+00	3.000000e+00	5	1.000000e+00
[2,]	-1.955901e-16	3.000000e+00	5.999975e-16	4	2.000000e+00
[3,]	5.000000e+00	4.000000e+00	3.000000e+00	3	-1.110223e-16
[4,]	-6.366435e-16	3.387048e-15	5.000000e+00	5	2.000000e+00
[5,]	5.000000e+00	1.949829e-15	-6.570265e-16	5	1.408595e-15

Reference

1. Slides in “Matrix Factorization and Collaborative Filtering”
 - By Matt Gormley (Carnegie Mellon Univ.)
2. Slides in “Recommender Systems”
 - By Jee-Hyong Lee (Sungkyunkwan Univ.)
3. Y. Koren *et al.*, “Matrix Factorization Techniques for Recommender Systems,” *Journal Computer*, 42(8), 2009
4. P. Symeonidis *et.al*, “Tag Recommendations based on Tensor Dimensionality Reduction,” *Recsys’08*