



# Chapter 14: Transactions

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Transaction Concept

- A **transaction** is a *unit* of program execution that accesses and possibly updates various data items
- E.g. transaction to transfer \$50 from account A to account B:
  1. **read**(A)
  2.  $A := A - 50$
  3. **write**(A)
  4. **read**(B)
  5.  $B := B + 50$
  6. **write**(B)
- Two main issues to deal with:
  - **Failures of various kinds**, such as hardware failures and system crashes
  - **Concurrent execution** of multiple transactions



# ACID Properties


To preserve the integrity of data the database system must ensure:

- **Atomicity.** Either **all** operations of the transaction are properly reflected in the database or **none** are.
- **Consistency.** Execution of a transaction in isolation preserves the consistency of the database.
- **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions.
- **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.



# Atomicity

- “All or nothing”
- The system should ensure that updates of a partially executed transaction are not reflected in the database
- Transaction to transfer \$50 from account A to account B:
  1. **read**(A)
  2.  $A := A - 50$
  3. **write**(A)
  4. **read**(B) ← System crash
  5.  $B := B + 50$
  6. **write**(B)
  - Money will be “lost” leading to an inconsistent database state





# Consistency

- When the transaction completes successfully, the database must be consistent
  - During transaction execution, the database may be temporarily inconsistent
  - A transaction must see a consistent database
  
- Transaction to transfer \$50 from account A to account B:
  1. **read**(A)
  2.  $A := A - 50$
  3. **write**(A)
  4. **read**(B)
  5.  $B := B + 50$
  6. **write**(B)
  - The sum of A and B is unchanged by the execution of the transaction
  
- In general, consistency requirements include
  - Explicitly specified integrity constraints such as primary keys and foreign keys
  - Implicit integrity constraints
    - ▶ e.g. sum of balances of all accounts, minus sum of loan amounts



# Isolation

- Each transaction must be unaware of other concurrently executing transactions
  - Intermediate transaction results must be hidden from other concurrently executed transactions
  - Isolation can be ensured trivially by running transactions **serially**
    - ▶ That is, one after the other.

- Transaction to transfer \$50 from account A to account B:

**T1**

1. **read**(A)
2.  $A := A - 50$
3. **write**(A)
4. **read**(B)
5.  $B := B + 50$
6. **write**(B)

**T2**

read(A), read(B), print(A+B)

If T2 accesses the partially updated database, it will see an inconsistent database (the sum  $A + B$  will be less than it should be).

- However, executing multiple transactions concurrently has significant benefits, as we will see later.



# Durability

- The updates to the database by the transaction must persist even if there are software or hardware failures
  
- Transaction to transfer \$50 from account A to account B:
  1. **read**(A)
  2.  $A := A - 50$
  3. **write**(A)
  4. **read**(B)
  5.  $B := B + 50$
  6. **write**(B)
  - Once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), it must persist



# Transaction State

- **Active** – the initial state; the transaction stays in this state while it is executing
- **Partially committed** – after the final statement has been executed
- **Failed** – after the discovery that normal execution can no longer proceed
- **Aborted** – after the transaction has been rolled back and the database restored to its state prior to the start of the transaction


Two options after it has been aborted:

- restart the transaction
    - ▶ can be done only if no internal logical error
  - kill the transaction
- **Committed** – after successful completion





# Transaction State (Cont.)





# Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system.  
Advantages are:
  - **Increased processor and disk utilization**, leading to better transaction *throughput*
  - **Reduced average response time** for transactions: short transactions need not wait behind long ones
- **Concurrency control schemes** – mechanisms to achieve isolation
  - To control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database



# Schedules

- **Schedule** – a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
  - a schedule for a set of transactions must consist of all instructions of those transactions
  - must preserve the order in which the instructions appear in each individual transaction
- **Serial schedule** – instruction sequences from one by one transactions
- Simplified view of transactions
  - Our simplified schedules consist of only **read** and **write** instructions
  - We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes



# Schedule 1

- Let  $T_1$  transfer \$50 from  $A$  to  $B$ , and  $T_2$  transfer 10% of the balance from  $A$  to  $B$
- A serial schedule in which  $T_1$  is followed by  $T_2$  :

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit



## Schedule 2

- A serial schedule in which  $T_2$  is followed by  $T_1$  :

$T_1$	$T_2$
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit



# Schedule 3

- The following schedule is not a serial schedule, but it is *equivalent* to Schedule 1
  - We call it a *serializable schedule*

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ )	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ )
read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

In Schedules 1, 2 and 3, the sum  $A + B$  is preserved.



# Schedule 4

- The following concurrent schedule does not preserve the value of  $(A + B)$ . The following schedule is **not serializable**.

$T_1$	$T_2$
read (A) $A := A - 50$	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
write (A) read (B) $B := B + 50$ write (B) commit	         $B := B + temp$ write (B) commit



# Serializability

- **Basic Assumption** – Each transaction preserves database consistency
- Thus serial execution of a set of transactions preserves database consistency
- A (possibly concurrent) schedule is **serializable** if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notions of:
  1. **conflict serializability**
  2. **view serializability**





# Conflicting Instructions

- Instructions  $I_i$  and  $I_j$  of transactions  $T_i$  and  $T_j$  respectively, **conflict** if and only if
  - There exists some item  $Q$  accessed by both  $I_i$  and  $I_j$ ,
  - and at least one of these instructions wrote  $Q$ .
    1.  $I_i = \text{read}(Q)$ ,  $I_j = \text{read}(Q)$ .  $I_i$  and  $I_j$  don't conflict.
    2.  $I_i = \text{read}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict.
    3.  $I_i = \text{write}(Q)$ ,  $I_j = \text{read}(Q)$ . They conflict
    4.  $I_i = \text{write}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict
- Intuitively, a conflict between  $I_i$  and  $I_j$  forces a (logical) temporal order between them
  - If  $I_i$  and  $I_j$  do not conflict, their results would remain the same even if they had been interchanged in the schedule



# Conflict Serializability

- Schedules  $S$  and  $S'$  are **conflict equivalent** if  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions
- A schedule  $S$  is **conflict serializable** if It is conflict equivalent to a serial schedule
- Example of a schedule that is **not conflict serializable**:

$T_3$	$T_4$
read(Q)	write(Q)
write(Q)	

We are unable to swap instructions in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$ .

$T_3$	$T_4$
read(Q)	
write(Q)	
	write(Q)

$T_3$	$T_4$
	write(Q)
read(Q)	
write(Q)	



# Conflict Serializability (Cont.)

- Schedule 3 can be transformed into Schedule 6, a serial schedule where  $T_2$  follows  $T_1$ , by series of **swaps of non-conflicting instructions**. Therefore Schedule 3 is conflict serializable.

**Schedule 3**

$T_1$	$T_2$
read(A)	
write(A)	
	read(A)
	write(A)
read(B)	
write(B)	
	read(B)
	write(B)

**Schedule 5 –**  
After Swapping a Pair of non  
conflicting Instructions in  
schedule 3

$T_1$	$T_2$
read(A)	
write(A)	
<u>read(B)</u>	read(A)
	<u>write(A)</u>
write(B)	
	read(B)
	write(B)

**Schedule 6 –**  
A Serial Schedule That is  
Equivalent to Schedule 3

$T_1$	$T_2$
read(A)	
write(A)	
read(B)	
write(B)	
	read(A)
	write(A)
	read(B)
	write(B)




# View Serializability

- Schedules  $S$  and  $S'$  are **view equivalent** if the following three conditions are met, for each data item  $Q$ :
  1. If in  $S$ , transaction  $T_i$  reads the initial value of  $Q$ , then in  $S'$  also transaction  $T_i$  must read the initial value of  $Q$ .
  2. If in  $S$ ,  $T_i$  executes **read**( $Q$ ), and that value was produced by  $T_j$  (if any), then in  $S'$  also  $T_i$  must read the value of  $Q$  that was produced by the same **write**( $Q$ ) operation of  $T_j$ .
  3. The transaction (if any) that performs the final **write**( $Q$ ) operation in  $S$  must also perform the final **write**( $Q$ ) operation in  $S'$ .
- A schedule  $S$  is **view serializable** if it is view equivalent to a serial schedule



# View Serializability (Cont.)

- Every conflict serializable schedule is also view serializable



- Below is a schedule which is **view-serializable** but *not* conflict serializable.

A view equivalent  
serial schedule

$T_{27}$	$T_{28}$	$T_{29}$
read (Q)	write (Q)	write (Q)
write (Q)		

$T_{27}$	$T_{28}$	$T_{29}$
read (Q)	write (Q)	write (Q)
write (Q)		


- Every view serializable schedule that is not conflict serializable has **blind writes**



# Testing for Serializability

- Consider some schedule of a set of transactions  $T_1, T_2, \dots, T_n$
- **Precedence graph** – a direct graph where the vertices are transactions(names)
  - draw an arc from  $T_i$  to  $T_j$  if the two transaction conflict, and  $T_i$  accessed the data item on which the conflict arose earlier
  - may label the arc by the item that was accessed
- Example




$T_1$	$T_2$
read(A) $A := A - 50$	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B)
write(A) read(B) $B := B + 50$ write(B)	$B := B + temp$ write(B)






# Test for Conflict Serializability

- A schedule is **conflict serializable** if and only if its **precedence graph** is **acyclic**
- Cycle-detection algorithms exist which take order  $n^2$  time, where  $n$  is the number of vertices in the graph
  - (Better algorithms take order  $n + e$  where  $e$  is the number of edges.)
- If precedence graph is acyclic, the serializability order can be obtained by a **topological sorting** of the graph
  - A linear order consistent with the partial order of the graph






# Precedence Graph for Serial Schedules



(a) Schedule 1



(b) Schedule 2

$T_1$	$T_2$
<code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code>	<code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code>

$T_1$	$T_2$
<code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code>	<code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code>





# Test for View Serializability

- The precedence graph test for conflict serializability cannot be used directly to test for view serializability
  - Extension to test for view serializability has cost exponential in the size of the precedence graph
- The problem of checking if a schedule is view serializable falls in the class of *NP-complete problems*
  - Thus existence of an efficient algorithm is *extremely* unlikely
- However, practical algorithms that just check some **sufficient conditions** for view serializability can still be used



# Recoverable Schedules

- Need to address the effect of transaction failures on concurrently running transactions
- **Recoverable schedule**
  - If a transaction  $T_j$  reads a data item previously written by a transaction  $T_i$ ,
  - then the commit operation of  $T_i$  appears before the commit operation of  $T_j$
- The following schedule is not recoverable if  $T_9$  commits immediately after the read

$T_8$	$T_9$
read (A) write (A)	
	read (A) commit
read (B)	

- DBMS must ensure that schedules are recoverable



# Cascading Rollbacks

- A single transaction failure leads to a series of transaction rollbacks
- Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

$T_{10}$	$T_{11}$	$T_{12}$
read (A) read (B) write (A)	read (A) write (A)	read (A)
abort		

- Can lead to **the undoing** of a significant amount of work



# Cascadeless Schedules

- **Cascadeless schedules** — cascading rollbacks cannot occur;
  - For each pair of transactions  $T_i$  and  $T_j$  such that  $T_j$  reads a data item previously written by  $T_i$ ,
  - The commit operation of  $T_i$  appears before the read operation of  $T_j$
- Every cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless
  
- Idea: **block other transactions until executing the commit instruction**
  - More concurrency → More cascading rollback
  - Less cascading rollback → Less concurrency



# Concurrency Control

- A database must provide a mechanism that will ensure that all possible schedules are
  - either conflict or view serializable, and
  - are recoverable and preferably cascadeless
- A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency
- Testing a schedule for serializability *after* it has executed is a little too late
- **Goal** – to develop concurrency control protocols that will assure serializability
  - Concurrency-control schemes tradeoff between the amount of concurrency they allow and the amount of overhead that they incur



# Weak Levels of Consistency

- Some applications are willing to live with weak levels of consistency, allowing schedules that are not serializable
  - Some transactions need not be serializable with respect to other transactions
    - ▶ E.g. a read-only transaction that wants to get an approximate total balance of all accounts
    - ▶ E.g. database statistics computed for query optimization can be approximate (why?)
- Tradeoff accuracy for performance



# Levels of Consistency in SQL-92

- **Serializable** — default
- **Repeatable read** — only committed records to be read, repeated reads of same record must return same value
  - However, a transaction may not be serializable – it may find some records inserted by a transaction but not find others
- **Read committed** — only committed records can be read, but successive reads of record may return different (but committed) values
- **Read uncommitted** — even uncommitted records may be read
  
- Warning: some database systems do not ensure serializable schedules by default
  - E.g. Oracle and PostgreSQL by default support a level of consistency called snapshot isolation (not part of the SQL standard)



# Transaction Definition in SQL

- Data manipulation language must include a construct for specifying the set of actions that comprise a transaction
- In SQL, **a transaction begins implicitly**
- A transaction in SQL **ends by:**
  - **Commit work** commits current transaction and begins a new one.
  - **Rollback work** causes current transaction to abort.
- In almost all database systems, by default, every SQL statement also commits implicitly if it executes successfully
  - Implicit commit can be turned off by a database directive
    - ▶ E.g. in JDBC, `connection.setAutoCommit(false);`





# End of Chapter 14

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use